

TextAugLLMEdge

June 16, 2025

```
[1]: # TextAugLLMEdge

# Supplementary Results

# Partha Pratim Ray, 15/6/2025

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Path to your CSV file (change if needed)
csv_path = r"D:\technical writing\textaugmentation\result.csv"

# Load data
df = pd.read_csv(csv_path)
print("Columns:", df.columns.tolist())
print(df.head())
```

Columns: ['timestamp', 'prompt_id', 'prompt', 'augmentation_type', 'model', 'augmented_text', 'total_duration_ns', 'load_duration_ns', 'prompt_eval_count', 'prompt_eval_duration_ns', 'eval_count', 'eval_duration_ns', 'tokens_per_second', 'levenshtein_similarity', 'jaccard_similarity', 'length_ratio', 'bleu', 'cosine_similarity', 'wer', 'char_diversity', 'type_token_ratio', 'bigram_overlap']

	timestamp	prompt_id	\
0	2025-06-15T09:50:22.829584	1	
1	2025-06-15T09:50:33.601336	2	
2	2025-06-15T09:50:41.541565	3	
3	2025-06-15T09:50:49.370159	4	
4	2025-06-15T09:50:57.445704	5	

	prompt	augmentation_type	\
0	Although agriculture remains the backbone of o...	paraphrase	
1	The rapid advancement of technology has not on...	paraphrase	
2	Despite receiving a prestigious scholarship to...	paraphrase	
3	In the aftermath of the devastating earthquake...	paraphrase	
4	The increasing adoption of renewable energy so...	paraphrase	

	model	augmented_text	\
0	qwen2:0.5b	Despite maintaining a substantial portion of i...	
1	qwen2:0.5b	The rapid technological progress over the last...	
2	qwen2:0.5b	Despite being offered an impressive scholarshi...	
3	qwen2:0.5b	In response to a severe natural disaster that ...	
4	qwen2:0.5b	Renewable energy sources like solar and wind a...	

	total_duration_ns	load_duration_ns	prompt_eval_count	\
0	10743499160	2422096170	62	
1	7918451598	53620750	58	
2	7807772750	51113778	54	
3	8053528241	49385191	62	
4	8027890100	52459502	61	

	prompt_eval_duration_ns	...	tokens_per_second	levenshtein_similarity	\
0	1835488062	...	9.253171	0.008955	
1	1415348964	...	9.304752	0.050445	
2	1302542405	...	9.300377	0.049261	
3	1530142101	...	9.269372	0.213873	
4	1501011644	...	9.268947	0.006299	

	jaccard_similarity	length_ratio	bleu	cosine_similarity	wer	\
0	0.338710	1.359155	0.216514	0.612094	1.054054	
1	0.178082	1.442029	0.023705	0.413481	1.394737	
2	0.164384	1.537500	0.018897	0.303488	1.500000	
3	0.148649	1.261438	0.024377	0.390257	1.250000	
4	0.219178	1.369403	0.084801	0.403823	1.119048	

	char_diversity	type_token_ratio	bigram_overlap
0	0.107143	0.924528	0.173333
1	0.038462	0.896552	0.032967
2	0.242424	0.981818	0.011628
3	0.068966	0.945455	0.033333
4	0.193548	0.980769	0.082353

[5 rows x 22 columns]

```
[2]: # BAR PLOT
df['total_duration_s'] = df['total_duration_ns'] / 1e9
agg = df.groupby(['augmentation_type', 'model'])['total_duration_s'].mean().
      ↪reset_index()
augmentation_order = sorted(df['augmentation_type'].unique())

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
```


```

        y='total_duration_s',
        hue='model',
        order=augmentation_order,
        palette='tab10'
    )
    for container in bar.containers:
        bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
            ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Total Duration (s)", fontsize=12)
plt.title("Mean Inference Duration by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(2.5,20.5)
plt.show()

print("Mean Total Duration (s) Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='total_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile" # Remove split=True!
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Total Duration (s)", fontsize=12)
plt.title("Distribution of Inference Duration by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

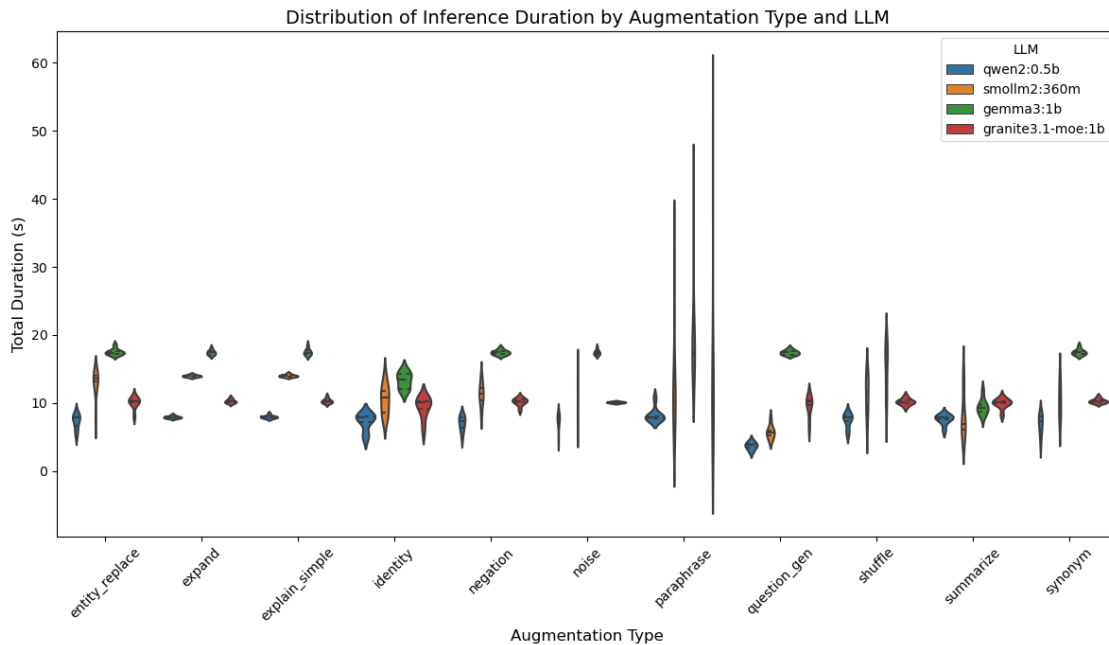
```



Mean Total Duration (s) Table:

augmentation_type	model	total_duration_s
entity_replace	gemma3:1b	17.519171
entity_replace	granite3.1-moe:1b	10.087887
entity_replace	qwen2:0.5b	7.306076
entity_replace	smollm2:360m	12.848805
expand	gemma3:1b	17.434952
expand	granite3.1-moe:1b	10.289341
expand	qwen2:0.5b	7.936603
expand	smollm2:360m	13.985258
explain_simple	gemma3:1b	17.521006
explain_simple	granite3.1-moe:1b	10.324334
explain_simple	qwen2:0.5b	7.984643
explain_simple	smollm2:360m	14.012370
identity	gemma3:1b	13.313809
identity	granite3.1-moe:1b	9.422094
identity	qwen2:0.5b	7.275286
identity	smollm2:360m	10.487829
negation	gemma3:1b	17.443342
negation	granite3.1-moe:1b	10.205199
negation	qwen2:0.5b	7.019680
negation	smollm2:360m	11.375481
noise	gemma3:1b	17.448303
noise	granite3.1-moe:1b	10.109262
noise	qwen2:0.5b	7.336289

noise	smollm2:360m	10.640014
paraphrase	gemma3:1b	19.976331
paraphrase	granite3.1-moe:1b	14.560397
paraphrase	qwen2:0.5b	8.228610
paraphrase	smollm2:360m	12.661435
question_gen	gemma3:1b	17.435018
question_gen	granite3.1-moe:1b	9.610098
question_gen	qwen2:0.5b	3.707023
question_gen	smollm2:360m	5.754444
shuffle	gemma3:1b	15.157816
shuffle	granite3.1-moe:1b	10.209583
shuffle	qwen2:0.5b	7.451169
shuffle	smollm2:360m	11.327010
summarize	gemma3:1b	9.429700
summarize	granite3.1-moe:1b	9.956269
summarize	qwen2:0.5b	7.627685
summarize	smollm2:360m	8.226404
synonym	gemma3:1b	17.496148
synonym	granite3.1-moe:1b	10.318432
synonym	qwen2:0.5b	6.892995
synonym	smollm2:360m	10.327454



```
[3]: # BAR PLOT
df['load_duration_s'] = df['load_duration_ns'] / 1e9
agg = df.groupby(['augmentation_type', 'model'])['load_duration_s'].mean().
      ↪reset_index()
```

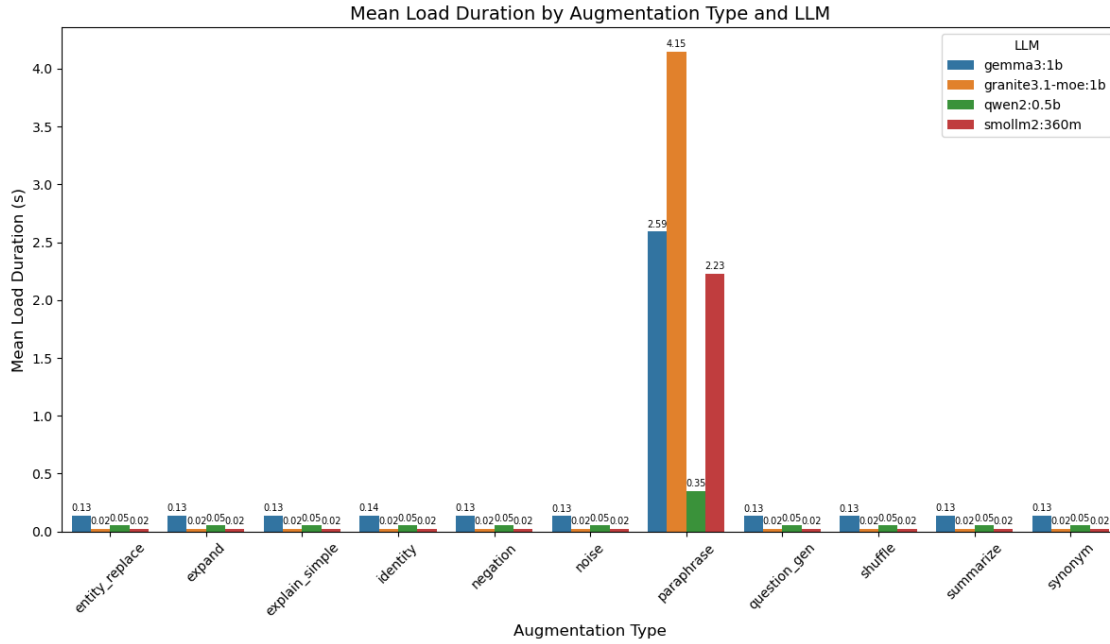
```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='load_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Load Duration (s)", fontsize=12)
plt.title("Mean Load Duration by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Load Duration (s) Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='load_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Load Duration (s)", fontsize=12)
plt.title("Distribution of Load Duration by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

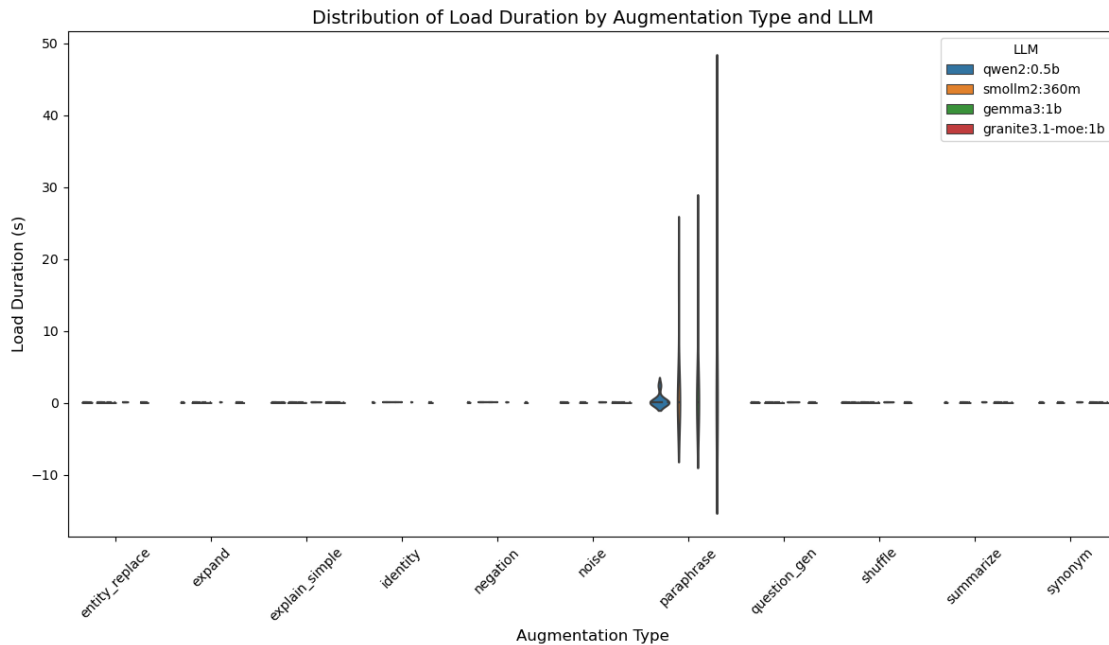
```



Mean Load Duration (s) Table:

augmentation_type	model	load_duration_s
entity_replace	gemma3:1b	0.134435
entity_replace	granite3.1-moe:1b	0.022135
entity_replace	qwen2:0.5b	0.051694
entity_replace	smollm2:360m	0.020649
expand	gemma3:1b	0.134711
expand	granite3.1-moe:1b	0.020546
expand	qwen2:0.5b	0.050972
expand	smollm2:360m	0.020424
explain_simple	gemma3:1b	0.134660
explain_simple	granite3.1-moe:1b	0.021282
explain_simple	qwen2:0.5b	0.051691
explain_simple	smollm2:360m	0.021615
identity	gemma3:1b	0.135201
identity	granite3.1-moe:1b	0.021608
identity	qwen2:0.5b	0.051934
identity	smollm2:360m	0.020424
negation	gemma3:1b	0.134433
negation	granite3.1-moe:1b	0.021027
negation	qwen2:0.5b	0.052432
negation	smollm2:360m	0.020509
noise	gemma3:1b	0.133162
noise	granite3.1-moe:1b	0.020497
noise	qwen2:0.5b	0.051306

noise	smollm2:360m	0.022082
paraphrase	gemma3:1b	2.591088
paraphrase	granite3.1-moe:1b	4.148439
paraphrase	qwen2:0.5b	0.348289
paraphrase	smollm2:360m	2.230602
question_gen	gemma3:1b	0.133863
question_gen	granite3.1-moe:1b	0.021682
question_gen	qwen2:0.5b	0.052184
question_gen	smollm2:360m	0.020965
shuffle	gemma3:1b	0.132988
shuffle	granite3.1-moe:1b	0.021605
shuffle	qwen2:0.5b	0.051072
shuffle	smollm2:360m	0.021318
summarize	gemma3:1b	0.134930
summarize	granite3.1-moe:1b	0.020477
summarize	qwen2:0.5b	0.052047
summarize	smollm2:360m	0.021119
synonym	gemma3:1b	0.134818
synonym	granite3.1-moe:1b	0.021068
synonym	qwen2:0.5b	0.052437
synonym	smollm2:360m	0.022486



```
[4]: # BAR PLOT
agg = df.groupby(['augmentation_type', 'model'])['prompt_eval_count'].mean().
      reset_index()
```



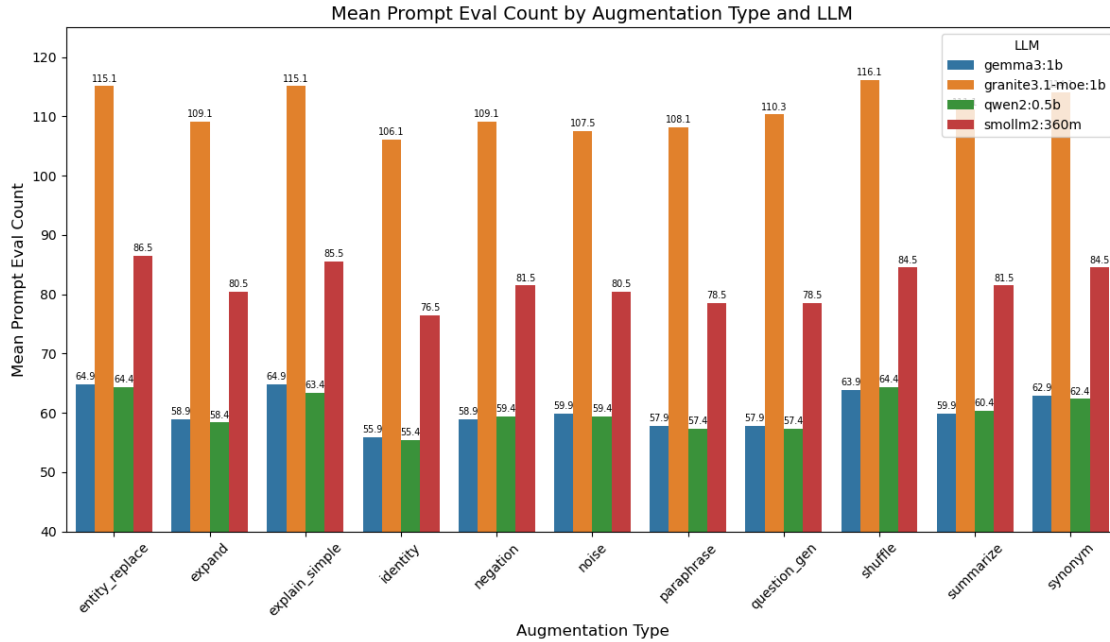
```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='prompt_eval_count',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.1f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Prompt Eval Count", fontsize=12)
plt.title("Mean Prompt Eval Count by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(40,125)
plt.show()

print("Mean Prompt Eval Count Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='prompt_eval_count',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Prompt Eval Count", fontsize=12)
plt.title("Distribution of Prompt Eval Count by Augmentation Type and LLM",
    ↪fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

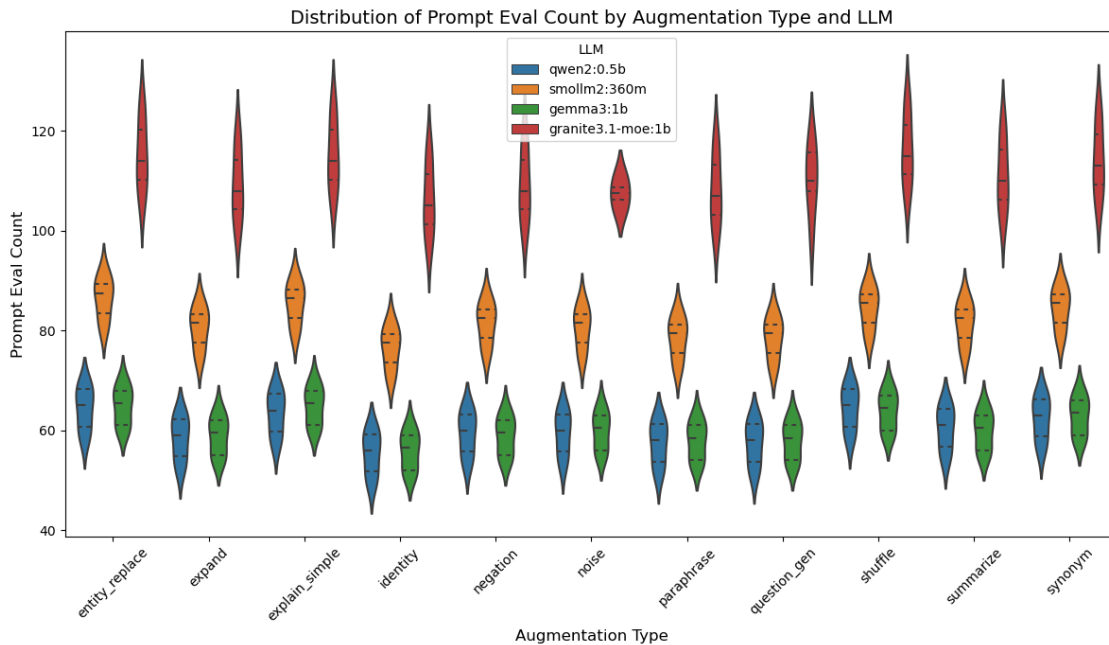
```



Mean Prompt Eval Count Table:

augmentation_type	model	prompt_eval_count
entity_replace	gemma3:1b	64.875000
entity_replace	granite3.1-moe:1b	115.125000
entity_replace	qwen2:0.5b	64.375000
entity_replace	smollm2:360m	86.500000
expand	gemma3:1b	58.875000
expand	granite3.1-moe:1b	109.125000
expand	qwen2:0.5b	58.375000
expand	smollm2:360m	80.500000
explain_simple	gemma3:1b	64.875000
explain_simple	granite3.1-moe:1b	115.125000
explain_simple	qwen2:0.5b	63.375000
explain_simple	smollm2:360m	85.500000
identity	gemma3:1b	55.875000
identity	granite3.1-moe:1b	106.125000
identity	qwen2:0.5b	55.375000
identity	smollm2:360m	76.500000
negation	gemma3:1b	58.875000
negation	granite3.1-moe:1b	109.125000
negation	qwen2:0.5b	59.375000
negation	smollm2:360m	81.500000
noise	gemma3:1b	59.875000
noise	granite3.1-moe:1b	107.500000
noise	qwen2:0.5b	59.375000

noise	smollm2:360m	80.500000
paraphrase	gemma3:1b	57.875000
paraphrase	granite3.1-moe:1b	108.125000
paraphrase	qwen2:0.5b	57.375000
paraphrase	smollm2:360m	78.500000
question_gen	gemma3:1b	57.875000
question_gen	granite3.1-moe:1b	110.333333
question_gen	qwen2:0.5b	57.375000
question_gen	smollm2:360m	78.500000
shuffle	gemma3:1b	63.875000
shuffle	granite3.1-moe:1b	116.125000
shuffle	qwen2:0.5b	64.375000
shuffle	smollm2:360m	84.500000
summarize	gemma3:1b	59.875000
summarize	granite3.1-moe:1b	111.125000
summarize	qwen2:0.5b	60.375000
summarize	smollm2:360m	81.500000
synonym	gemma3:1b	62.875000
synonym	granite3.1-moe:1b	114.125000
synonym	qwen2:0.5b	62.375000
synonym	smollm2:360m	84.500000



```
[5]: # BAR PLOT
df['prompt_eval_duration_s'] = df['prompt_eval_duration_ns'] / 1e9
agg = df.groupby(['augmentation_type', 'model'])['prompt_eval_duration_s'].
      ↪mean().reset_index()
```

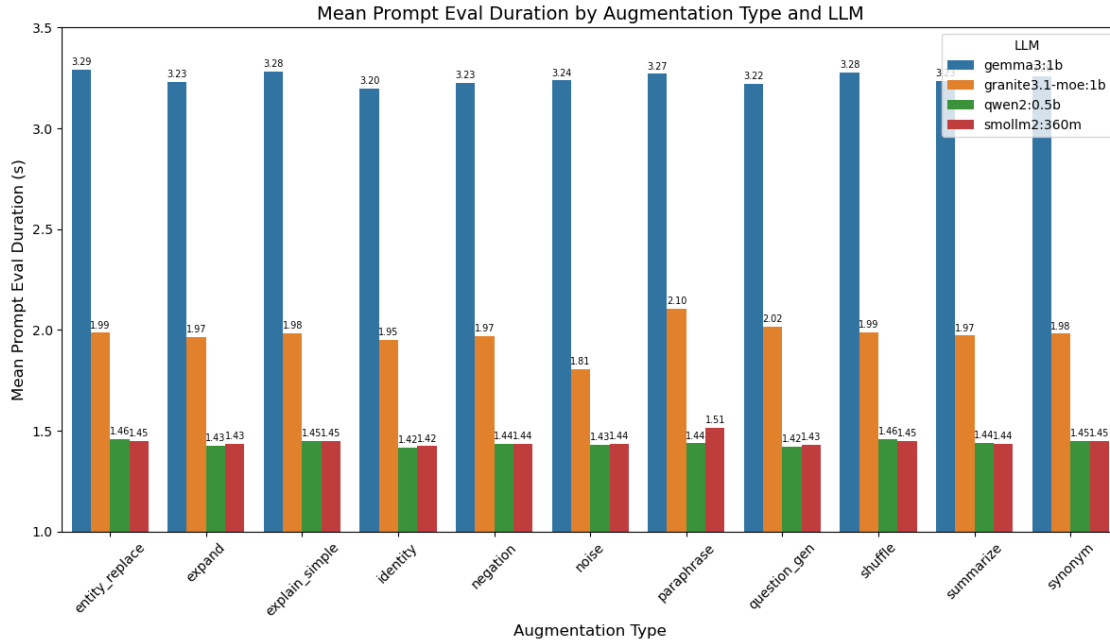
```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='prompt_eval_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Prompt Eval Duration (s)", fontsize=12)
plt.title("Mean Prompt Eval Duration by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(1,3.5)
plt.show()

print("Mean Prompt Eval Duration (s) Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='prompt_eval_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Prompt Eval Duration (s)", fontsize=12)
plt.title("Distribution of Prompt Eval Duration by Augmentation Type and LLM",
    ↪fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

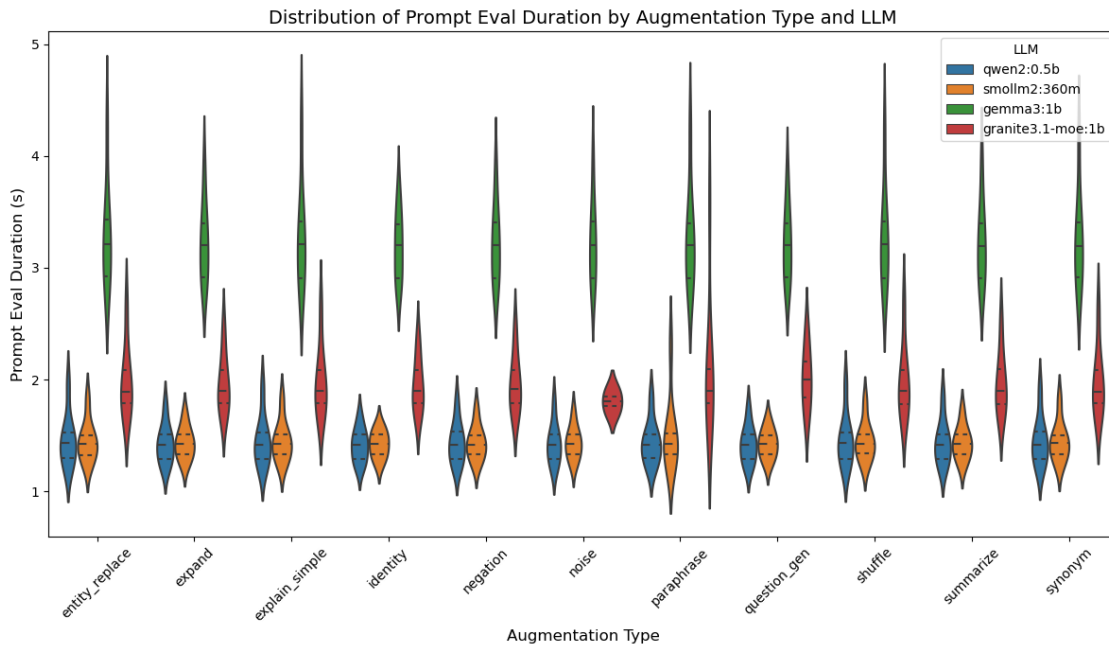
```



Mean Prompt Eval Duration (s) Table:

augmentation_type	model	prompt_eval_duration_s
entity_replace	gemma3:1b	3.289920
entity_replace	granite3.1-moe:1b	1.986179
entity_replace	qwen2:0.5b	1.458739
entity_replace	smollm2:360m	1.448993
expand	gemma3:1b	3.230149
expand	granite3.1-moe:1b	1.965222
expand	qwen2:0.5b	1.425312
expand	smollm2:360m	1.434685
explain_simple	gemma3:1b	3.282262
explain_simple	granite3.1-moe:1b	1.984986
explain_simple	qwen2:0.5b	1.450086
explain_simple	smollm2:360m	1.450368
identity	gemma3:1b	3.198626
identity	granite3.1-moe:1b	1.952127
identity	qwen2:0.5b	1.415143
identity	smollm2:360m	1.423411
negation	gemma3:1b	3.226121
negation	granite3.1-moe:1b	1.969167
negation	qwen2:0.5b	1.436152
negation	smollm2:360m	1.435887
noise	gemma3:1b	3.237751
noise	granite3.1-moe:1b	1.806109
noise	qwen2:0.5b	1.429667

noise	smollm2:360m	1.435769
paraphrase	gemma3:1b	3.269960
paraphrase	granite3.1-moe:1b	2.104556
paraphrase	qwen2:0.5b	1.437472
paraphrase	smollm2:360m	1.514472
question_gen	gemma3:1b	3.219902
question_gen	granite3.1-moe:1b	2.017061
question_gen	qwen2:0.5b	1.422472
question_gen	smollm2:360m	1.428817
shuffle	gemma3:1b	3.277141
shuffle	granite3.1-moe:1b	1.988787
shuffle	qwen2:0.5b	1.459124
shuffle	smollm2:360m	1.448311
summarize	gemma3:1b	3.234136
summarize	granite3.1-moe:1b	1.971836
summarize	qwen2:0.5b	1.438362
summarize	smollm2:360m	1.436426
synonym	gemma3:1b	3.260029
synonym	granite3.1-moe:1b	1.981388
synonym	qwen2:0.5b	1.450563
synonym	smollm2:360m	1.449356



```
[6]: # BAR PLOT
agg = df.groupby(['augmentation_type', 'model'])['eval_count'].mean().
      reset_index()
```

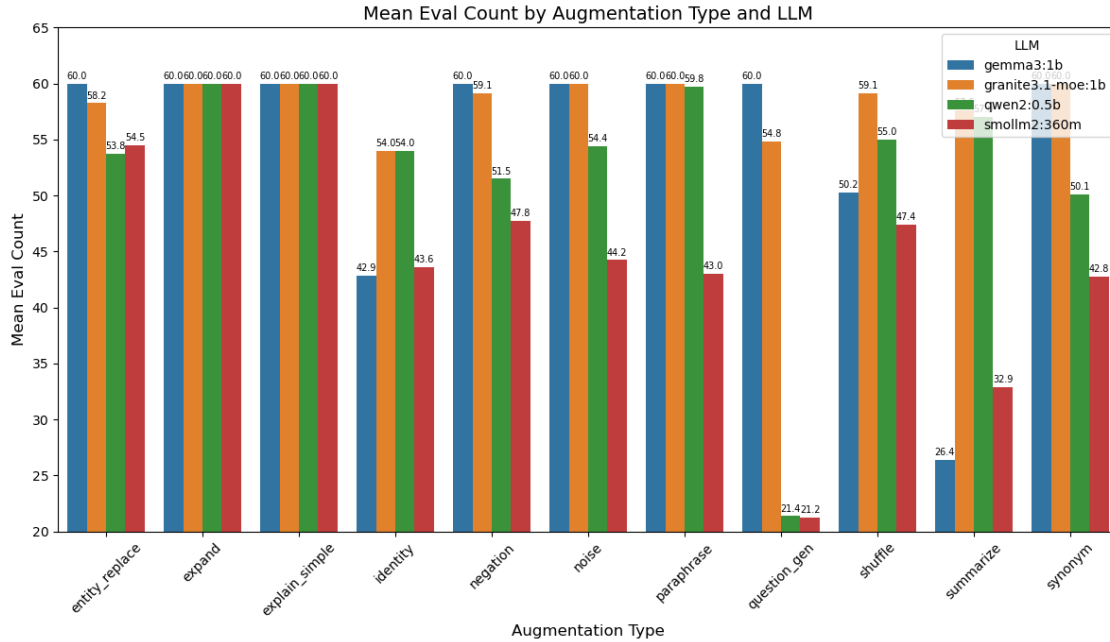
```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='eval_count',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.1f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Eval Count", fontsize=12)
plt.title("Mean Eval Count by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(20,65)
plt.show()

print("Mean Eval Count Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='eval_count',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Eval Count", fontsize=12)
plt.title("Distribution of Eval Count by Augmentation Type and LLM",
    ↪fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

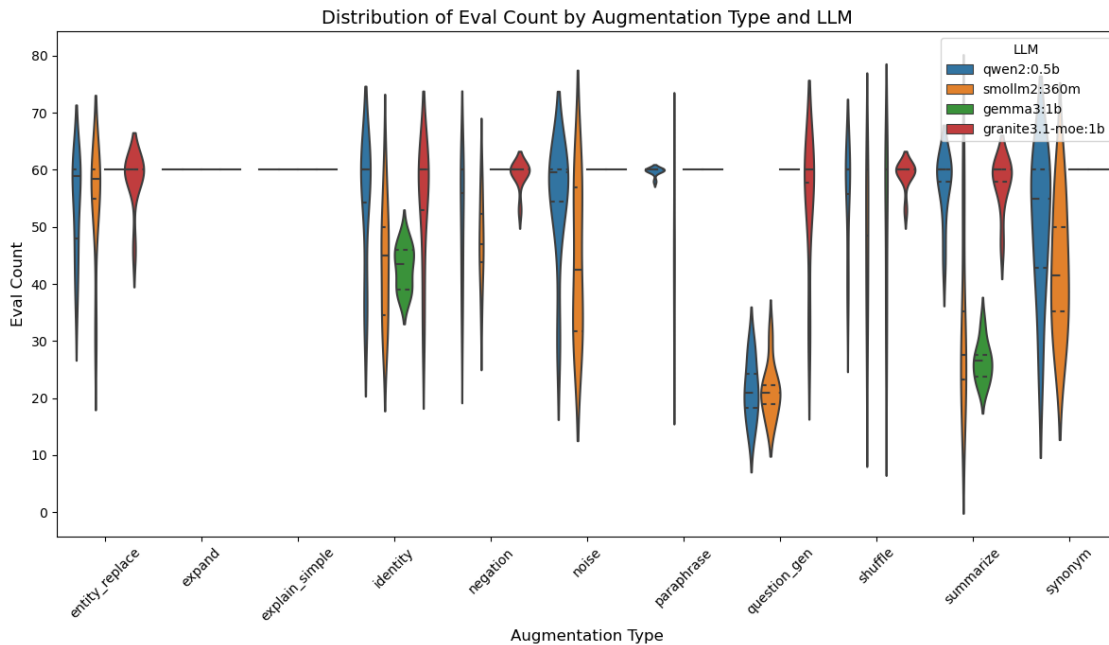
```



Mean Eval Count Table:

augmentation_type	model	eval_count
entity_replace	gemma3:1b	60.000000
entity_replace	granite3.1-moe:1b	58.250000
entity_replace	qwen2:0.5b	53.750000
entity_replace	smollm2:360m	54.500000
expand	gemma3:1b	60.000000
expand	granite3.1-moe:1b	60.000000
expand	qwen2:0.5b	60.000000
expand	smollm2:360m	60.000000
explain_simple	gemma3:1b	60.000000
explain_simple	granite3.1-moe:1b	60.000000
explain_simple	qwen2:0.5b	60.000000
explain_simple	smollm2:360m	60.000000
identity	gemma3:1b	42.875000
identity	granite3.1-moe:1b	54.000000
identity	qwen2:0.5b	54.000000
identity	smollm2:360m	43.625000
negation	gemma3:1b	60.000000
negation	granite3.1-moe:1b	59.125000
negation	qwen2:0.5b	51.500000
negation	smollm2:360m	47.750000
noise	gemma3:1b	60.000000
noise	granite3.1-moe:1b	60.000000
noise	qwen2:0.5b	54.375000

noise	smollm2:360m	44.250000
paraphrase	gemma3:1b	60.000000
paraphrase	granite3.1-moe:1b	60.000000
paraphrase	qwen2:0.5b	59.750000
paraphrase	smollm2:360m	43.000000
question_gen	gemma3:1b	60.000000
question_gen	granite3.1-moe:1b	54.833333
question_gen	qwen2:0.5b	21.375000
question_gen	smollm2:360m	21.250000
shuffle	gemma3:1b	50.250000
shuffle	granite3.1-moe:1b	59.125000
shuffle	qwen2:0.5b	55.000000
shuffle	smollm2:360m	47.375000
summarize	gemma3:1b	26.375000
summarize	granite3.1-moe:1b	57.625000
summarize	qwen2:0.5b	57.000000
summarize	smollm2:360m	32.875000
synonym	gemma3:1b	60.000000
synonym	granite3.1-moe:1b	60.000000
synonym	qwen2:0.5b	50.125000
synonym	smollm2:360m	42.750000



```
[7]: # BAR PLOT
df['eval_duration_s'] = df['eval_duration_ns'] / 1e9
agg = df.groupby(['augmentation_type', 'model'])['eval_duration_s'].mean().
      ↪reset_index()
```

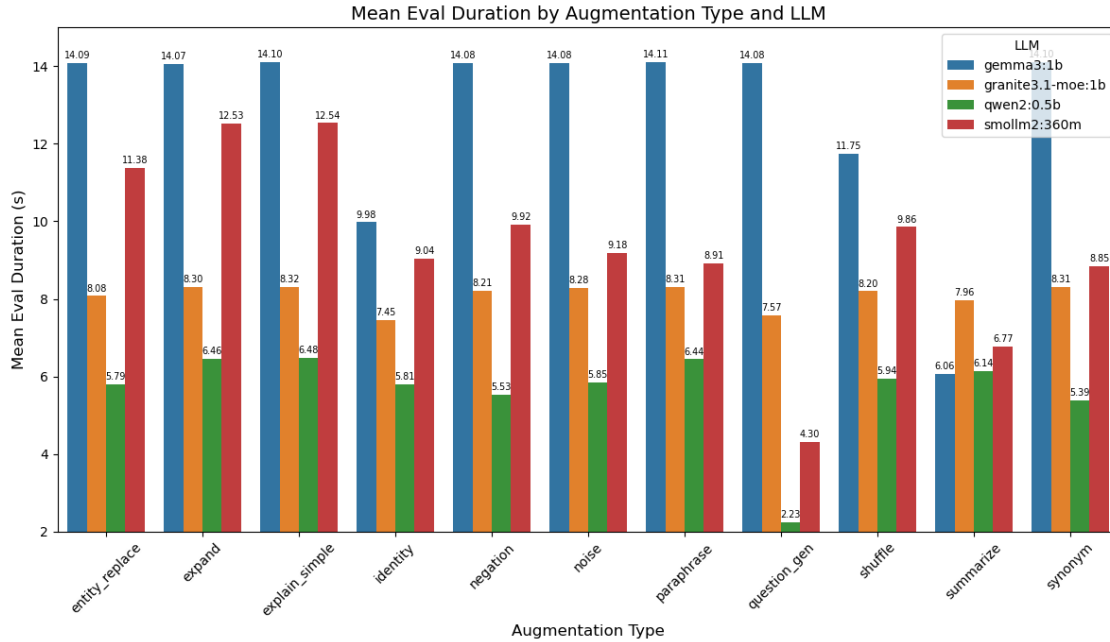
```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='eval_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Eval Duration (s)", fontsize=12)
plt.title("Mean Eval Duration by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(2,15)
plt.show()

print("Mean Eval Duration (s) Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='eval_duration_s',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Eval Duration (s)", fontsize=12)
plt.title("Distribution of Eval Duration by Augmentation Type and LLM",
    ↪fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()


```



Mean Eval Duration (s) Table:

augmentation_type	model	eval_duration_s
entity_replace	gemma3:1b	14.093728
entity_replace	granite3.1-moe:1b	8.077532
entity_replace	qwen2:0.5b	5.794070
entity_replace	smollm2:360m	11.377713
expand	gemma3:1b	14.069005
expand	granite3.1-moe:1b	8.301877
expand	qwen2:0.5b	6.459140
expand	smollm2:360m	12.528693
explain_simple	gemma3:1b	14.102958
explain_simple	granite3.1-moe:1b	8.316195
explain_simple	qwen2:0.5b	6.481418
explain_simple	smollm2:360m	12.538618
identity	gemma3:1b	9.978718
identity	granite3.1-moe:1b	7.446409
identity	qwen2:0.5b	5.806615
identity	smollm2:360m	9.042256
negation	gemma3:1b	14.081699
negation	granite3.1-moe:1b	8.212668
negation	qwen2:0.5b	5.529797
negation	smollm2:360m	9.917460
noise	gemma3:1b	14.076263
noise	granite3.1-moe:1b	8.280945
noise	qwen2:0.5b	5.853783

noise	smollm2:360m	9.180232
paraphrase	gemma3:1b	14.113944
paraphrase	granite3.1-moe:1b	8.305655
paraphrase	qwen2:0.5b	6.441398
paraphrase	smollm2:360m	8.914208
question_gen	gemma3:1b	14.080092
question_gen	granite3.1-moe:1b	7.569508
question_gen	qwen2:0.5b	2.231112
question_gen	smollm2:360m	4.302042
shuffle	gemma3:1b	11.746517
shuffle	granite3.1-moe:1b	8.197225
shuffle	qwen2:0.5b	5.939719
shuffle	smollm2:360m	9.855855
summarize	gemma3:1b	6.059394
summarize	granite3.1-moe:1b	7.962053
summarize	qwen2:0.5b	6.136115
summarize	smollm2:360m	6.767067
synonym	gemma3:1b	14.100174
synonym	granite3.1-moe:1b	8.314038
synonym	qwen2:0.5b	5.388443
synonym	smollm2:360m	8.853510



```
[8]: # BAR PLOT
agg = df.groupby(['augmentation_type', 'model'])['tokens_per_second'].mean().
      ↪reset_index()
```

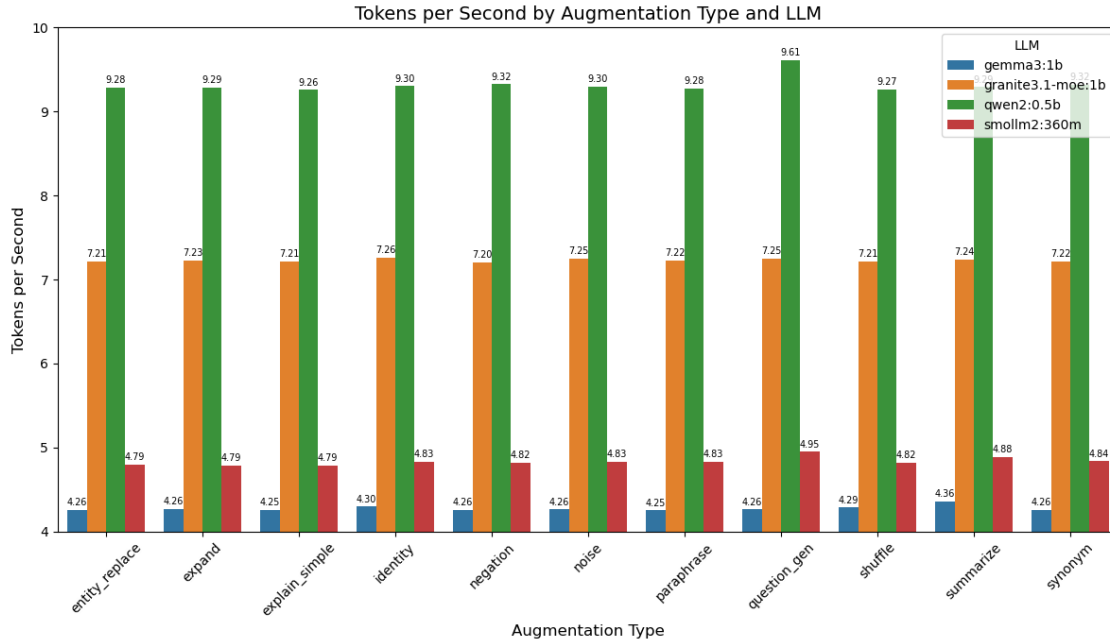
```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='tokens_per_second',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Tokens per Second", fontsize=12)
plt.title("Tokens per Second by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(4,10)
plt.show()

print("Tokens per Second Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
violin = sns.violinplot(
    data=df,
    x='augmentation_type',
    y='tokens_per_second',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Tokens per Second", fontsize=12)
plt.title("Distribution of Tokens per Second by Augmentation Type and LLM",
    ↪fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

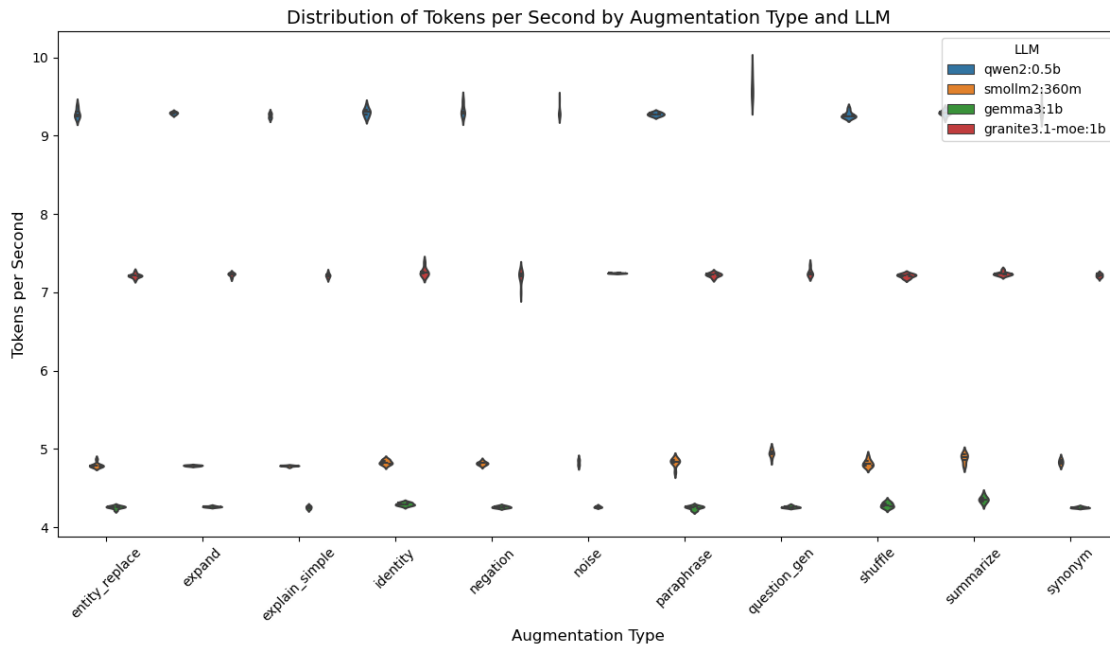
```



Tokens per Second Table:

augmentation_type	model	tokens_per_second
entity_replace	gemma3:1b	4.257286
entity_replace	granite3.1-moe:1b	7.212979
entity_replace	qwen2:0.5b	9.283693
entity_replace	smollm2:360m	4.794924
expand	gemma3:1b	4.264705
expand	granite3.1-moe:1b	7.227345
expand	qwen2:0.5b	9.289185
expand	smollm2:360m	4.789015
explain_simple	gemma3:1b	4.254498
explain_simple	granite3.1-moe:1b	7.214922
explain_simple	qwen2:0.5b	9.257323
explain_simple	smollm2:360m	4.785226
identity	gemma3:1b	4.297115
identity	granite3.1-moe:1b	7.261247
identity	qwen2:0.5b	9.301166
identity	smollm2:360m	4.830443
negation	gemma3:1b	4.260880
negation	granite3.1-moe:1b	7.201331
negation	qwen2:0.5b	9.324690
negation	smollm2:360m	4.817815
noise	gemma3:1b	4.262515
noise	granite3.1-moe:1b	7.245551
noise	qwen2:0.5b	9.299473

noise	smollm2:360m	4.829795
paraphrase	gemma3:1b	4.251231
paraphrase	granite3.1-moe:1b	7.224071
paraphrase	qwen2:0.5b	9.276019
paraphrase	smollm2:360m	4.830580
question_gen	gemma3:1b	4.261365
question_gen	granite3.1-moe:1b	7.252370
question_gen	qwen2:0.5b	9.610416
question_gen	smollm2:360m	4.948760
shuffle	gemma3:1b	4.285316
shuffle	granite3.1-moe:1b	7.212863
shuffle	qwen2:0.5b	9.265605
shuffle	smollm2:360m	4.817816
summarize	gemma3:1b	4.356848
summarize	granite3.1-moe:1b	7.238769
summarize	qwen2:0.5b	9.291481
summarize	smollm2:360m	4.882819
synonym	gemma3:1b	4.255283
synonym	granite3.1-moe:1b	7.216768
synonym	qwen2:0.5b	9.320334
synonym	smollm2:360m	4.836850



```
[9]: agg = df.groupby(['augmentation_type', 'model'])['levenshtein_similarity'].
      ↪mean().reset_index()
      augmentation_order = sorted(df['augmentation_type'].unique())
```


```

plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='levenshtein_similarity',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Levenshtein Similarity", fontsize=12)
plt.title("Levenshtein Similarity by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Levenshtein Similarity Table:\n")
print(agg.to_string(index=False))

# VIOLIN PLOT
plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='levenshtein_similarity',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Levenshtein Similarity", fontsize=12)
plt.title("Distribution of Levenshtein Similarity by Augmentation Type and ↪
    ↪ LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

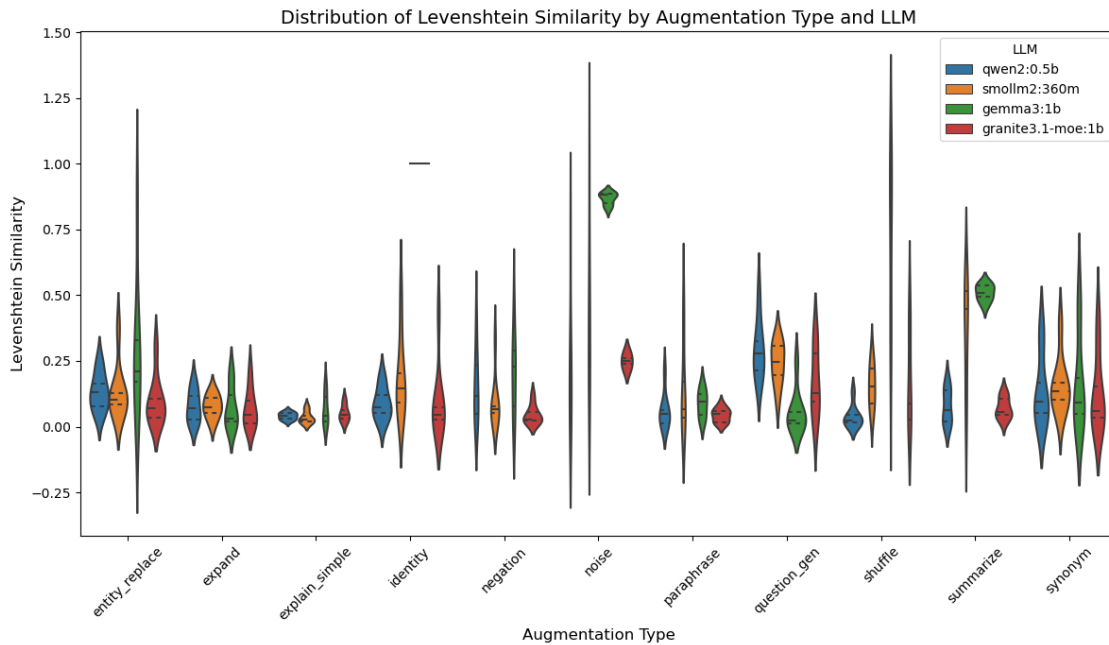
```

Mean Levenshtein Similarity Table:

augmentation_type	model	levenshtein_similarity
entity_replace	gemma3:1b	0.290348
entity_replace	granite3.1-moe:1b	0.097073
entity_replace	qwen2:0.5b	0.132908
entity_replace	smollm2:360m	0.132863
expand	gemma3:1b	0.075134
expand	granite3.1-moe:1b	0.069932
expand	qwen2:0.5b	0.078037
expand	smollm2:360m	0.077145
explain_simple	gemma3:1b	0.067784
explain_simple	granite3.1-moe:1b	0.051873
explain_simple	qwen2:0.5b	0.040759
explain_simple	smollm2:360m	0.036005
identity	gemma3:1b	1.000000
identity	granite3.1-moe:1b	0.094511
identity	qwen2:0.5b	0.089047
identity	smollm2:360m	0.190436
negation	gemma3:1b	0.205183
negation	granite3.1-moe:1b	0.045812
negation	qwen2:0.5b	0.154084
negation	smollm2:360m	0.093264
noise	gemma3:1b	0.868107
noise	granite3.1-moe:1b	0.249687
noise	qwen2:0.5b	0.214436

noise	smollm2:360m	0.400869
paraphrase	gemma3:1b	0.088593
paraphrase	granite3.1-moe:1b	0.044057
paraphrase	qwen2:0.5b	0.060148
paraphrase	smollm2:360m	0.144007
question_gen	gemma3:1b	0.057572
question_gen	granite3.1-moe:1b	0.168588
question_gen	qwen2:0.5b	0.293763
question_gen	smollm2:360m	0.241747
shuffle	gemma3:1b	0.721207
shuffle	granite3.1-moe:1b	0.179992
shuffle	qwen2:0.5b	0.038025
shuffle	smollm2:360m	0.154309
summarize	gemma3:1b	0.509782
summarize	granite3.1-moe:1b	0.071765
summarize	qwen2:0.5b	0.078576
summarize	smollm2:360m	0.367185
synonym	gemma3:1b	0.163364
synonym	granite3.1-moe:1b	0.126834
synonym	qwen2:0.5b	0.138962
synonym	smollm2:360m	0.156136



```
[10]: agg = df.groupby(['augmentation_type', 'model'])['jaccard_similarity'].mean().
      ↪reset_index()

plt.figure(figsize=(12,7))
```

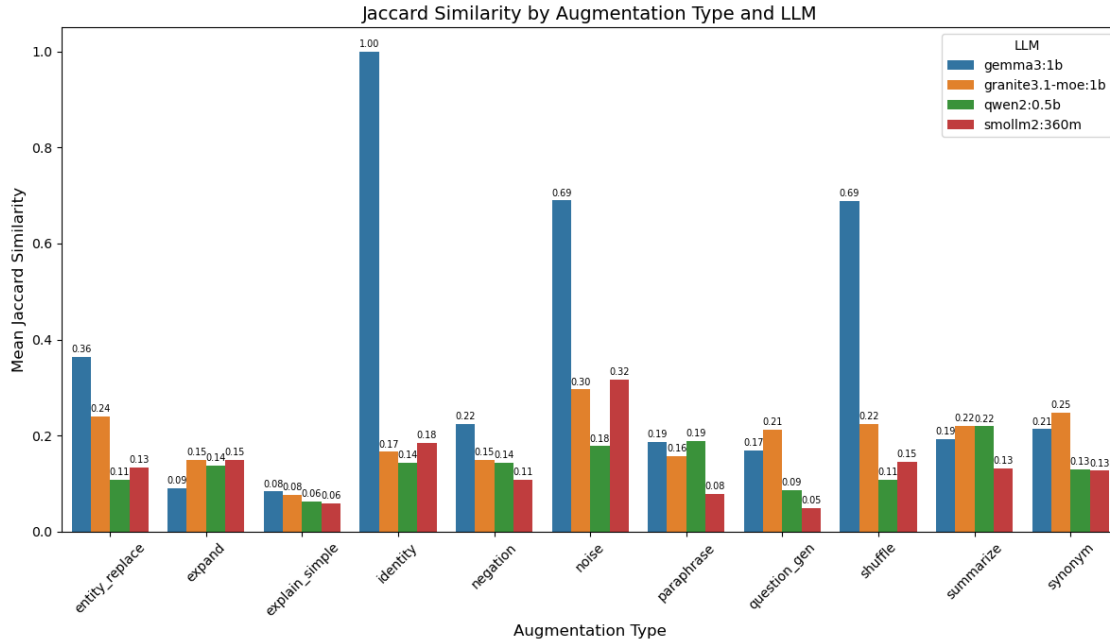
```

bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='jaccard_similarity',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Jaccard Similarity", fontsize=12)
plt.title("Jaccard Similarity by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Jaccard Similarity Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='jaccard_similarity',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Jaccard Similarity", fontsize=12)
plt.title("Distribution of Jaccard Similarity by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

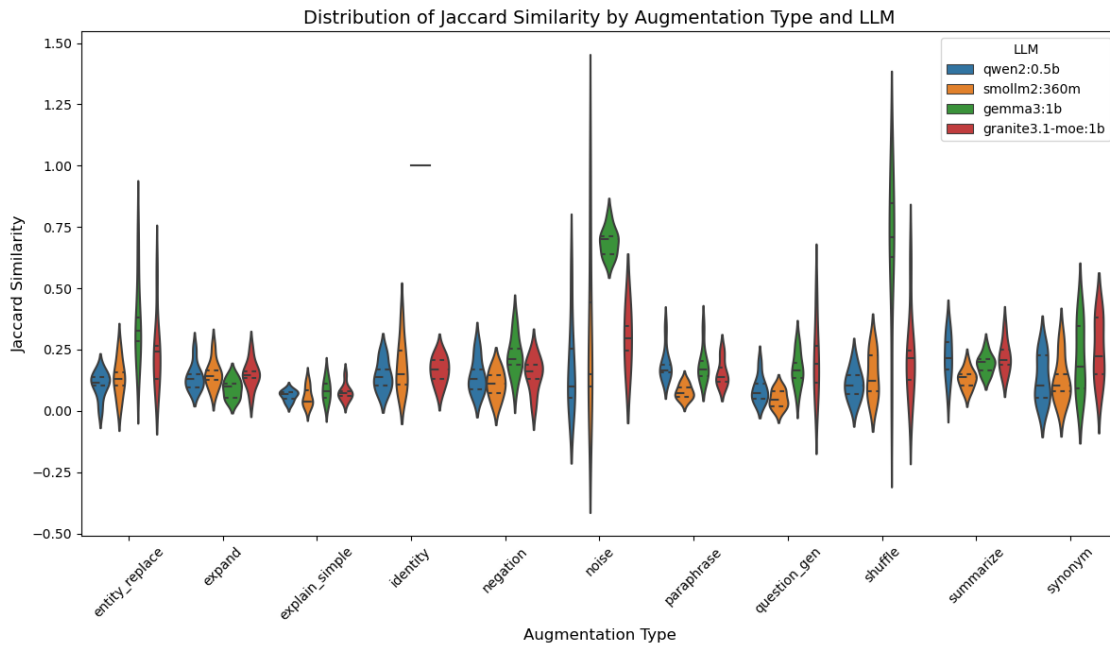
```



Mean Jaccard Similarity Table:

augmentation_type	model	jaccard_similarity
entity_replace	gemma3:1b	0.364388
entity_replace	granite3.1-moe:1b	0.239889
entity_replace	qwen2:0.5b	0.107304
entity_replace	smollm2:360m	0.132787
expand	gemma3:1b	0.090450
expand	granite3.1-moe:1b	0.148763
expand	qwen2:0.5b	0.136878
expand	smollm2:360m	0.149658
explain_simple	gemma3:1b	0.083227
explain_simple	granite3.1-moe:1b	0.076358
explain_simple	qwen2:0.5b	0.062661
explain_simple	smollm2:360m	0.058572
identity	gemma3:1b	1.000000
identity	granite3.1-moe:1b	0.166029
identity	qwen2:0.5b	0.143552
identity	smollm2:360m	0.184490
negation	gemma3:1b	0.223519
negation	granite3.1-moe:1b	0.149267
negation	qwen2:0.5b	0.143467
negation	smollm2:360m	0.108570
noise	gemma3:1b	0.689598
noise	granite3.1-moe:1b	0.295714
noise	qwen2:0.5b	0.177881

noise	smollm2:360m	0.316545
paraphrase	gemma3:1b	0.186556
paraphrase	granite3.1-moe:1b	0.156475
paraphrase	qwen2:0.5b	0.188939
paraphrase	smollm2:360m	0.078176
question_gen	gemma3:1b	0.169422
question_gen	granite3.1-moe:1b	0.211565
question_gen	qwen2:0.5b	0.086832
question_gen	smollm2:360m	0.049405
shuffle	gemma3:1b	0.688624
shuffle	granite3.1-moe:1b	0.224040
shuffle	qwen2:0.5b	0.108495
shuffle	smollm2:360m	0.146016
summarize	gemma3:1b	0.191970
summarize	granite3.1-moe:1b	0.219977
summarize	qwen2:0.5b	0.219177
summarize	smollm2:360m	0.131202
synonym	gemma3:1b	0.213119
synonym	granite3.1-moe:1b	0.247364
synonym	qwen2:0.5b	0.129532
synonym	smollm2:360m	0.126717



```
[11]: agg = df.groupby(['augmentation_type', 'model'])['length_ratio'].mean().
      ↪reset_index()

plt.figure(figsize=(12,7))
```

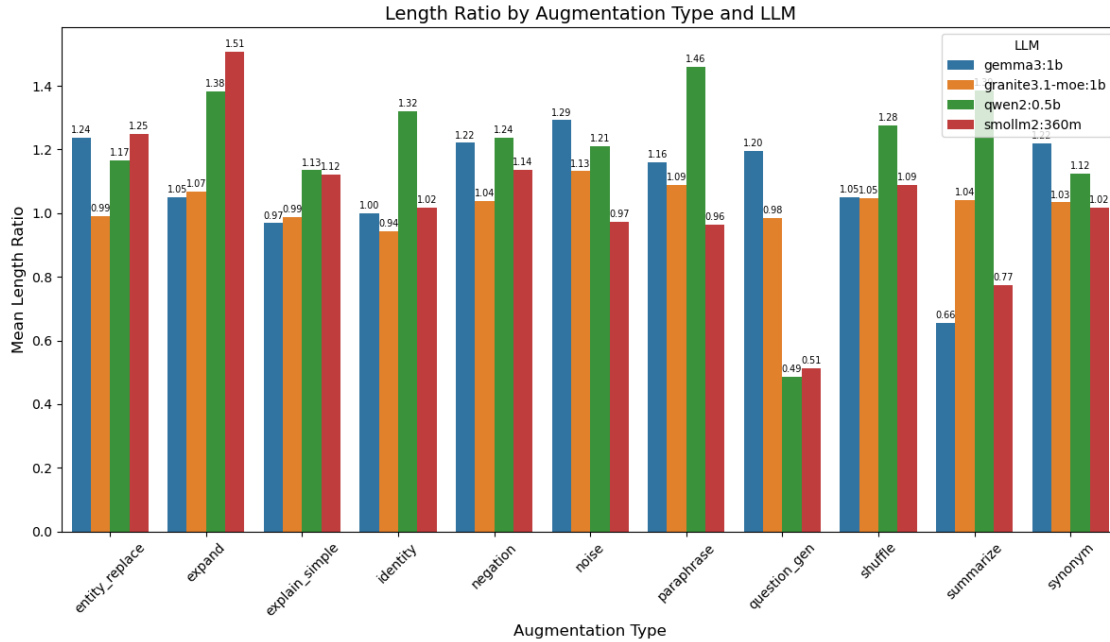
```

bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='length_ratio',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Length Ratio", fontsize=12)
plt.title("Length Ratio by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Length Ratio Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='length_ratio',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Length Ratio", fontsize=12)
plt.title("Distribution of Length Ratio by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(.4, 1.6)
plt.show()

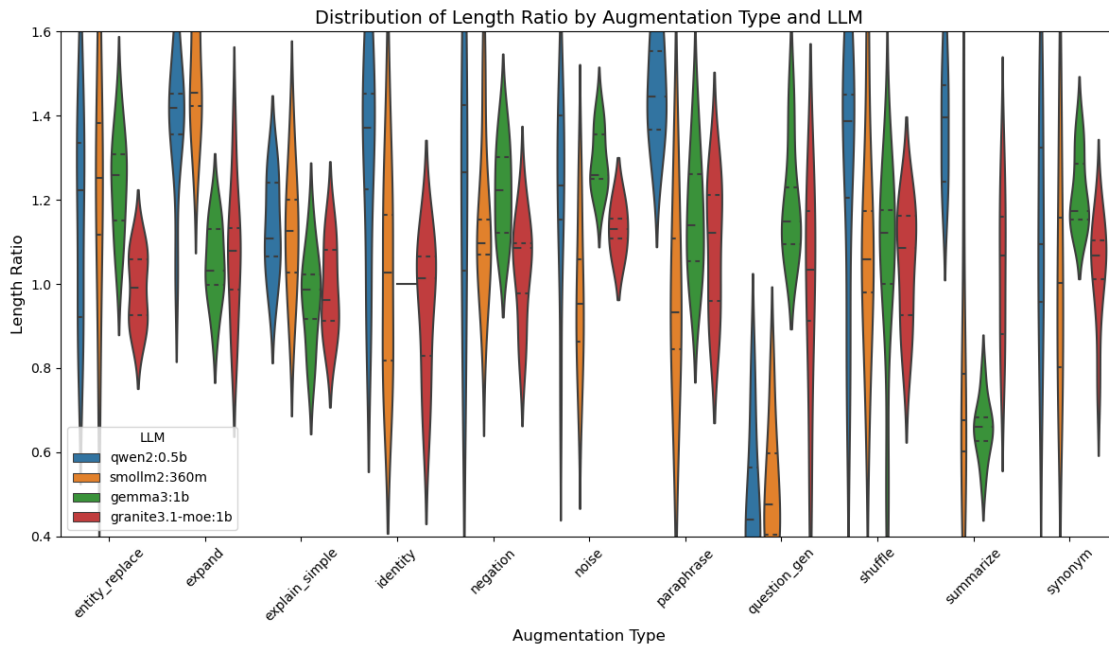
```



Mean Length Ratio Table:

augmentation_type	model	length_ratio
entity_replace	gemma3:1b	1.236984
entity_replace	granite3.1-moe:1b	0.990239
entity_replace	qwen2:0.5b	1.166508
entity_replace	smollm2:360m	1.249348
expand	gemma3:1b	1.049079
expand	granite3.1-moe:1b	1.067523
expand	qwen2:0.5b	1.382795
expand	smollm2:360m	1.507696
explain_simple	gemma3:1b	0.968811
explain_simple	granite3.1-moe:1b	0.988082
explain_simple	qwen2:0.5b	1.134455
explain_simple	smollm2:360m	1.122041
identity	gemma3:1b	1.000000
identity	granite3.1-moe:1b	0.943784
identity	qwen2:0.5b	1.320280
identity	smollm2:360m	1.016836
negation	gemma3:1b	1.220938
negation	granite3.1-moe:1b	1.038055
negation	qwen2:0.5b	1.237072
negation	smollm2:360m	1.137183
noise	gemma3:1b	1.292165
noise	granite3.1-moe:1b	1.131486
noise	qwen2:0.5b	1.211395

noise	smollm2:360m	0.973453
paraphrase	gemma3:1b	1.159645
paraphrase	granite3.1-moe:1b	1.089286
paraphrase	qwen2:0.5b	1.458942
paraphrase	smollm2:360m	0.963935
question_gen	gemma3:1b	1.195989
question_gen	granite3.1-moe:1b	0.984485
question_gen	qwen2:0.5b	0.487344
question_gen	smollm2:360m	0.513460
shuffle	gemma3:1b	1.050232
shuffle	granite3.1-moe:1b	1.047250
shuffle	qwen2:0.5b	1.275040
shuffle	smollm2:360m	1.087972
summarize	gemma3:1b	0.656102
summarize	granite3.1-moe:1b	1.041472
summarize	qwen2:0.5b	1.385018
summarize	smollm2:360m	0.773369
synonym	gemma3:1b	1.219601
synonym	granite3.1-moe:1b	1.033883
synonym	qwen2:0.5b	1.123736
synonym	smollm2:360m	1.018208



```
[12]: agg = df.groupby(['augmentation_type', 'model'])['bleu'].mean().reset_index()

plt.figure(figsize=(12,7))
bar = sns.barplot(
```



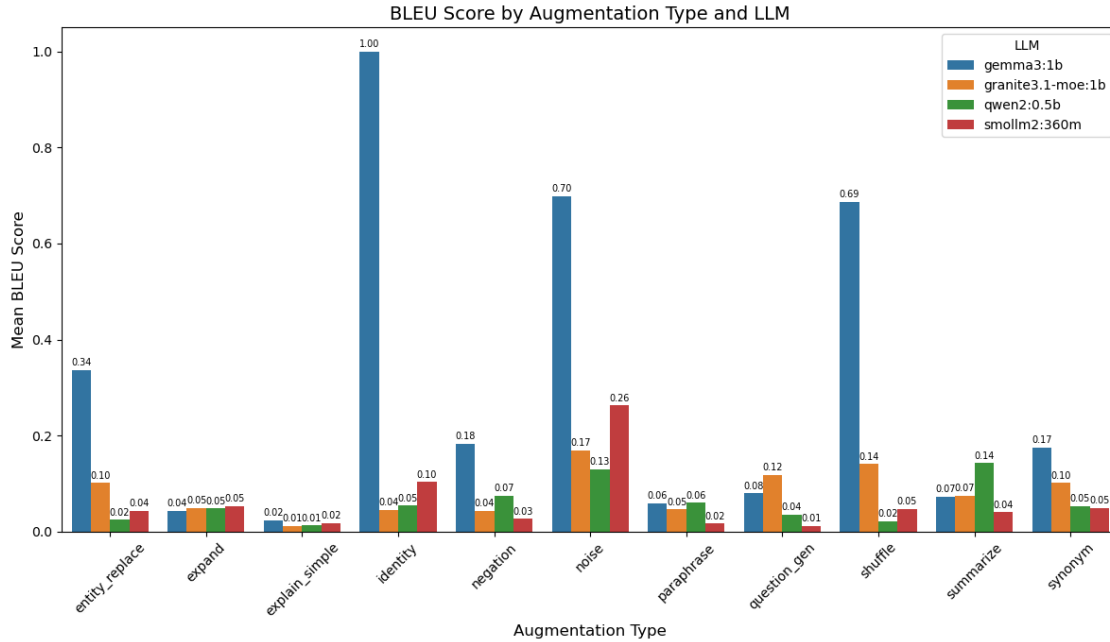
```

    data=agg,
    x='augmentation_type',
    y='bleu',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean BLEU Score", fontsize=12)
plt.title("BLEU Score by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean BLEU Score Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='bleu',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("BLEU Score", fontsize=12)
plt.title("Distribution of BLEU Score by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

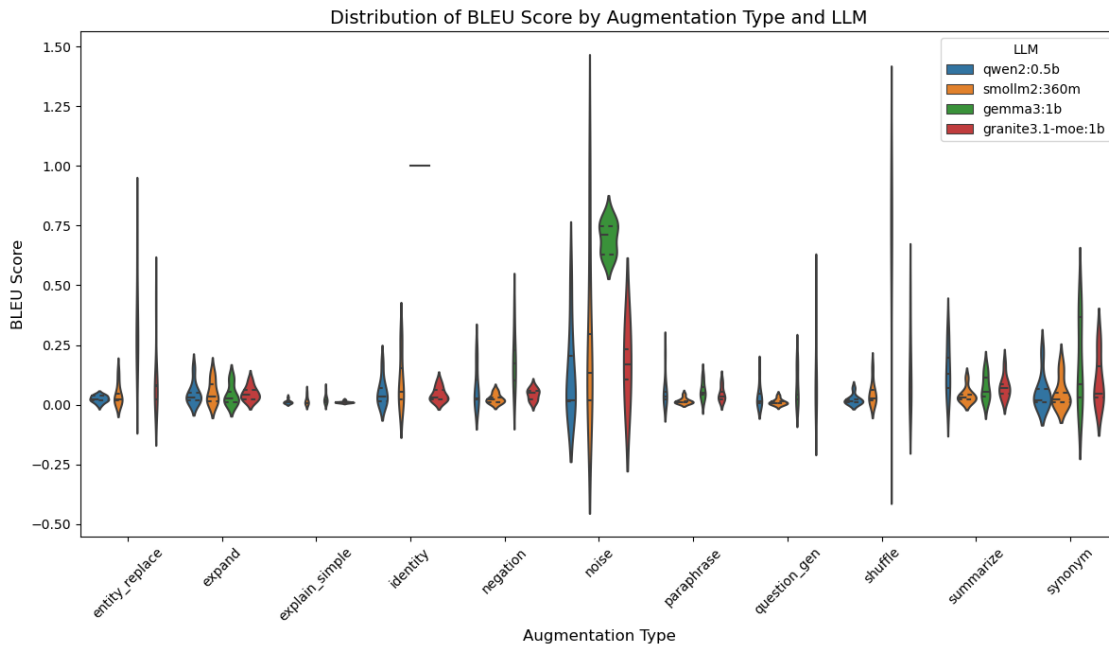
```



Mean BLEU Score Table:

augmentation_type	model	bleu
entity_replace	gemma3:1b	0.336007
entity_replace	granite3.1-moe:1b	0.100904
entity_replace	qwen2:0.5b	0.024408
entity_replace	smollm2:360m	0.042657
expand	gemma3:1b	0.042257
expand	granite3.1-moe:1b	0.048258
expand	qwen2:0.5b	0.047998
expand	smollm2:360m	0.053188
explain_simple	gemma3:1b	0.023838
explain_simple	granite3.1-moe:1b	0.011359
explain_simple	qwen2:0.5b	0.012483
explain_simple	smollm2:360m	0.017801
identity	gemma3:1b	1.000000
identity	granite3.1-moe:1b	0.044345
identity	qwen2:0.5b	0.053984
identity	smollm2:360m	0.102974
negation	gemma3:1b	0.182515
negation	granite3.1-moe:1b	0.043616
negation	qwen2:0.5b	0.074165
negation	smollm2:360m	0.026137
noise	gemma3:1b	0.698746
noise	granite3.1-moe:1b	0.168748
noise	qwen2:0.5b	0.130235

noise	smollm2:360m	0.262345
paraphrase	gemma3:1b	0.059311
paraphrase	granite3.1-moe:1b	0.046070
paraphrase	qwen2:0.5b	0.059510
paraphrase	smollm2:360m	0.016563
question_gen	gemma3:1b	0.079287
question_gen	granite3.1-moe:1b	0.118030
question_gen	qwen2:0.5b	0.035539
question_gen	smollm2:360m	0.012183
shuffle	gemma3:1b	0.687404
shuffle	granite3.1-moe:1b	0.140264
shuffle	qwen2:0.5b	0.021556
shuffle	smollm2:360m	0.047154
summarize	gemma3:1b	0.072126
summarize	granite3.1-moe:1b	0.074009
summarize	qwen2:0.5b	0.142341
summarize	smollm2:360m	0.039963
synonym	gemma3:1b	0.174645
synonym	granite3.1-moe:1b	0.101196
synonym	qwen2:0.5b	0.052701
synonym	smollm2:360m	0.049614



```
[13]: agg = df.groupby(['augmentation_type', 'model'])['cosine_similarity'].mean().
      ↪reset_index()

plt.figure(figsize=(12,7))
```

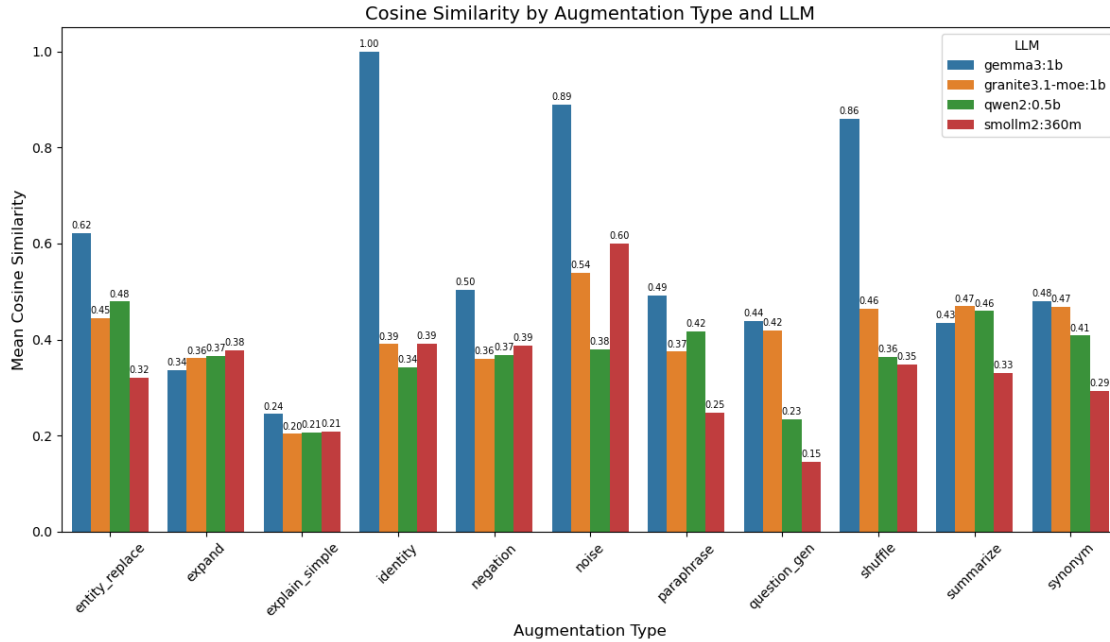
```

bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='cosine_similarity',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Cosine Similarity", fontsize=12)
plt.title("Cosine Similarity by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Cosine Similarity Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='cosine_similarity',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Cosine Similarity", fontsize=12)
plt.title("Distribution of Cosine Similarity by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

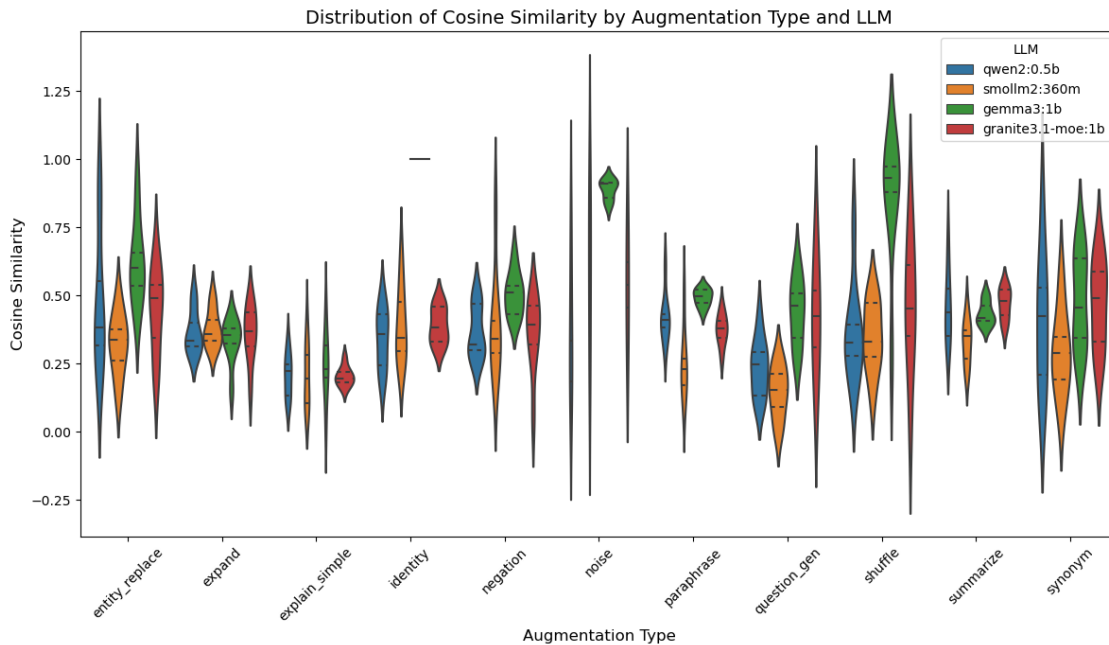
```



Mean Cosine Similarity Table:

augmentation_type	model	cosine_similarity
entity_replace	gemma3:1b	0.622095
entity_replace	granite3.1-moe:1b	0.445180
entity_replace	qwen2:0.5b	0.479066
entity_replace	smollm2:360m	0.321294
expand	gemma3:1b	0.335939
expand	granite3.1-moe:1b	0.361077
expand	qwen2:0.5b	0.365863
expand	smollm2:360m	0.377856
explain_simple	gemma3:1b	0.244547
explain_simple	granite3.1-moe:1b	0.203341
explain_simple	qwen2:0.5b	0.206513
explain_simple	smollm2:360m	0.208020
identity	gemma3:1b	1.000000
identity	granite3.1-moe:1b	0.390596
identity	qwen2:0.5b	0.342094
identity	smollm2:360m	0.391382
negation	gemma3:1b	0.504206
negation	granite3.1-moe:1b	0.360369
negation	qwen2:0.5b	0.367078
negation	smollm2:360m	0.386901
noise	gemma3:1b	0.889123
noise	granite3.1-moe:1b	0.539366
noise	qwen2:0.5b	0.379859

noise	smollm2:360m	0.600188
paraphrase	gemma3:1b	0.492292
paraphrase	granite3.1-moe:1b	0.374676
paraphrase	qwen2:0.5b	0.417718
paraphrase	smollm2:360m	0.247645
question_gen	gemma3:1b	0.439280
question_gen	granite3.1-moe:1b	0.419343
question_gen	qwen2:0.5b	0.234261
question_gen	smollm2:360m	0.145924
shuffle	gemma3:1b	0.859389
shuffle	granite3.1-moe:1b	0.464123
shuffle	qwen2:0.5b	0.364328
shuffle	smollm2:360m	0.348602
summarize	gemma3:1b	0.434552
summarize	granite3.1-moe:1b	0.469117
summarize	qwen2:0.5b	0.459336
summarize	smollm2:360m	0.330073
synonym	gemma3:1b	0.480473
synonym	granite3.1-moe:1b	0.467419
synonym	qwen2:0.5b	0.408312
synonym	smollm2:360m	0.293056



```
[14]: agg = df.groupby(['augmentation_type', 'model'])['wer'].mean().reset_index()

plt.figure(figsize=(12,7))
bar = sns.barplot(
```

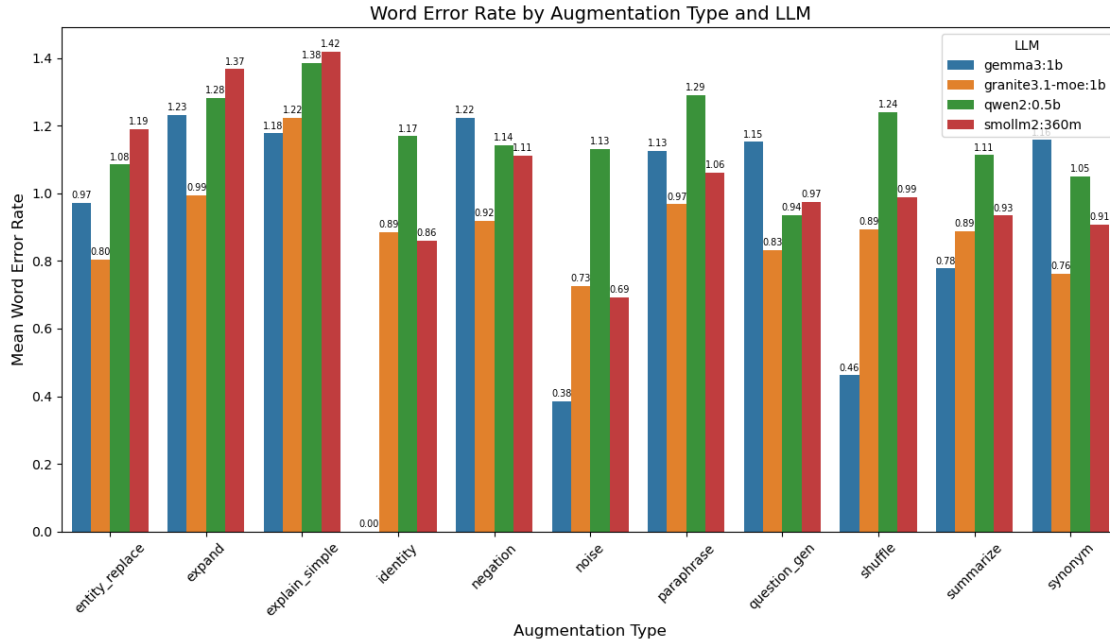
```

    data=agg,
    x='augmentation_type',
    y='wer',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Word Error Rate", fontsize=12)
plt.title("Word Error Rate by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Word Error Rate Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='wer',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Word Error Rate", fontsize=12)
plt.title("Distribution of Word Error Rate by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

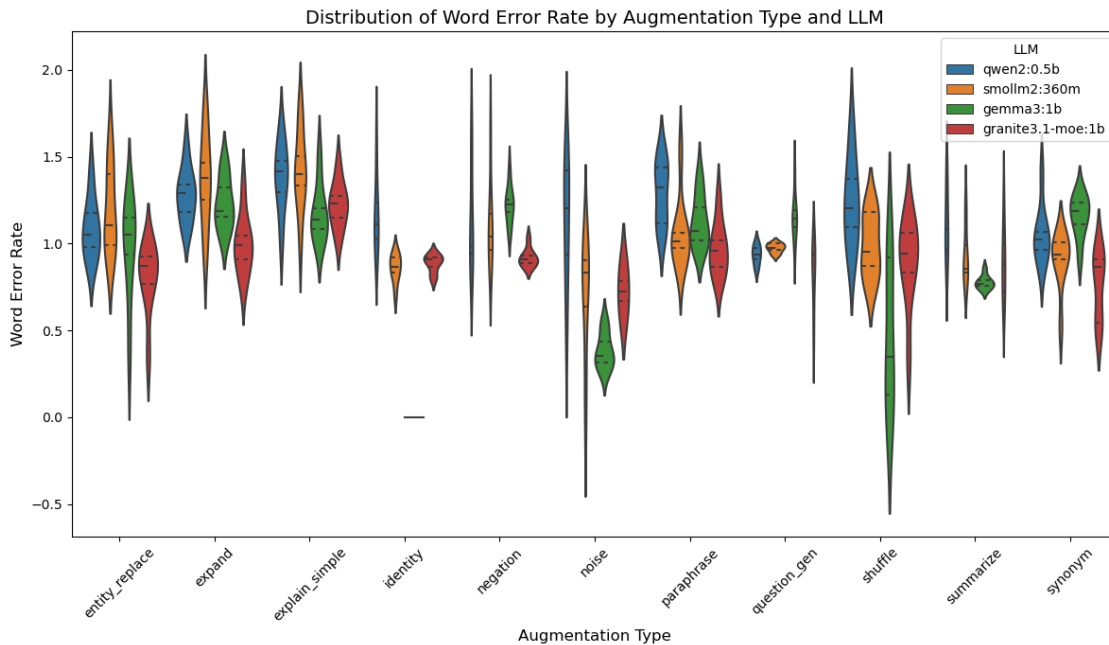
```



Mean Word Error Rate Table:

augmentation_type	model	wer
entity_replace	gemma3:1b	0.970939
entity_replace	granite3.1-moe:1b	0.804647
entity_replace	qwen2:0.5b	1.084771
entity_replace	smollm2:360m	1.190478
expand	gemma3:1b	1.230516
expand	granite3.1-moe:1b	0.993683
expand	qwen2:0.5b	1.280821
expand	smollm2:360m	1.367130
explain_simple	gemma3:1b	1.177038
explain_simple	granite3.1-moe:1b	1.223285
explain_simple	qwen2:0.5b	1.384482
explain_simple	smollm2:360m	1.419171
identity	gemma3:1b	0.000000
identity	granite3.1-moe:1b	0.885898
identity	qwen2:0.5b	1.168764
identity	smollm2:360m	0.859945
negation	gemma3:1b	1.222498
negation	granite3.1-moe:1b	0.919054
negation	qwen2:0.5b	1.141896
negation	smollm2:360m	1.111944
noise	gemma3:1b	0.384143
noise	granite3.1-moe:1b	0.725806
noise	qwen2:0.5b	1.131913

noise	smollm2:360m	0.692547
paraphrase	gemma3:1b	1.125937
paraphrase	granite3.1-moe:1b	0.967314
paraphrase	qwen2:0.5b	1.290791
paraphrase	smollm2:360m	1.060419
question_gen	gemma3:1b	1.151979
question_gen	granite3.1-moe:1b	0.833109
question_gen	qwen2:0.5b	0.935138
question_gen	smollm2:360m	0.974449
shuffle	gemma3:1b	0.462177
shuffle	granite3.1-moe:1b	0.893410
shuffle	qwen2:0.5b	1.240501
shuffle	smollm2:360m	0.989375
summarize	gemma3:1b	0.777333
summarize	granite3.1-moe:1b	0.887743
summarize	qwen2:0.5b	1.113069
summarize	smollm2:360m	0.933830
synonym	gemma3:1b	1.157594
synonym	granite3.1-moe:1b	0.763213
synonym	qwen2:0.5b	1.049095
synonym	smollm2:360m	0.906838



```
[15]: agg = df.groupby(['augmentation_type', 'model'])['char_diversity'].mean().
      ↪reset_index()

plt.figure(figsize=(12,7))
```

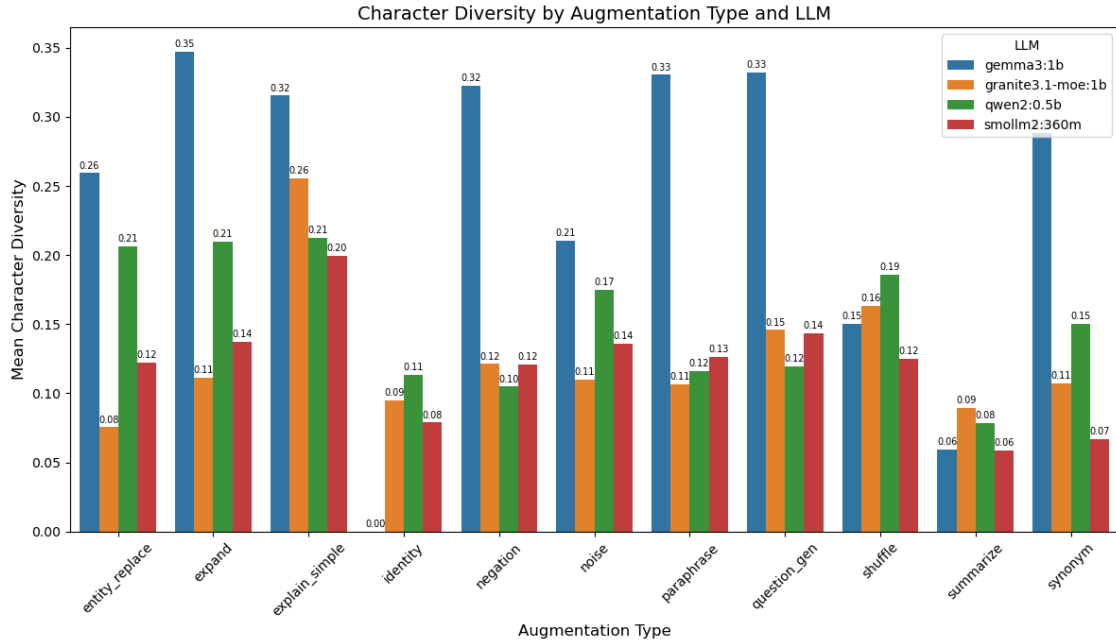
```

bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='char_diversity',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Character Diversity", fontsize=12)
plt.title("Character Diversity by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Character Diversity Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='char_diversity',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Character Diversity", fontsize=12)
plt.title("Distribution of Character Diversity by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

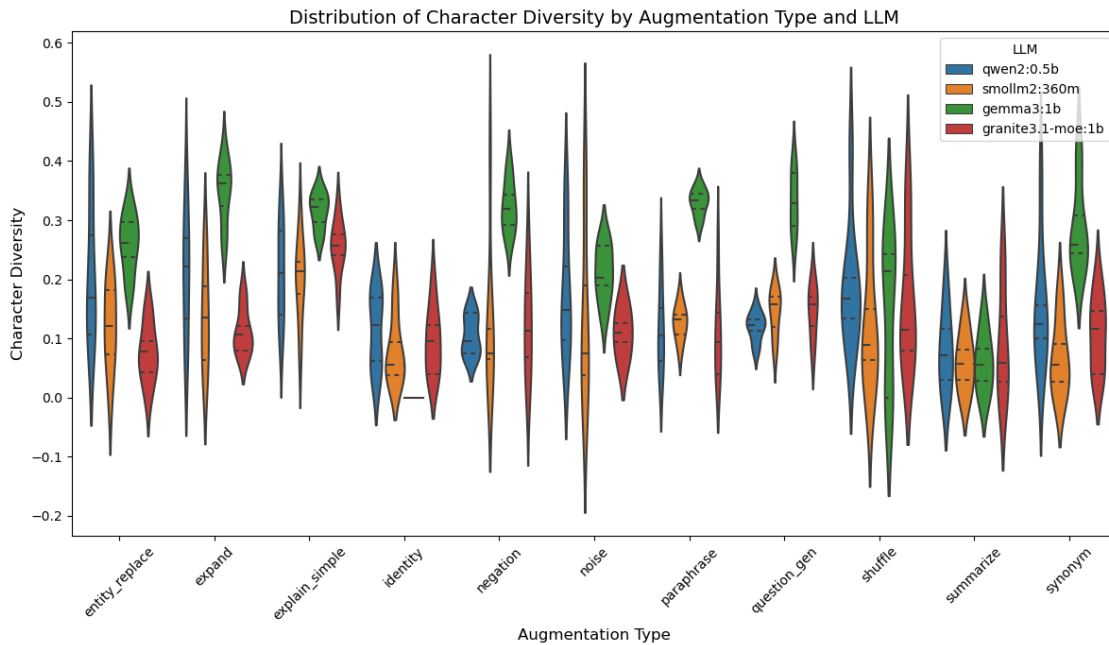
```



Mean Character Diversity Table:

augmentation_type	model	char_diversity
entity_replace	gemma3:1b	0.259408
entity_replace	granite3.1-moe:1b	0.075425
entity_replace	qwen2:0.5b	0.206214
entity_replace	smollm2:360m	0.122106
expand	gemma3:1b	0.347327
expand	granite3.1-moe:1b	0.111536
expand	qwen2:0.5b	0.209976
expand	smollm2:360m	0.137446
explain_simple	gemma3:1b	0.315454
explain_simple	granite3.1-moe:1b	0.255851
explain_simple	qwen2:0.5b	0.212264
explain_simple	smollm2:360m	0.199737
identity	gemma3:1b	0.000000
identity	granite3.1-moe:1b	0.094899
identity	qwen2:0.5b	0.113213
identity	smollm2:360m	0.078906
negation	gemma3:1b	0.322866
negation	granite3.1-moe:1b	0.121212
negation	qwen2:0.5b	0.104828
negation	smollm2:360m	0.120919
noise	gemma3:1b	0.210512
noise	granite3.1-moe:1b	0.109890
noise	qwen2:0.5b	0.174991

noise	smollm2:360m	0.135840
paraphrase	gemma3:1b	0.330481
paraphrase	granite3.1-moe:1b	0.106477
paraphrase	qwen2:0.5b	0.116298
paraphrase	smollm2:360m	0.126579
question_gen	gemma3:1b	0.332157
question_gen	granite3.1-moe:1b	0.145877
question_gen	qwen2:0.5b	0.119569
question_gen	smollm2:360m	0.143274
shuffle	gemma3:1b	0.150104
shuffle	granite3.1-moe:1b	0.163039
shuffle	qwen2:0.5b	0.186026
shuffle	smollm2:360m	0.124883
summarize	gemma3:1b	0.059246
summarize	granite3.1-moe:1b	0.089673
summarize	qwen2:0.5b	0.078301
summarize	smollm2:360m	0.058926
synonym	gemma3:1b	0.288397
synonym	granite3.1-moe:1b	0.107062
synonym	qwen2:0.5b	0.150192
synonym	smollm2:360m	0.067023



```
[16]: agg = df.groupby(['augmentation_type', 'model'])['type_token_ratio'].mean().
      ↪reset_index()

plt.figure(figsize=(12,7))
```


```

bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='type_token_ratio',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Type Token Ratio", fontsize=12)
plt.title("Type Token Ratio by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.ylim(.8,1.01)
plt.show()

print("Mean Type Token Ratio Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='type_token_ratio',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Type Token Ratio", fontsize=12)
plt.title("Distribution of Type Token Ratio by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

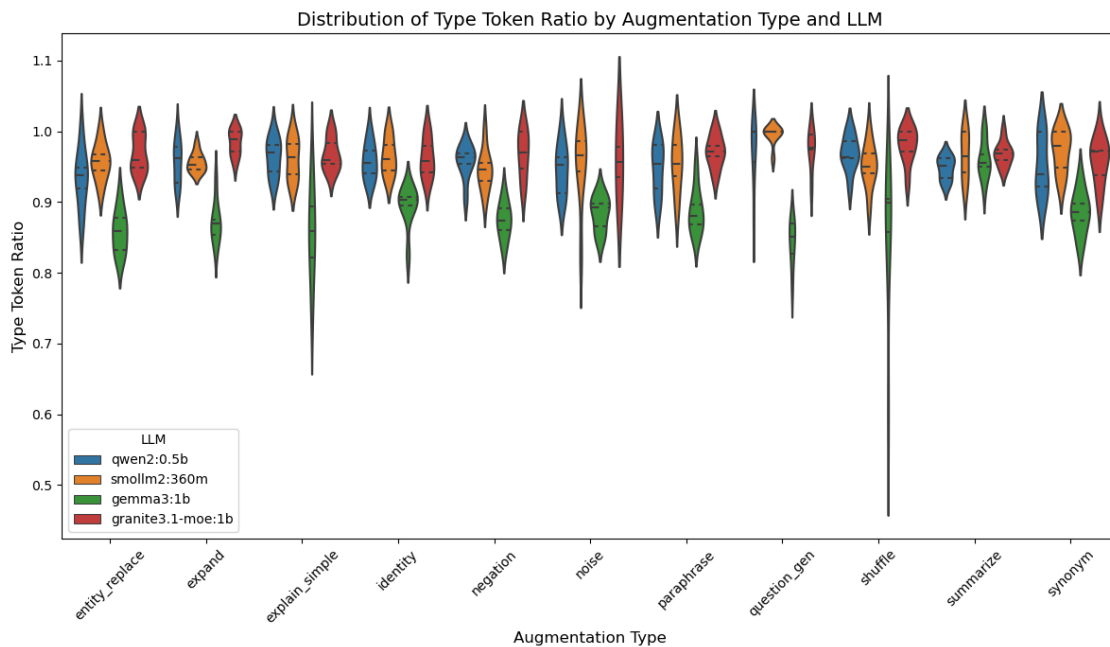
```



Mean Type Token Ratio Table:

augmentation_type	model	type_token_ratio
entity_replace	gemma3:1b	0.857892
entity_replace	granite3.1-moe:1b	0.969279
entity_replace	qwen2:0.5b	0.933791
entity_replace	smollm2:360m	0.956961
expand	gemma3:1b	0.871875
expand	granite3.1-moe:1b	0.983983
expand	qwen2:0.5b	0.956655
expand	smollm2:360m	0.956069
explain_simple	gemma3:1b	0.856441
explain_simple	granite3.1-moe:1b	0.967255
explain_simple	qwen2:0.5b	0.963499
explain_simple	smollm2:360m	0.960689
identity	gemma3:1b	0.894546
identity	granite3.1-moe:1b	0.962174
identity	qwen2:0.5b	0.957552
identity	smollm2:360m	0.964159
negation	gemma3:1b	0.874110
negation	granite3.1-moe:1b	0.967172
negation	qwen2:0.5b	0.956909
negation	smollm2:360m	0.945347
noise	gemma3:1b	0.884083
noise	granite3.1-moe:1b	0.957143
noise	qwen2:0.5b	0.945767
noise	smollm2:360m	0.953060

paraphrase	gemma3:1b	0.887194
paraphrase	granite3.1-moe:1b	0.971228
paraphrase	qwen2:0.5b	0.947000
paraphrase	smollm2:360m	0.953424
question_gen	gemma3:1b	0.843012
question_gen	granite3.1-moe:1b	0.974845
question_gen	qwen2:0.5b	0.972431
question_gen	smollm2:360m	0.995192
shuffle	gemma3:1b	0.854640
shuffle	granite3.1-moe:1b	0.980633
shuffle	qwen2:0.5b	0.969498
shuffle	smollm2:360m	0.951689
summarize	gemma3:1b	0.960620
summarize	granite3.1-moe:1b	0.967855
summarize	qwen2:0.5b	0.947847
summarize	smollm2:360m	0.965432
synonym	gemma3:1b	0.885274
synonym	granite3.1-moe:1b	0.957266
synonym	qwen2:0.5b	0.951701
synonym	smollm2:360m	0.971937



```
[17]: agg = df.groupby(['augmentation_type', 'model'])['bigram_overlap'].mean().
      ↪reset_index()

plt.figure(figsize=(12,7))
bar = sns.barplot(
```

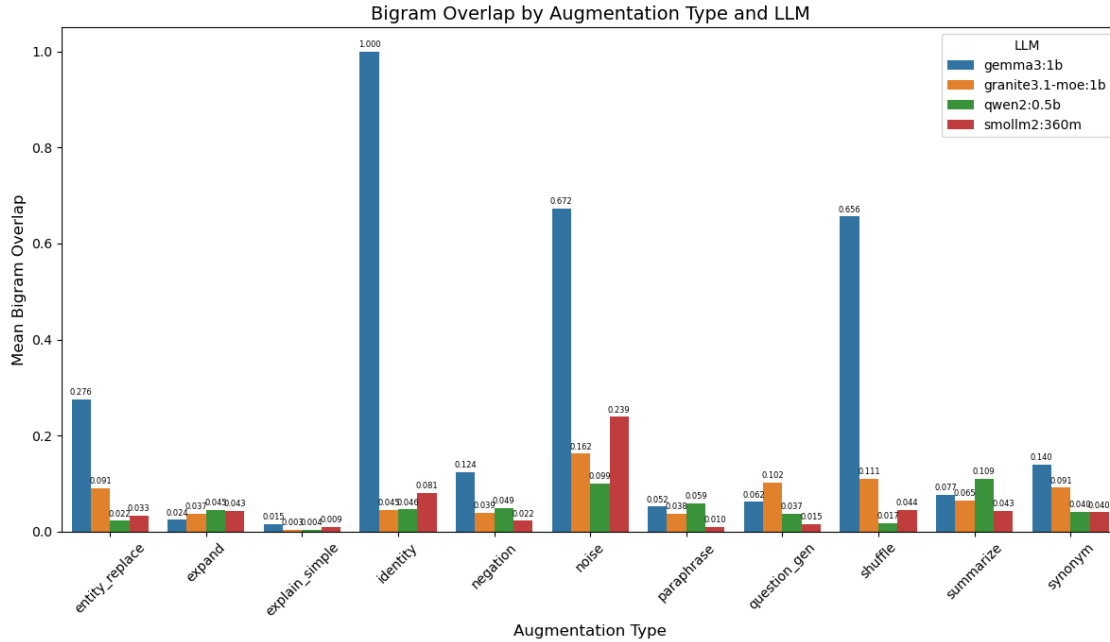
```

    data=agg,
    x='augmentation_type',
    y='bigram_overlap',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.3f', label_type='edge', padding=2,
    ↪ fontsize=6)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Bigram Overlap", fontsize=12)
plt.title("Bigram Overlap by Augmentation Type and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Bigram Overlap Table:\n")
print(agg.to_string(index=False))

plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='bigram_overlap',
    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Bigram Overlap", fontsize=12)
plt.title("Distribution of Bigram Overlap by Augmentation Type and LLM",
    ↪ fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

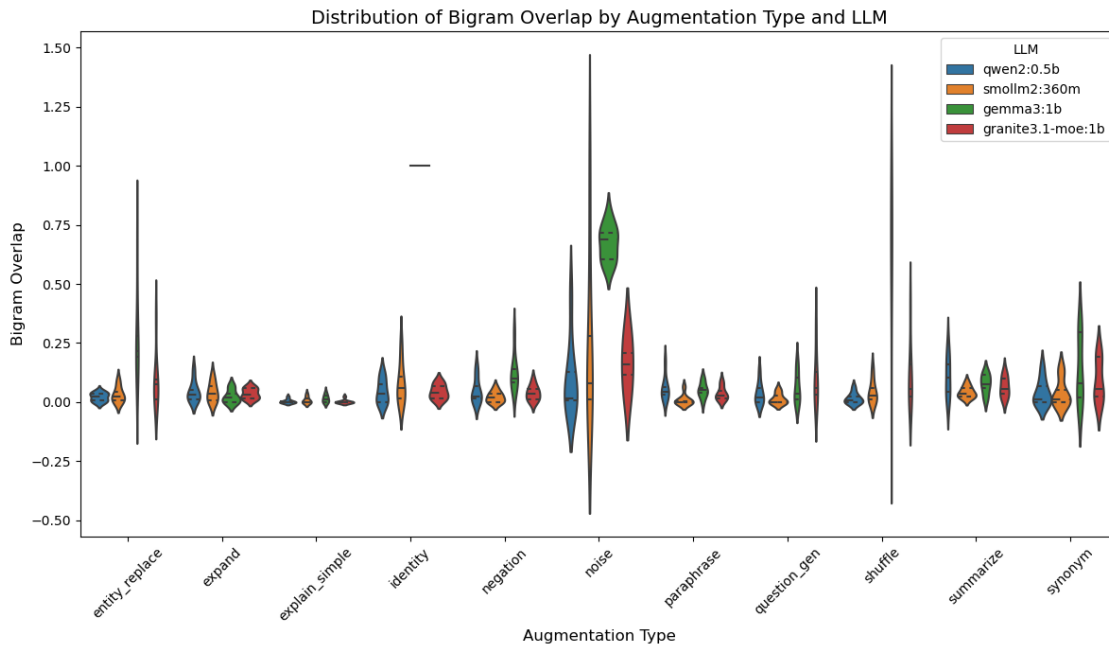
```

Mean Bigram Overlap Table:

augmentation_type	model	bigram_overlap
entity_replace	gemma3:1b	0.276008
entity_replace	granite3.1-moe:1b	0.090576
entity_replace	qwen2:0.5b	0.022345
entity_replace	smollm2:360m	0.032654
expand	gemma3:1b	0.024316
expand	granite3.1-moe:1b	0.036781
expand	qwen2:0.5b	0.045052
expand	smollm2:360m	0.043334
explain_simple	gemma3:1b	0.015429
explain_simple	granite3.1-moe:1b	0.003049
explain_simple	qwen2:0.5b	0.004320
explain_simple	smollm2:360m	0.008956
identity	gemma3:1b	1.000000
identity	granite3.1-moe:1b	0.044572
identity	qwen2:0.5b	0.046083
identity	smollm2:360m	0.081047
negation	gemma3:1b	0.124283
negation	granite3.1-moe:1b	0.039028
negation	qwen2:0.5b	0.049130
negation	smollm2:360m	0.022086
noise	gemma3:1b	0.672246
noise	granite3.1-moe:1b	0.161934
noise	qwen2:0.5b	0.099316

noise	smollm2:360m	0.238900
paraphrase	gemma3:1b	0.051615
paraphrase	granite3.1-moe:1b	0.037508
paraphrase	qwen2:0.5b	0.058800
paraphrase	smollm2:360m	0.010150
question_gen	gemma3:1b	0.061532
question_gen	granite3.1-moe:1b	0.102173
question_gen	qwen2:0.5b	0.036526
question_gen	smollm2:360m	0.014815
shuffle	gemma3:1b	0.655958
shuffle	granite3.1-moe:1b	0.110579
shuffle	qwen2:0.5b	0.016694
shuffle	smollm2:360m	0.044234
summarize	gemma3:1b	0.076663
summarize	granite3.1-moe:1b	0.064816
summarize	qwen2:0.5b	0.109437
summarize	smollm2:360m	0.042743
synonym	gemma3:1b	0.139534
synonym	granite3.1-moe:1b	0.091272
synonym	qwen2:0.5b	0.040477
synonym	smollm2:360m	0.040125




```
[18]: metrics = [
    'levenshtein_similarity', 'jaccard_similarity', 'length_ratio', 'bleu',
    'cosine_similarity', 'wer', 'char_diversity', 'type_token_ratio',
    ↪ 'bigram_overlap',
```

```

    'total_duration_s', 'load_duration_s', 'prompt_eval_duration_s',
    'tokens_per_second'
]
corr = df[metrics].corr()
plt.figure(figsize=(12,10))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Correlation Heatmap of All Metrics")
plt.show()

```



```

[19]: import numpy as np

# 1. Normalize tokens_per_second to 0-1 for fair combination
tps_min = df['tokens_per_second'].min()
tps_max = df['tokens_per_second'].max()

```

```

df['tokens_per_second_norm'] = (df['tokens_per_second'] - tps_min) / (tps_max -
↳tps_min)

# 2. Compute composite score
# You may change these weights (must sum to 1 for interpretation)
w1, w2, w3 = 0.4, 0.3, 0.3
df['composite_score'] = (
    w1 * df['levenshtein_similarity'] +
    w2 * (1 - df['bigram_overlap']) +
    w3 * df['tokens_per_second_norm']
)

# 3. Aggregate by augmentation_type and model
agg = df.groupby(['augmentation_type', 'model'])['composite_score'].mean().
↳reset_index()

# 4. Bar Plot
plt.figure(figsize=(12,7))
bar = sns.barplot(
    data=agg,
    x='augmentation_type',
    y='composite_score',
    hue='model',
    order=augmentation_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='% .2f', label_type='edge', padding=2,
↳fontsize=7)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Mean Composite Score", fontsize=12)
plt.title("Composite Score (Fidelity, Novelty, Speed) by Augmentation and LLM",
↳fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Composite Score Table:\n")
print(agg.to_string(index=False))


# 5. Violin Plot for distribution
plt.figure(figsize=(12,7))
sns.violinplot(
    data=df,
    x='augmentation_type',
    y='composite_score',

```

```

    hue='model',
    order=augmentation_order,
    palette='tab10',
    inner="quartile"
)
plt.xlabel("Augmentation Type", fontsize=12)
plt.ylabel("Composite Score", fontsize=12)
plt.title("Composite Score Distribution by Augmentation and LLM", fontsize=14)
plt.xticks(rotation=45)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()


```



Mean Composite Score Table:

augmentation_type	model	composite_score
entity_replace	gemma3:1b	0.335912
entity_replace	granite3.1-moe:1b	0.471249
entity_replace	qwen2:0.5b	0.616057
entity_replace	smollm2:360m	0.374485
expand	gemma3:1b	0.325728
expand	granite3.1-moe:1b	0.477295
expand	qwen2:0.5b	0.587588
expand	smollm2:360m	0.348680
explain_simple	gemma3:1b	0.324912
explain_simple	granite3.1-moe:1b	0.479531


explain_simple	qwen2:0.5b	0.583204
explain_simple	smollm2:360m	0.342336
identity	gemma3:1b	0.404691
identity	granite3.1-moe:1b	0.486590
identity	qwen2:0.5b	0.592319
identity	smollm2:360m	0.384884
negation	gemma3:1b	0.347554
negation	granite3.1-moe:1b	0.465591
negation	qwen2:0.5b	0.618670
negation	smollm2:360m	0.363032
noise	gemma3:1b	0.448422
noise	granite3.1-moe:1b	0.512618
noise	qwen2:0.5b	0.626415
noise	smollm2:360m	0.421667
paraphrase	gemma3:1b	0.322206
paraphrase	granite3.1-moe:1b	0.466552
paraphrase	qwen2:0.5b	0.575609
paraphrase	smollm2:360m	0.387589
question_gen	gemma3:1b	0.307361
question_gen	granite3.1-moe:1b	0.498469
question_gen	qwen2:0.5b	0.693501
question_gen	smollm2:360m	0.431563
shuffle	gemma3:1b	0.395759
shuffle	granite3.1-moe:1b	0.498410
shuffle	qwen2:0.5b	0.578838
shuffle	smollm2:360m	0.380806
summarize	gemma3:1b	0.488778
summarize	granite3.1-moe:1b	0.470224
summarize	qwen2:0.5b	0.568610
summarize	smollm2:360m	0.469857
synonym	gemma3:1b	0.325954
synonym	granite3.1-moe:1b	0.483146
synonym	qwen2:0.5b	0.614986
synonym	smollm2:360m	0.383781



```
[20]: metric_col = 'bleu'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean BLEU", fontsize=12)
plt.title("Mean BLEU by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean BLEU Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean BLEU Table by Prompt and Model:

prompt_id	model	bleu
1	gemma3:1b	0.254414
1	granite3.1-moe:1b	0.052784
1	qwen2:0.5b	0.044227
1	smollm2:360m	0.025847
2	gemma3:1b	0.374889
2	granite3.1-moe:1b	0.033505
2	qwen2:0.5b	0.049927
2	smollm2:360m	0.060754
3	gemma3:1b	0.367525
3	granite3.1-moe:1b	0.074351
3	qwen2:0.5b	0.076296
3	smollm2:360m	0.056945
4	gemma3:1b	0.300483
4	granite3.1-moe:1b	0.102675
4	qwen2:0.5b	0.067952
4	smollm2:360m	0.113136
5	gemma3:1b	0.287023
5	granite3.1-moe:1b	0.110838
5	qwen2:0.5b	0.106843
5	smollm2:360m	0.109319
6	gemma3:1b	0.231133
6	granite3.1-moe:1b	0.080830
6	qwen2:0.5b	0.056787


```

6      smollm2:360m 0.056255
7      gemma3:1b 0.317204
7 granite3.1-moe:1b 0.078031
7      qwen2:0.5b 0.050871
7      smollm2:360m 0.024435
8      gemma3:1b 0.308155
8 granite3.1-moe:1b 0.059570
8      qwen2:0.5b 0.023400
8      smollm2:360m 0.041003

```


```

[21]: metric_col = 'levenshtein_similarity'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Levenshtein Similarity", fontsize=12)
plt.title("Mean Levenshtein Similarity by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Levenshtein Similarity Table by Prompt and Model:\n")
print(agg.to_string(index=False))

```



Mean Levenshtein Similarity Table by Prompt and Model:


prompt_id	model	levenshtein_similarity
1	gemma3:1b	0.336949
1	granite3.1-moe:1b	0.022555
1	qwen2:0.5b	0.042446
1	smollm2:360m	0.088772
2	gemma3:1b	0.449745
2	granite3.1-moe:1b	0.058244
2	qwen2:0.5b	0.106182
2	smollm2:360m	0.104981
3	gemma3:1b	0.401415
3	granite3.1-moe:1b	0.090178
3	qwen2:0.5b	0.119815
3	smollm2:360m	0.206051
4	gemma3:1b	0.342165
4	granite3.1-moe:1b	0.110445
4	qwen2:0.5b	0.134457
4	smollm2:360m	0.208253
5	gemma3:1b	0.349078
5	granite3.1-moe:1b	0.133652
5	qwen2:0.5b	0.155324
5	smollm2:360m	0.213184
6	gemma3:1b	0.333702
6	granite3.1-moe:1b	0.101035
6	qwen2:0.5b	0.176270

6	smollm2:360m	0.165891
7	gemma3:1b	0.391250
7	granite3.1-moe:1b	0.129848
7	qwen2:0.5b	0.095614
7	smollm2:360m	0.188834
8	gemma3:1b	0.339023
8	granite3.1-moe:1b	0.130611
8	qwen2:0.5b	0.128982
8	smollm2:360m	0.274192

```
[22]: metric_col = 'jaccard_similarity'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Jaccard Similarity", fontsize=12)
plt.title("Mean Jaccard Similarity by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Jaccard Similarity Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Jaccard Similarity Table by Prompt and Model:


prompt_id	model	jaccard_similarity
1	gemma3:1b	0.307219
1	granite3.1-moe:1b	0.183757
1	qwen2:0.5b	0.127423
1	smollm2:360m	0.089247
2	gemma3:1b	0.412321
2	granite3.1-moe:1b	0.157192
2	qwen2:0.5b	0.153279
2	smollm2:360m	0.157645
3	gemma3:1b	0.424638
3	granite3.1-moe:1b	0.218134
3	qwen2:0.5b	0.154935
3	smollm2:360m	0.133321
4	gemma3:1b	0.329740
4	granite3.1-moe:1b	0.225126
4	qwen2:0.5b	0.133691
4	smollm2:360m	0.166607
5	gemma3:1b	0.341227
5	granite3.1-moe:1b	0.221673
5	qwen2:0.5b	0.184770
5	smollm2:360m	0.216383
6	gemma3:1b	0.295691
6	granite3.1-moe:1b	0.153952
6	qwen2:0.5b	0.145783

6	smollm2:360m	0.143711
7	gemma3:1b	0.359963
7	granite3.1-moe:1b	0.183441
7	qwen2:0.5b	0.096504
7	smollm2:360m	0.078461
8	gemma3:1b	0.366199
8	granite3.1-moe:1b	0.145337
8	qwen2:0.5b	0.097956
8	smollm2:360m	0.092543

```
[23]: metric_col = 'length_ratio'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Length Ratio", fontsize=12)
plt.title("Mean Length Ratio by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Length Ratio Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Length Ratio Table by Prompt and Model:


prompt_id	model	length_ratio
1	gemma3:1b	1.031050
1	granite3.1-moe:1b	0.902465
1	qwen2:0.5b	1.180218
1	smollm2:360m	1.015045
2	gemma3:1b	1.090909
2	granite3.1-moe:1b	1.121014
2	qwen2:0.5b	1.253623
2	smollm2:360m	1.094203
3	gemma3:1b	1.181061
3	granite3.1-moe:1b	1.120000
3	qwen2:0.5b	1.307576
3	smollm2:360m	1.118561
4	gemma3:1b	1.002674
4	granite3.1-moe:1b	0.832026
4	qwen2:0.5b	1.011586
4	smollm2:360m	0.903446
5	gemma3:1b	1.045455
5	granite3.1-moe:1b	0.995522
5	qwen2:0.5b	1.201832
5	smollm2:360m	0.983039
6	gemma3:1b	1.062328
6	granite3.1-moe:1b	1.011742
6	qwen2:0.5b	1.231061

6	smollm2:360m	1.068526
7	gemma3:1b	1.179189
7	granite3.1-moe:1b	1.107025
7	qwen2:0.5b	1.202855
7	smollm2:360m	1.084899
8	gemma3:1b	1.170641
8	granite3.1-moe:1b	1.118852
8	qwen2:0.5b	1.198584
8	smollm2:360m	0.996647

```
[24]: metric_col = 'cosine_similarity'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Cosine Similarity", fontsize=12)
plt.title("Mean Cosine Similarity by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Cosine Similarity Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Cosine Similarity Table by Prompt and Model:


prompt_id	model	cosine_similarity
1	gemma3:1b	0.518970
1	granite3.1-moe:1b	0.382670
1	qwen2:0.5b	0.328296
1	smollm2:360m	0.322072
2	gemma3:1b	0.642691
2	granite3.1-moe:1b	0.413292
2	qwen2:0.5b	0.355949
2	smollm2:360m	0.355266
3	gemma3:1b	0.633511
3	granite3.1-moe:1b	0.430231
3	qwen2:0.5b	0.369297
3	smollm2:360m	0.307824
4	gemma3:1b	0.582984
4	granite3.1-moe:1b	0.447602
4	qwen2:0.5b	0.353659
4	smollm2:360m	0.347627
5	gemma3:1b	0.564296
5	granite3.1-moe:1b	0.412355
5	qwen2:0.5b	0.363213
5	smollm2:360m	0.415689
6	gemma3:1b	0.521634
6	granite3.1-moe:1b	0.341044
6	qwen2:0.5b	0.405859

6	smollm2:360m	0.327513
7	gemma3:1b	0.555280
7	granite3.1-moe:1b	0.429507
7	qwen2:0.5b	0.477016
7	smollm2:360m	0.370274
8	gemma3:1b	0.563831
8	granite3.1-moe:1b	0.331496
8	qwen2:0.5b	0.273568
8	smollm2:360m	0.208965

```
[25]: metric_col = 'wer'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Word Error Rate", fontsize=12)
plt.title("Mean Word Error Rate by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Word Error Rate Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Word Error Rate Table by Prompt and Model:

prompt_id	model	wer
1	gemma3:1b	0.926290
1	granite3.1-moe:1b	0.900000
1	qwen2:0.5b	1.211302
1	smollm2:360m	1.095823
2	gemma3:1b	0.736842
2	granite3.1-moe:1b	0.976316
2	qwen2:0.5b	1.188995
2	smollm2:360m	1.086124
3	gemma3:1b	0.847594
3	granite3.1-moe:1b	0.967647
3	qwen2:0.5b	1.237968
3	smollm2:360m	1.122995
4	gemma3:1b	0.804545
4	granite3.1-moe:1b	0.792500
4	qwen2:0.5b	1.072727
4	smollm2:360m	0.961364
5	gemma3:1b	0.863636
5	granite3.1-moe:1b	0.776190
5	qwen2:0.5b	0.969697
5	smollm2:360m	0.861472
6	gemma3:1b	0.925837
6	granite3.1-moe:1b	0.886842
6	qwen2:0.5b	1.083732

```

6      smollm2:360m 0.961722
7      gemma3:1b 0.950147
7 granite3.1-moe:1b 0.912903
7      qwen2:0.5b 1.284457
7      smollm2:360m 1.143695
8      gemma3:1b 0.970674
8 granite3.1-moe:1b 1.103226
8      qwen2:0.5b 1.275660
8      smollm2:360m 1.134897

```


```

[26]: metric_col = 'char_diversity'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Character Diversity", fontsize=12)
plt.title("Mean Character Diversity by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Character Diversity Table by Prompt and Model:\n")
print(agg.to_string(index=False))

```



Mean Character Diversity Table by Prompt and Model:


prompt_id	model	char_diversity
1	gemma3:1b	0.270020
1	granite3.1-moe:1b	0.129324
1	qwen2:0.5b	0.132288
1	smollm2:360m	0.143191
2	gemma3:1b	0.209836
2	granite3.1-moe:1b	0.148893
2	qwen2:0.5b	0.130082
2	smollm2:360m	0.095010
3	gemma3:1b	0.241285
3	granite3.1-moe:1b	0.092397
3	qwen2:0.5b	0.120813
3	smollm2:360m	0.111742
4	gemma3:1b	0.233601
4	granite3.1-moe:1b	0.098295
4	qwen2:0.5b	0.128973
4	smollm2:360m	0.105373
5	gemma3:1b	0.254620
5	granite3.1-moe:1b	0.113102
5	qwen2:0.5b	0.132592
5	smollm2:360m	0.103992
6	gemma3:1b	0.212050
6	granite3.1-moe:1b	0.128421
6	qwen2:0.5b	0.149011

6	smollm2:360m	0.079962
7	gemma3:1b	0.234534
7	granite3.1-moe:1b	0.086330
7	qwen2:0.5b	0.238583
7	smollm2:360m	0.165616
8	gemma3:1b	0.246564
8	granite3.1-moe:1b	0.212882
8	qwen2:0.5b	0.183565
8	smollm2:360m	0.151941

```
[27]: metric_col = 'type_token_ratio'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Type Token Ratio", fontsize=12)
plt.title("Mean Type Token Ratio by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Type Token Ratio Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Type Token Ratio Table by Prompt and Model:


prompt_id	model	type_token_ratio
1	gemma3:1b	0.893512
1	granite3.1-moe:1b	0.984078
1	qwen2:0.5b	0.954470
1	smollm2:360m	0.964942
2	gemma3:1b	0.877512
2	granite3.1-moe:1b	0.958874
2	qwen2:0.5b	0.957332
2	smollm2:360m	0.943816
3	gemma3:1b	0.880468
3	granite3.1-moe:1b	0.969680
3	qwen2:0.5b	0.942743
3	smollm2:360m	0.963048
4	gemma3:1b	0.864202
4	granite3.1-moe:1b	0.962176
4	qwen2:0.5b	0.948129
4	smollm2:360m	0.948189
5	gemma3:1b	0.890984
5	granite3.1-moe:1b	0.962562
5	qwen2:0.5b	0.959796
5	smollm2:360m	0.970213
6	gemma3:1b	0.863806
6	granite3.1-moe:1b	0.972466
6	qwen2:0.5b	0.952460

6	smollm2:360m	0.958267
7	gemma3:1b	0.901537
7	granite3.1-moe:1b	0.959901
7	qwen2:0.5b	0.953417
7	smollm2:360m	0.957546
8	gemma3:1b	0.860478
8	granite3.1-moe:1b	0.988074
8	qwen2:0.5b	0.969944
8	smollm2:360m	0.984129

```
[28]: metric_col = 'bigram_overlap'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.3f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Bigram Overlap", fontsize=12)
plt.title("Mean Bigram Overlap by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Bigram Overlap Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Bigram Overlap Table by Prompt and Model:


prompt_id	model	bigram_overlap
1	gemma3:1b	0.227948
1	granite3.1-moe:1b	0.047773
1	qwen2:0.5b	0.033781
1	smollm2:360m	0.016246
2	gemma3:1b	0.351405
2	granite3.1-moe:1b	0.035467
2	qwen2:0.5b	0.045038
2	smollm2:360m	0.052751
3	gemma3:1b	0.333533
3	granite3.1-moe:1b	0.062834
3	qwen2:0.5b	0.052062
3	smollm2:360m	0.039926
4	gemma3:1b	0.277089
4	granite3.1-moe:1b	0.092721
4	qwen2:0.5b	0.048092
4	smollm2:360m	0.107319
5	gemma3:1b	0.254168
5	granite3.1-moe:1b	0.093665
5	qwen2:0.5b	0.094014
5	smollm2:360m	0.106278
6	gemma3:1b	0.218936
6	granite3.1-moe:1b	0.063029
6	qwen2:0.5b	0.057376

6	smollm2:360m	0.046410
7	gemma3:1b	0.296233
7	granite3.1-moe:1b	0.063329
7	qwen2:0.5b	0.040025
7	smollm2:360m	0.018833
8	gemma3:1b	0.293477
8	granite3.1-moe:1b	0.049417
8	qwen2:0.5b	0.013744
8	smollm2:360m	0.033360

```
[29]: metric_col = 'total_duration_s'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Total Duration (s)", fontsize=12)
plt.title("Mean Total Duration (s) by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Total Duration (s) Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Total Duration (s) Table by Prompt and Model:


prompt_id	model	total_duration_s
1	gemma3:1b	19.188144
1	granite3.1-moe:1b	14.308235
1	qwen2:0.5b	7.998199
1	smollm2:360m	13.717496
2	gemma3:1b	16.380431
2	granite3.1-moe:1b	9.874609
2	qwen2:0.5b	7.288003
2	smollm2:360m	11.835116
3	gemma3:1b	16.000983
3	granite3.1-moe:1b	10.246667
3	qwen2:0.5b	7.114733
3	smollm2:360m	10.895771
4	gemma3:1b	16.743796
4	granite3.1-moe:1b	10.095100
4	qwen2:0.5b	7.329973
4	smollm2:360m	11.602065
5	gemma3:1b	16.463278
5	granite3.1-moe:1b	10.051996
5	qwen2:0.5b	7.319335
5	smollm2:360m	11.066141
6	gemma3:1b	15.499161
6	granite3.1-moe:1b	9.947697
6	qwen2:0.5b	7.036435

6	smollm2:360m	10.453246
7	gemma3:1b	15.390212
7	granite3.1-moe:1b	10.223190
7	qwen2:0.5b	6.846961
7	smollm2:360m	9.418833
8	gemma3:1b	15.370792
8	granite3.1-moe:1b	9.339246
8	qwen2:0.5b	6.350768
8	smollm2:360m	9.481517

```
[30]: df['load_duration_s'] = df['load_duration_ns'] / 1e9
metric_col = 'load_duration_s'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Load Duration (s)", fontsize=12)
plt.title("Mean Load Duration (s) by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Load Duration (s) Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Load Duration (s) Table by Prompt and Model:


prompt_id	model	load_duration_s
1	gemma3:1b	1.920152
1	granite3.1-moe:1b	3.322415
1	qwen2:0.5b	0.267021
1	smollm2:360m	1.627052
2	gemma3:1b	0.134766
2	granite3.1-moe:1b	0.021508
2	qwen2:0.5b	0.051730
2	smollm2:360m	0.021827
3	gemma3:1b	0.134548
3	granite3.1-moe:1b	0.021360
3	qwen2:0.5b	0.051673
3	smollm2:360m	0.021094
4	gemma3:1b	0.134743
4	granite3.1-moe:1b	0.020843
4	qwen2:0.5b	0.051226
4	smollm2:360m	0.021520
5	gemma3:1b	0.134060
5	granite3.1-moe:1b	0.021788
5	qwen2:0.5b	0.052652
5	smollm2:360m	0.021343
6	gemma3:1b	0.133452
6	granite3.1-moe:1b	0.021591
6	qwen2:0.5b	0.051767

6	smollm2:360m	0.021137
7	gemma3:1b	0.135561
7	granite3.1-moe:1b	0.021237
7	qwen2:0.5b	0.052055
7	smollm2:360m	0.020865
8	gemma3:1b	0.134019
8	granite3.1-moe:1b	0.020916
8	qwen2:0.5b	0.051736
8	smollm2:360m	0.021303

```
[31]: metric_col = 'prompt_eval_count'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.1f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Prompt Eval Count", fontsize=12)
plt.title("Mean Prompt Eval Count by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Prompt Eval Count Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Prompt Eval Count Table by Prompt and Model:


prompt_id	model	prompt_eval_count
1	gemma3:1b	63.636364
1	granite3.1-moe:1b	121.100000
1	qwen2:0.5b	64.818182
1	smollm2:360m	84.181818
2	gemma3:1b	61.636364
2	granite3.1-moe:1b	102.100000
2	qwen2:0.5b	60.818182
2	smollm2:360m	83.181818
3	gemma3:1b	56.636364
3	granite3.1-moe:1b	111.100000
3	qwen2:0.5b	56.818182
3	smollm2:360m	79.181818
4	gemma3:1b	65.636364
4	granite3.1-moe:1b	120.100000
4	qwen2:0.5b	64.818182
4	smollm2:360m	87.181818
5	gemma3:1b	63.636364
5	granite3.1-moe:1b	115.100000
5	qwen2:0.5b	63.818182
5	smollm2:360m	85.181818
6	gemma3:1b	60.636364
6	granite3.1-moe:1b	107.500000
6	qwen2:0.5b	60.818182

6	smollm2:360m	82.181818
7	gemma3:1b	56.636364
7	granite3.1-moe:1b	109.500000
7	qwen2:0.5b	55.818182
7	smollm2:360m	75.181818
8	gemma3:1b	55.636364
8	granite3.1-moe:1b	104.500000
8	qwen2:0.5b	53.818182
8	smollm2:360m	77.181818

```
[32]: df['prompt_eval_duration_s'] = df['prompt_eval_duration_ns'] / 1e9
metric_col = 'prompt_eval_duration_s'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Prompt Eval Duration (s)", fontsize=12)
plt.title("Mean Prompt Eval Duration (s) by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Prompt Eval Duration (s) Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Prompt Eval Duration (s) Table by Prompt and Model:


prompt_id	model	prompt_eval_duration_s
1	gemma3:1b	4.032869
1	granite3.1-moe:1b	2.649561
1	qwen2:0.5b	1.831723
1	smollm2:360m	1.786276
2	gemma3:1b	3.236039
2	granite3.1-moe:1b	1.662311
2	qwen2:0.5b	1.418149
2	smollm2:360m	1.450954
3	gemma3:1b	2.910115
3	granite3.1-moe:1b	1.926919
3	qwen2:0.5b	1.301053
3	smollm2:360m	1.338266
4	gemma3:1b	3.505857
4	granite3.1-moe:1b	2.198182
4	qwen2:0.5b	1.537797
4	smollm2:360m	1.537981
5	gemma3:1b	3.373812
5	granite3.1-moe:1b	2.052317
5	qwen2:0.5b	1.512005
5	smollm2:360m	1.497412
6	gemma3:1b	3.169544
6	granite3.1-moe:1b	1.810106
6	qwen2:0.5b	1.418276

6	smollm2:360m	1.402899
7	gemma3:1b	2.909144
7	granite3.1-moe:1b	1.873770
7	qwen2:0.5b	1.273654
7	smollm2:360m	1.237074
8	gemma3:1b	2.845163
8	granite3.1-moe:1b	1.721691
8	qwen2:0.5b	1.215047
8	smollm2:360m	1.317496

```
[33]: metric_col = 'eval_count'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.1f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Eval Count", fontsize=12)
plt.title("Mean Eval Count by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Eval Count Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Eval Count Table by Prompt and Model:


prompt_id	model	eval_count
1	gemma3:1b	56.272727
1	granite3.1-moe:1b	60.000000
1	qwen2:0.5b	54.636364
1	smollm2:360m	49.545455
2	gemma3:1b	55.545455
2	granite3.1-moe:1b	59.300000
2	qwen2:0.5b	54.090909
2	smollm2:360m	49.818182
3	gemma3:1b	55.272727
3	granite3.1-moe:1b	60.000000
3	qwen2:0.5b	53.454545
3	smollm2:360m	45.909091
4	gemma3:1b	55.909091
4	granite3.1-moe:1b	56.700000
4	qwen2:0.5b	53.363636
4	smollm2:360m	48.181818
5	gemma3:1b	55.363636
5	granite3.1-moe:1b	57.600000
5	qwen2:0.5b	53.454545
5	smollm2:360m	45.909091
6	gemma3:1b	52.000000
6	granite3.1-moe:1b	58.700000
6	qwen2:0.5b	51.818182

6	smollm2:360m	43.545455
7	gemma3:1b	52.727273
7	granite3.1-moe:1b	60.000000
7	qwen2:0.5b	51.363636
7	smollm2:360m	39.454545
8	gemma3:1b	52.909091
8	granite3.1-moe:1b	55.100000
8	qwen2:0.5b	47.363636
8	smollm2:360m	39.363636

```
[34]: df['eval_duration_s'] = df['eval_duration_ns'] / 1e9
metric_col = 'eval_duration_s'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Eval Duration (s)", fontsize=12)
plt.title("Mean Eval Duration (s) by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Eval Duration (s) Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Eval Duration (s) Table by Prompt and Model:


prompt_id	model	eval_duration_s
1	gemma3:1b	13.233804
1	granite3.1-moe:1b	8.333593
1	qwen2:0.5b	5.898187
1	smollm2:360m	10.302506
2	gemma3:1b	13.008460
2	granite3.1-moe:1b	8.188722
2	qwen2:0.5b	5.816410
2	smollm2:360m	10.360541
3	gemma3:1b	12.955185
3	granite3.1-moe:1b	8.296848
3	qwen2:0.5b	5.760703
3	smollm2:360m	9.534744
4	gemma3:1b	13.102083
4	granite3.1-moe:1b	7.874251
4	qwen2:0.5b	5.739488
4	smollm2:360m	10.040729
5	gemma3:1b	12.954279
5	granite3.1-moe:1b	7.976062
5	qwen2:0.5b	5.753469
5	smollm2:360m	9.545345
6	gemma3:1b	12.195011
6	granite3.1-moe:1b	8.114146
6	qwen2:0.5b	5.564658

6	smollm2:360m	9.027609
7	gemma3:1b	12.344349
7	granite3.1-moe:1b	8.326613
7	qwen2:0.5b	5.520027
7	smollm2:360m	8.158739
8	gemma3:1b	12.390459
8	granite3.1-moe:1b	7.594581
8	qwen2:0.5b	5.082774
8	smollm2:360m	8.140808

```
[35]: metric_col = 'tokens_per_second'
agg = df.groupby(['prompt_id', 'model'])[metric_col].mean().reset_index()
prompt_order = sorted(df['prompt_id'].unique())

plt.figure(figsize=(10,6))
bar = sns.barplot(
    data=agg,
    x='prompt_id',
    y=metric_col,
    hue='model',
    order=prompt_order,
    palette='tab10'
)
for container in bar.containers:
    bar.bar_label(container, fmt='%.2f', label_type='edge', padding=2,
    ↪ fontsize=7)
plt.xlabel("Prompt ID", fontsize=12)
plt.ylabel("Mean Tokens per Second", fontsize=12)
plt.title("Mean Tokens per Second by Prompt and Model", fontsize=14)
plt.legend(title="LLM")
plt.tight_layout()
plt.show()

print("Mean Tokens per Second Table by Prompt and Model:\n")
print(agg.to_string(index=False))
```



Mean Tokens per Second Table by Prompt and Model:

prompt_id	model	tokens_per_second
1	gemma3:1b	4.254656
1	granite3.1-moe:1b	7.199806
1	qwen2:0.5b	9.286185
1	smollm2:360m	4.818912
2	gemma3:1b	4.274260
2	granite3.1-moe:1b	7.241827
2	qwen2:0.5b	9.324803
2	smollm2:360m	4.825393
3	gemma3:1b	4.272150
3	granite3.1-moe:1b	7.231682
3	qwen2:0.5b	9.291864
3	smollm2:360m	4.827106
4	gemma3:1b	4.271214
4	granite3.1-moe:1b	7.207941
4	qwen2:0.5b	9.314725
4	smollm2:360m	4.813005
5	gemma3:1b	4.280101
5	granite3.1-moe:1b	7.223503
5	qwen2:0.5b	9.322892
5	smollm2:360m	4.829543
6	gemma3:1b	4.274892
6	granite3.1-moe:1b	7.235559
6	qwen2:0.5b	9.331980

6	smollm2:360m	4.842847
7	gemma3:1b	4.283511
7	granite3.1-moe:1b	7.206529
7	qwen2:0.5b	9.347458
7	smollm2:360m	4.854049
8	gemma3:1b	4.276154
8	granite3.1-moe:1b	7.261920
8	qwen2:0.5b	9.339646
8	smollm2:360m	4.853902

```
[36]: from scipy.stats import f_oneway
import pandas as pd

metric_col = 'bleu'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                              'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                              float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (BLEU):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean BLEU across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
                  difference in mean BLEU across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
                  for ANOVA.")
```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	12.031344	3.074649e-05	True
1	expand	0.088237	9.659122e-01	False
2	explain_simple	1.439620	2.522929e-01	False

3	identity	418.946728	2.337463e-23	True
4	negation	7.310758	9.090059e-04	True
5	noise	9.195276	3.905748e-04	True
6	paraphrase	1.888940	1.543782e-01	False
7	question_gen	2.234300	1.080536e-01	False
8	shuffle	24.313615	5.957550e-08	True
9	summarize	3.417430	3.085063e-02	True
10	synonym	2.062183	1.279123e-01	False

Significance Interpretation (BLEU):

Augmentation 'entity_replace': Significant difference in mean BLEU across LLMs (p=3.075e-05).

Augmentation 'expand': No significant difference in mean BLEU across LLMs (p=0.9659).

Augmentation 'explain_simple': No significant difference in mean BLEU across LLMs (p=0.2523).

Augmentation 'identity': Significant difference in mean BLEU across LLMs (p=2.337e-23).

Augmentation 'negation': Significant difference in mean BLEU across LLMs (p=0.000909).

Augmentation 'noise': Significant difference in mean BLEU across LLMs (p=0.0003906).

Augmentation 'paraphrase': No significant difference in mean BLEU across LLMs (p=0.1544).

Augmentation 'question_gen': No significant difference in mean BLEU across LLMs (p=0.1081).

Augmentation 'shuffle': Significant difference in mean BLEU across LLMs (p=5.958e-08).

Augmentation 'summarize': Significant difference in mean BLEU across LLMs (p=0.03085).

Augmentation 'synonym': No significant difference in mean BLEU across LLMs (p=0.1279).

```
[37]: metric_col = 'levenshtein_similarity'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):
          subset = df[df['augmentation_type'] == aug]
          groups = [subset[subset['model'] == m][metric_col].dropna().values
                     for m in subset['model'].unique()]
          if sum([len(g) > 1 for g in groups]) >= 2:
              stat, pval = f_oneway(*groups)
              anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                                     'p-value': pval})
          else:
              anova_results.append({'augmentation_type': aug, 'F-statistic':
                                     float('nan'), 'p-value': float('nan')})
```



```

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Levenshtein Similarity):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant_
↪difference in mean Levenshtein similarity across LLMs (p={row['p-value']:.
↪4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant_
↪difference in mean Levenshtein similarity across LLMs (p={row['p-value']:.
↪4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data_
↪for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	2.625844	6.997952e-02	False
1	expand	0.024014	9.948452e-01	False
2	explain_simple	1.190140	3.314228e-01	False
3	identity	132.481680	1.200935e-16	True
4	negation	2.854926	5.501458e-02	False
5	noise	14.193399	2.299913e-05	True
6	paraphrase	1.649674	2.004514e-01	False
7	question_gen	8.240009	5.109847e-04	True
8	shuffle	20.934390	2.570053e-07	True
9	summarize	31.819413	3.665627e-09	True
10	synonym	0.103508	9.573280e-01	False

Significance Interpretation (Levenshtein Similarity):

Augmentation 'entity_replace': No significant difference in mean Levenshtein similarity across LLMs (p=0.06998).

Augmentation 'expand': No significant difference in mean Levenshtein similarity across LLMs (p=0.9948).

Augmentation 'explain_simple': No significant difference in mean Levenshtein similarity across LLMs (p=0.3314).

Augmentation 'identity': Significant difference in mean Levenshtein similarity across LLMs (p=1.201e-16).

Augmentation 'negation': No significant difference in mean Levenshtein similarity across LLMs (p=0.05501).

Augmentation 'noise': Significant difference in mean Levenshtein similarity across LLMs (p=2.3e-05).

Augmentation 'paraphrase': No significant difference in mean Levenshtein similarity across LLMs (p=0.2005).

Augmentation 'question_gen': Significant difference in mean Levenshtein similarity across LLMs (p=0.000511).
 Augmentation 'shuffle': Significant difference in mean Levenshtein similarity across LLMs (p=2.57e-07).
 Augmentation 'summarize': Significant difference in mean Levenshtein similarity across LLMs (p=3.666e-09).
 Augmentation 'synonym': No significant difference in mean Levenshtein similarity across LLMs (p=0.9573).

```
[38]: metric_col = 'jaccard_similarity'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                              'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                              float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Jaccard Similarity):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean Jaccard similarity across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
                  difference in mean Jaccard similarity across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
                  for ANOVA.")
```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	7.304946	9.132850e-04	True
1	expand	2.343734	9.446474e-02	False
2	explain_simple	0.829942	4.886380e-01	False
3	identity	316.158954	1.086267e-21	True
4	negation	3.746211	2.218612e-02	True
5	noise	7.561195	1.180207e-03	True
6	paraphrase	6.759171	1.429236e-03	True

7	question_gen	5.763377	3.678599e-03	True
8	shuffle	17.924767	1.088418e-06	True
9	summarize	3.590329	2.591959e-02	True
10	synonym	2.171286	1.136909e-01	False

Significance Interpretation (Jaccard Similarity):

Augmentation 'entity_replace': Significant difference in mean Jaccard similarity across LLMs (p=0.0009133).

Augmentation 'expand': No significant difference in mean Jaccard similarity across LLMs (p=0.09446).

Augmentation 'explain_simple': No significant difference in mean Jaccard similarity across LLMs (p=0.4886).

Augmentation 'identity': Significant difference in mean Jaccard similarity across LLMs (p=1.086e-21).

Augmentation 'negation': Significant difference in mean Jaccard similarity across LLMs (p=0.02219).

Augmentation 'noise': Significant difference in mean Jaccard similarity across LLMs (p=0.00118).

Augmentation 'paraphrase': Significant difference in mean Jaccard similarity across LLMs (p=0.001429).

Augmentation 'question_gen': Significant difference in mean Jaccard similarity across LLMs (p=0.003679).

Augmentation 'shuffle': Significant difference in mean Jaccard similarity across LLMs (p=1.088e-06).

Augmentation 'summarize': Significant difference in mean Jaccard similarity across LLMs (p=0.02592).

Augmentation 'synonym': No significant difference in mean Jaccard similarity across LLMs (p=0.1137).

```
[39]: metric_col = 'length_ratio'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                               'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                               float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)
```

```

print("\nSignificance Interpretation (Length Ratio):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
↳difference in mean Length Ratio across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
↳difference in mean Length Ratio across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
↳for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	2.566639	7.450275e-02	False
1	expand	19.571886	4.850757e-07	True
2	explain_simple	3.906114	1.894207e-02	True
3	identity	5.717529	3.494327e-03	True
4	negation	1.480470	2.412523e-01	False
5	noise	4.947161	8.949647e-03	True
6	paraphrase	11.201796	5.277643e-05	True
7	question_gen	30.531595	1.125467e-08	True
8	shuffle	1.587574	2.145492e-01	False
9	summarize	20.727578	2.825033e-07	True
10	synonym	1.586254	2.148594e-01	False

Significance Interpretation (Length Ratio):

Augmentation 'entity_replace': No significant difference in mean Length Ratio across LLMs (p=0.0745).

Augmentation 'expand': Significant difference in mean Length Ratio across LLMs (p=4.851e-07).

Augmentation 'explain_simple': Significant difference in mean Length Ratio across LLMs (p=0.01894).

Augmentation 'identity': Significant difference in mean Length Ratio across LLMs (p=0.003494).

Augmentation 'negation': No significant difference in mean Length Ratio across LLMs (p=0.2413).

Augmentation 'noise': Significant difference in mean Length Ratio across LLMs (p=0.00895).

Augmentation 'paraphrase': Significant difference in mean Length Ratio across LLMs (p=5.278e-05).

Augmentation 'question_gen': Significant difference in mean Length Ratio across LLMs (p=1.125e-08).

Augmentation 'shuffle': No significant difference in mean Length Ratio across LLMs (p=0.2145).

Augmentation 'summarize': Significant difference in mean Length Ratio across LLMs (p=2.825e-07).

Augmentation 'synonym': No significant difference in mean Length Ratio across LLMs (p=0.2149).

```
[40]: metric_col = 'cosine_similarity'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                               'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                               float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Cosine Similarity):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean Cosine Similarity across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
                  difference in mean Cosine Similarity across LLMs (p={row['p-value']:.4g}).")
    else:
        print(f"Augmentation '{row['augmentation_type']}': Insufficient data
              for ANOVA.")
```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	3.811782	2.078983e-02	True
1	expand	0.350688	7.889392e-01	False
2	explain_simple	0.315831	8.137917e-01	False
3	identity	88.400634	2.170128e-14	True
4	negation	1.984077	1.392127e-01	False
5	noise	6.981303	1.795858e-03	True
6	paraphrase	11.692212	3.825179e-05	True
7	question_gen	8.245998	5.087193e-04	True
8	shuffle	10.959296	6.205005e-05	True
9	summarize	4.278674	1.318119e-02	True
10	synonym	1.635780	2.035221e-01	False

Significance Interpretation (Cosine Similarity):

Augmentation 'entity_replace': Significant difference in mean Cosine Similarity across LLMs (p=0.02079).

Augmentation 'expand': No significant difference in mean Cosine Similarity across LLMs (p=0.7889).

Augmentation 'explain_simple': No significant difference in mean Cosine Similarity across LLMs (p=0.8138).

Augmentation 'identity': Significant difference in mean Cosine Similarity across LLMs (p=2.17e-14).

Augmentation 'negation': No significant difference in mean Cosine Similarity across LLMs (p=0.1392).

Augmentation 'noise': Significant difference in mean Cosine Similarity across LLMs (p=0.001796).

Augmentation 'paraphrase': Significant difference in mean Cosine Similarity across LLMs (p=3.825e-05).

Augmentation 'question_gen': Significant difference in mean Cosine Similarity across LLMs (p=0.0005087).

Augmentation 'shuffle': Significant difference in mean Cosine Similarity across LLMs (p=6.205e-05).

Augmentation 'summarize': Significant difference in mean Cosine Similarity across LLMs (p=0.01318).

Augmentation 'synonym': No significant difference in mean Cosine Similarity across LLMs (p=0.2035).

```
[41]: metric_col = 'wer'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):
          subset = df[df['augmentation_type'] == aug]
          groups = [subset[subset['model'] == m][metric_col].dropna().values
                     for m in subset['model'].unique()]
          if sum([len(g) > 1 for g in groups]) >= 2:
              stat, pval = f_oneway(*groups)
              anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                                   ↪ 'p-value': pval})
          else:
              anova_results.append({'augmentation_type': aug, 'F-statistic':
                                   ↪ float('nan'), 'p-value': float('nan')})

      anova_df = pd.DataFrame(anova_results)
      anova_df['Significant'] = anova_df['p-value'] < 0.05
      display(anova_df)

      print("\nSignificance Interpretation (Word Error Rate):")
      for idx, row in anova_df.iterrows():
          if pd.notna(row['p-value']):
              if row['Significant']:
```

```

        print(f"Augmentation '{row['augmentation_type']}': Significant_
↳difference in mean WER across LLMs (p={row['p-value']:.4g}).")
    else:
        print(f"Augmentation '{row['augmentation_type']}': No significant_
↳difference in mean WER across LLMs (p={row['p-value']:.4g}).")
    else:
        print(f"Augmentation '{row['augmentation_type']}': Insufficient data_
↳for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	4.048903	1.646887e-02	True
1	expand	5.770317	3.335208e-03	True
2	explain_simple	3.311658	3.434822e-02	True
3	identity	141.124096	5.255988e-17	True
4	negation	3.193739	3.874627e-02	True
5	noise	9.000695	4.430880e-04	True
6	paraphrase	4.997235	6.697015e-03	True
7	question_gen	9.699730	1.803215e-04	True
8	shuffle	10.068562	1.142896e-04	True
9	summarize	5.380230	4.722528e-03	True
10	synonym	8.927697	2.600684e-04	True

Significance Interpretation (Word Error Rate):

Augmentation 'entity_replace': Significant difference in mean WER across LLMs (p=0.01647).

Augmentation 'expand': Significant difference in mean WER across LLMs (p=0.003335).

Augmentation 'explain_simple': Significant difference in mean WER across LLMs (p=0.03435).

Augmentation 'identity': Significant difference in mean WER across LLMs (p=5.256e-17).

Augmentation 'negation': Significant difference in mean WER across LLMs (p=0.03875).

Augmentation 'noise': Significant difference in mean WER across LLMs (p=0.0004431).

Augmentation 'paraphrase': Significant difference in mean WER across LLMs (p=0.006697).

Augmentation 'question_gen': Significant difference in mean WER across LLMs (p=0.0001803).

Augmentation 'shuffle': Significant difference in mean WER across LLMs (p=0.0001143).

Augmentation 'summarize': Significant difference in mean WER across LLMs (p=0.004723).

Augmentation 'synonym': Significant difference in mean WER across LLMs (p=0.0002601).

```
[42]: metric_col = 'char_diversity'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                               'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                               float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Character Diversity):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean Character Diversity across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
                  difference in mean Character Diversity across LLMs (p={row['p-value']:.4g}).")
    else:
        print(f"Augmentation '{row['augmentation_type']}': Insufficient data
              for ANOVA.")
```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	9.366203	1.885479e-04	True
1	expand	15.506217	3.904275e-06	True
2	explain_simple	6.145506	2.404549e-03	True
3	identity	7.637759	6.995811e-04	True
4	negation	13.222636	1.464390e-05	True
5	noise	0.882166	4.655683e-01	False
6	paraphrase	29.943815	6.997905e-09	True
7	question_gen	46.338710	1.425922e-10	True
8	shuffle	0.392716	7.591815e-01	False
9	summarize	0.401780	7.528100e-01	False
10	synonym	11.899835	3.345085e-05	True

Significance Interpretation (Character Diversity):

Augmentation 'entity_replace': Significant difference in mean Character Diversity across LLMs ($p=0.0001885$).

Augmentation 'expand': Significant difference in mean Character Diversity across LLMs ($p=3.904e-06$).

Augmentation 'explain_simple': Significant difference in mean Character Diversity across LLMs ($p=0.002405$).

Augmentation 'identity': Significant difference in mean Character Diversity across LLMs ($p=0.0006996$).

Augmentation 'negation': Significant difference in mean Character Diversity across LLMs ($p=1.464e-05$).

Augmentation 'noise': No significant difference in mean Character Diversity across LLMs ($p=0.4656$).

Augmentation 'paraphrase': Significant difference in mean Character Diversity across LLMs ($p=6.998e-09$).

Augmentation 'question_gen': Significant difference in mean Character Diversity across LLMs ($p=1.426e-10$).

Augmentation 'shuffle': No significant difference in mean Character Diversity across LLMs ($p=0.7592$).

Augmentation 'summarize': No significant difference in mean Character Diversity across LLMs ($p=0.7528$).

Augmentation 'synonym': Significant difference in mean Character Diversity across LLMs ($p=3.345e-05$).

```
[43]: metric_col = 'type_token_ratio'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):
          subset = df[df['augmentation_type'] == aug]
          groups = [subset[subset['model'] == m][metric_col].dropna().values
                     for m in subset['model'].unique()]
          if sum([len(g) > 1 for g in groups]) >= 2:
              stat, pval = f_oneway(*groups)
              anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                                   'p-value': pval})
          else:
              anova_results.append({'augmentation_type': aug, 'F-statistic':
                                   float('nan'), 'p-value': float('nan')})

      anova_df = pd.DataFrame(anova_results)
      anova_df['Significant'] = anova_df['p-value'] < 0.05
      display(anova_df)

      print("\nSignificance Interpretation (Type Token Ratio):")
      for idx, row in anova_df.iterrows():
          if pd.notna(row['p-value']):
              if row['Significant']:
                  print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean Type Token Ratio across LLMs ( $p={row['p-value']:.4g}$ ).")
```

```

else:
    print(f"Augmentation '{row['augmentation_type']}': No significant
↪difference in mean Type Token Ratio across LLMs (p={row['p-value']:.4g}).")
else:
    print(f"Augmentation '{row['augmentation_type']}': Insufficient data
↪for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	20.346446	3.368696e-07	True
1	expand	33.260904	2.273976e-09	True
2	explain_simple	14.251822	7.948207e-06	True
3	identity	12.132673	2.882260e-05	True
4	negation	18.305278	8.993081e-07	True
5	noise	4.669504	1.134595e-02	True
6	paraphrase	9.965862	1.228372e-04	True
7	question_gen	36.764313	1.681870e-09	True
8	shuffle	7.630062	7.038690e-04	True
9	summarize	1.061596	3.811208e-01	False
10	synonym	10.262494	9.983533e-05	True

Significance Interpretation (Type Token Ratio):

Augmentation 'entity_replace': Significant difference in mean Type Token Ratio across LLMs (p=3.369e-07).

Augmentation 'expand': Significant difference in mean Type Token Ratio across LLMs (p=2.274e-09).

Augmentation 'explain_simple': Significant difference in mean Type Token Ratio across LLMs (p=7.948e-06).

Augmentation 'identity': Significant difference in mean Type Token Ratio across LLMs (p=2.882e-05).

Augmentation 'negation': Significant difference in mean Type Token Ratio across LLMs (p=8.993e-07).

Augmentation 'noise': Significant difference in mean Type Token Ratio across LLMs (p=0.01135).

Augmentation 'paraphrase': Significant difference in mean Type Token Ratio across LLMs (p=0.0001228).

Augmentation 'question_gen': Significant difference in mean Type Token Ratio across LLMs (p=1.682e-09).

Augmentation 'shuffle': Significant difference in mean Type Token Ratio across LLMs (p=0.0007039).

Augmentation 'summarize': No significant difference in mean Type Token Ratio across LLMs (p=0.3811).

Augmentation 'synonym': Significant difference in mean Type Token Ratio across LLMs (p=9.984e-05).

```

[44]: metric_col = 'bigram_overlap'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):

```

```

subset = df[df['augmentation_type'] == aug]
groups = [subset[subset['model'] == m][metric_col].dropna().values
           for m in subset['model'].unique()]
if sum([len(g) > 1 for g in groups]) >= 2:
    stat, pval = f_oneway(*groups)
    anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
↳ 'p-value': pval})
else:
    anova_results.append({'augmentation_type': aug, 'F-statistic':
↳ float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Bigram Overlap):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
↳ difference in mean Bigram Overlap across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
↳ difference in mean Bigram Overlap across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
↳ for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	8.484103	3.627449e-04	True
1	expand	0.573622	6.370908e-01	False
2	explain_simple	1.664026	1.973288e-01	False
3	identity	645.730192	6.116771e-26	True
4	negation	6.273466	2.154121e-03	True
5	noise	9.714903	2.807608e-04	True
6	paraphrase	2.946594	5.000476e-02	False
7	question_gen	2.092720	1.256118e-01	False
8	shuffle	23.042056	1.014598e-07	True
9	summarize	2.147089	1.166971e-01	False
10	synonym	2.112927	1.210820e-01	False

Significance Interpretation (Bigram Overlap):

Augmentation 'entity_replace': Significant difference in mean Bigram Overlap across LLMs (p=0.0003627).

Augmentation 'expand': No significant difference in mean Bigram Overlap across LLMs (p=0.6371).

Augmentation 'explain_simple': No significant difference in mean Bigram Overlap

across LLMs ($p=0.1973$).

Augmentation 'identity': Significant difference in mean Bigram Overlap across LLMs ($p=6.117e-26$).

Augmentation 'negation': Significant difference in mean Bigram Overlap across LLMs ($p=0.002154$).

Augmentation 'noise': Significant difference in mean Bigram Overlap across LLMs ($p=0.0002808$).

Augmentation 'paraphrase': No significant difference in mean Bigram Overlap across LLMs ($p=0.05$).

Augmentation 'question_gen': No significant difference in mean Bigram Overlap across LLMs ($p=0.1256$).

Augmentation 'shuffle': Significant difference in mean Bigram Overlap across LLMs ($p=1.015e-07$).

Augmentation 'summarize': No significant difference in mean Bigram Overlap across LLMs ($p=0.1167$).

Augmentation 'synonym': No significant difference in mean Bigram Overlap across LLMs ($p=0.1211$).

```
[45]: metric_col = 'total_duration_s'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):
          subset = df[df['augmentation_type'] == aug]
          groups = [subset[subset['model'] == m][metric_col].dropna().values
                     for m in subset['model'].unique()]
          if sum([len(g) > 1 for g in groups]) >= 2:
              stat, pval = f_oneway(*groups)
              anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                                   'p-value': pval})
          else:
              anova_results.append({'augmentation_type': aug, 'F-statistic':
                                   float('nan'), 'p-value': float('nan')})

      anova_df = pd.DataFrame(anova_results)
      anova_df['Significant'] = anova_df['p-value'] < 0.05
      display(anova_df)

      print("\nSignificance Interpretation (Total Duration (s)):")
      for idx, row in anova_df.iterrows():
          if pd.notna(row['p-value']):
              if row['Significant']:
                  print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean total duration (s) across LLMs (p={row['p-value']:.4g}).")
              else:
                  print(f"Augmentation '{row['augmentation_type']}': No significant
                  difference in mean total duration (s) across LLMs (p={row['p-value']:.4g}).")
              else:
```

```
print(f"Augmentation '{row['augmentation_type']}': Insufficient data_
↳for ANOVA.")
```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	88.881725	2.026440e-14	True
1	expand	2069.061530	5.865327e-33	True
2	explain_simple	1272.918072	5.061652e-30	True
3	identity	19.903110	4.145660e-07	True
4	negation	136.157393	8.400300e-17	True
5	noise	44.356838	1.690792e-09	True
6	paraphrase	2.900621	5.245446e-02	False
7	question_gen	349.253167	4.069736e-21	True
8	shuffle	14.522880	6.795688e-06	True
9	summarize	2.706015	6.430778e-02	False
10	synonym	70.628896	3.564170e-13	True

Significance Interpretation (Total Duration (s)):

Augmentation 'entity_replace': Significant difference in mean total duration (s) across LLMs (p=2.026e-14).

Augmentation 'expand': Significant difference in mean total duration (s) across LLMs (p=5.865e-33).

Augmentation 'explain_simple': Significant difference in mean total duration (s) across LLMs (p=5.062e-30).

Augmentation 'identity': Significant difference in mean total duration (s) across LLMs (p=4.146e-07).

Augmentation 'negation': Significant difference in mean total duration (s) across LLMs (p=8.4e-17).

Augmentation 'noise': Significant difference in mean total duration (s) across LLMs (p=1.691e-09).

Augmentation 'paraphrase': No significant difference in mean total duration (s) across LLMs (p=0.05245).

Augmentation 'question_gen': Significant difference in mean total duration (s) across LLMs (p=4.07e-21).

Augmentation 'shuffle': Significant difference in mean total duration (s) across LLMs (p=6.796e-06).

Augmentation 'summarize': No significant difference in mean total duration (s) across LLMs (p=0.06431).

Augmentation 'synonym': Significant difference in mean total duration (s) across LLMs (p=3.564e-13).

```
[46]: metric_col = 'prompt_eval_count'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):
          subset = df[df['augmentation_type'] == aug]
          groups = [subset[subset['model'] == m][metric_col].dropna().values
                     for m in subset['model'].unique()]
          if sum([len(g) > 1 for g in groups]) >= 2:
```

```

        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
        ↪ 'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
        ↪ float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Prompt Eval Count):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
            ↪ difference in mean Prompt Eval Count across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
            ↪ difference in mean Prompt Eval Count across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
            ↪ for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	184.551363	1.519865e-18	True
1	expand	184.551363	1.519865e-18	True
2	explain_simple	187.784906	1.206164e-18	True
3	identity	183.820126	1.602289e-18	True
4	negation	181.532495	1.892597e-18	True
5	noise	109.941443	2.146687e-13	True
6	paraphrase	183.820126	1.602289e-18	True
7	question_gen	180.483303	1.606188e-17	True
8	shuffle	194.251992	7.681570e-19	True
9	summarize	187.677568	1.215382e-18	True
10	synonym	191.588679	9.233930e-19	True

Significance Interpretation (Prompt Eval Count):

Augmentation 'entity_replace': Significant difference in mean Prompt Eval Count across LLMs (p=1.52e-18).

Augmentation 'expand': Significant difference in mean Prompt Eval Count across LLMs (p=1.52e-18).

Augmentation 'explain_simple': Significant difference in mean Prompt Eval Count across LLMs (p=1.206e-18).

Augmentation 'identity': Significant difference in mean Prompt Eval Count across LLMs (p=1.602e-18).

Augmentation 'negation': Significant difference in mean Prompt Eval Count across

LLMs ($p=1.893e-18$).
 Augmentation 'noise': Significant difference in mean Prompt Eval Count across LLMs ($p=2.147e-13$).
 Augmentation 'paraphrase': Significant difference in mean Prompt Eval Count across LLMs ($p=1.602e-18$).
 Augmentation 'question_gen': Significant difference in mean Prompt Eval Count across LLMs ($p=1.606e-17$).
 Augmentation 'shuffle': Significant difference in mean Prompt Eval Count across LLMs ($p=7.682e-19$).
 Augmentation 'summarize': Significant difference in mean Prompt Eval Count across LLMs ($p=1.215e-18$).
 Augmentation 'synonym': Significant difference in mean Prompt Eval Count across LLMs ($p=9.234e-19$).

```
[47]: metric_col = 'prompt_eval_duration_ns'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                              'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                              float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Prompt Eval Duration [ns]):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
difference in mean Prompt Eval Duration (ns) across LLMs (p={row['p-value']:.
4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
difference in mean Prompt Eval Duration (ns) across LLMs (p={row['p-value']:.
4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
for ANOVA.")
```

augmentation_type	F-statistic	p-value	Significant
-------------------	-------------	---------	-------------

0	entity_replace	58.697083	3.386839e-12	True
1	expand	93.546078	1.060680e-14	True
2	explain_simple	58.958322	3.210801e-12	True
3	identity	117.871546	5.481998e-16	True
4	negation	90.102476	1.705638e-14	True
5	noise	92.650919	1.234457e-12	True
6	paraphrase	30.775348	5.234439e-09	True
7	question_gen	99.581542	2.236007e-14	True
8	shuffle	59.311762	2.988070e-12	True
9	summarize	81.265498	6.247786e-14	True
10	synonym	64.419395	1.099528e-12	True

Significance Interpretation (Prompt Eval Duration [ns]):

Augmentation 'entity_replace': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=3.387e-12$).

Augmentation 'expand': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=1.061e-14$).

Augmentation 'explain_simple': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=3.211e-12$).

Augmentation 'identity': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=5.482e-16$).

Augmentation 'negation': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=1.706e-14$).

Augmentation 'noise': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=1.234e-12$).

Augmentation 'paraphrase': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=5.234e-09$).

Augmentation 'question_gen': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=2.236e-14$).

Augmentation 'shuffle': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=2.988e-12$).

Augmentation 'summarize': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=6.248e-14$).

Augmentation 'synonym': Significant difference in mean Prompt Eval Duration (ns) across LLMs ($p=1.1e-12$).

```
[48]: metric_col = 'eval_count'
      anova_results = []
      for aug in sorted(df['augmentation_type'].unique()):
          subset = df[df['augmentation_type'] == aug]
          groups = [subset[subset['model'] == m][metric_col].dropna().values
                    for m in subset['model'].unique()]
          if sum([len(g) > 1 for g in groups]) >= 2:
              stat, pval = f_oneway(*groups)
              anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                                     'p-value': pval})
          else:
```



```

        anova_results.append({'augmentation_type': aug, 'F-statistic':
↪float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Eval Count):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
↪difference in mean Eval Count across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
↪difference in mean Eval Count across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data
↪for ANOVA.")

```

C:\Users\parth\anaconda3\Lib\site-packages\scipy\stats_stats_py.py:4167:
ConstantInputWarning: Each of the input arrays is constant;the F statistic is
not defined or infinite

warnings.warn(stats.ConstantInputWarning(msg))

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	1.449674	2.495296e-01	False
1	expand	NaN	NaN	False
2	explain_simple	NaN	NaN	False
3	identity	3.540955	2.723647e-02	True
4	negation	6.485087	1.799039e-03	True
5	noise	4.094457	1.883038e-02	True
6	paraphrase	21.558619	1.939062e-07	True
7	question_gen	88.272611	9.341158e-14	True
8	shuffle	1.874725	1.567859e-01	False
9	summarize	27.612183	1.633497e-08	True
10	synonym	7.764514	6.328400e-04	True

Significance Interpretation (Eval Count):

Augmentation 'entity_replace': No significant difference in mean Eval Count
across LLMs (p=0.2495).

Augmentation 'expand': Insufficient data for ANOVA.

Augmentation 'explain_simple': Insufficient data for ANOVA.

Augmentation 'identity': Significant difference in mean Eval Count across LLMs
(p=0.02724).

Augmentation 'negation': Significant difference in mean Eval Count across LLMs
(p=0.001799).

Augmentation 'noise': Significant difference in mean Eval Count across LLMs (p=0.01883).

Augmentation 'paraphrase': Significant difference in mean Eval Count across LLMs (p=1.939e-07).

Augmentation 'question_gen': Significant difference in mean Eval Count across LLMs (p=9.341e-14).

Augmentation 'shuffle': No significant difference in mean Eval Count across LLMs (p=0.1568).

Augmentation 'summarize': Significant difference in mean Eval Count across LLMs (p=1.633e-08).

Augmentation 'synonym': Significant difference in mean Eval Count across LLMs (p=0.0006328).

```
[49]: metric_col = 'eval_duration_ns'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                              'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                              float('nan'), 'p-value': float('nan')})

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Eval Duration [ns]):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant
                  difference in mean Eval Duration (ns) across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant
                  difference in mean Eval Duration (ns) across LLMs (p={row['p-value']:.4g}).")
    else:
        print(f"Augmentation '{row['augmentation_type']}': Insufficient data
              for ANOVA.")
```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	72.447491	2.605049e-13	True
1	expand	239890.205769	7.888191e-62	True
2	explain_simple	73165.378112	1.307125e-54	True

3	identity	11.974073	3.189430e-05	True
4	negation	99.381020	4.911820e-15	True
5	noise	31.464285	3.869458e-08	True
6	paraphrase	72.540123	2.564261e-13	True
7	question_gen	254.463979	2.211799e-19	True
8	shuffle	9.738025	1.443391e-04	True
9	summarize	2.031106	1.322907e-01	False
10	synonym	52.952672	1.156670e-11	True

Significance Interpretation (Eval Duration [ns]):

Augmentation 'entity_replace': Significant difference in mean Eval Duration (ns) across LLMs (p=2.605e-13).

Augmentation 'expand': Significant difference in mean Eval Duration (ns) across LLMs (p=7.888e-62).

Augmentation 'explain_simple': Significant difference in mean Eval Duration (ns) across LLMs (p=1.307e-54).

Augmentation 'identity': Significant difference in mean Eval Duration (ns) across LLMs (p=3.189e-05).

Augmentation 'negation': Significant difference in mean Eval Duration (ns) across LLMs (p=4.912e-15).

Augmentation 'noise': Significant difference in mean Eval Duration (ns) across LLMs (p=3.869e-08).

Augmentation 'paraphrase': Significant difference in mean Eval Duration (ns) across LLMs (p=2.564e-13).

Augmentation 'question_gen': Significant difference in mean Eval Duration (ns) across LLMs (p=2.212e-19).

Augmentation 'shuffle': Significant difference in mean Eval Duration (ns) across LLMs (p=0.0001443).

Augmentation 'summarize': No significant difference in mean Eval Duration (ns) across LLMs (p=0.1323).

Augmentation 'synonym': Significant difference in mean Eval Duration (ns) across LLMs (p=1.157e-11).

```
[50]: metric_col = 'tokens_per_second'
anova_results = []
for aug in sorted(df['augmentation_type'].unique()):
    subset = df[df['augmentation_type'] == aug]
    groups = [subset[subset['model'] == m][metric_col].dropna().values
               for m in subset['model'].unique()]
    if sum([len(g) > 1 for g in groups]) >= 2:
        stat, pval = f_oneway(*groups)
        anova_results.append({'augmentation_type': aug, 'F-statistic': stat,
                              'p-value': pval})
    else:
        anova_results.append({'augmentation_type': aug, 'F-statistic':
                              float('nan'), 'p-value': float('nan')})
```

```

anova_df = pd.DataFrame(anova_results)
anova_df['Significant'] = anova_df['p-value'] < 0.05
display(anova_df)

print("\nSignificance Interpretation (Tokens per Second):")
for idx, row in anova_df.iterrows():
    if pd.notna(row['p-value']):
        if row['Significant']:
            print(f"Augmentation '{row['augmentation_type']}': Significant_
↳difference in mean Tokens per Second across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': No significant_
↳difference in mean Tokens per Second across LLMs (p={row['p-value']:.4g}).")
        else:
            print(f"Augmentation '{row['augmentation_type']}': Insufficient data_
↳for ANOVA.")

```

	augmentation_type	F-statistic	p-value	Significant
0	entity_replace	30649.552539	2.544171e-49	True
1	expand	198606.280785	1.109686e-60	True
2	explain_simple	84203.953624	1.828294e-55	True
3	identity	24097.032617	7.370114e-48	True
4	negation	12476.274974	7.370716e-44	True
5	noise	21347.177609	3.029594e-38	True
6	paraphrase	37493.323921	1.515228e-50	True
7	question_gen	7603.696943	2.258063e-38	True
8	shuffle	31266.724459	1.924724e-49	True
9	summarize	23577.888763	9.996357e-48	True
10	synonym	19457.175842	1.469829e-46	True

Significance Interpretation (Tokens per Second):

Augmentation 'entity_replace': Significant difference in mean Tokens per Second across LLMs (p=2.544e-49).

Augmentation 'expand': Significant difference in mean Tokens per Second across LLMs (p=1.11e-60).

Augmentation 'explain_simple': Significant difference in mean Tokens per Second across LLMs (p=1.828e-55).

Augmentation 'identity': Significant difference in mean Tokens per Second across LLMs (p=7.37e-48).

Augmentation 'negation': Significant difference in mean Tokens per Second across LLMs (p=7.371e-44).

Augmentation 'noise': Significant difference in mean Tokens per Second across LLMs (p=3.03e-38).

Augmentation 'paraphrase': Significant difference in mean Tokens per Second across LLMs (p=1.515e-50).

Augmentation 'question_gen': Significant difference in mean Tokens per Second across LLMs (p=2.258e-38).

Augmentation 'shuffle': Significant difference in mean Tokens per Second across LLMs ($p=1.925e-49$).

Augmentation 'summarize': Significant difference in mean Tokens per Second across LLMs ($p=9.996e-48$).

Augmentation 'synonym': Significant difference in mean Tokens per Second across LLMs ($p=1.47e-46$).

```
[51]: # This code runs the Kruskal-Wallis test for each metric, grouped by
      ↪ augmentation_type, and does pairwise Mann-Whitney U tests if significant.
      # It prints results for each metric and group.

      # Interpretation:
      # If  $p < 0.05$  in Kruskal-Wallis, at least one model is different for that
      ↪ metric/augmentation.
      # Mann-Whitney U will tell you which LLM pairs differ significantly.

      from scipy.stats import kruskal, mannwhitneyu
      from itertools import combinations
      import numpy as np

      metrics = [
          'levenshtein_similarity', 'jaccard_similarity', 'length_ratio', 'bleu',
          'cosine_similarity', 'wer', 'char_diversity', 'type_token_ratio',
          ↪ 'bigram_overlap',
          'total_duration_ns', 'load_duration_ns', 'prompt_eval_count',
          ↪ 'prompt_eval_duration_ns',
          'eval_count', 'eval_duration_ns', 'tokens_per_second'
      ]

      group_var = 'augmentation_type' # Or 'prompt_id'

      for metric_col in metrics:
          print(f"\n\n==== Kruskal-Wallis and Mann-Whitney U for {metric_col} (by
          ↪ {group_var}) =====")
          for group_val in sorted(df[group_var].unique()):
              sub = df[df[group_var] == group_val]
              models = sub['model'].unique()
              data = [sub[sub['model'] == m][metric_col].dropna().values for m in
              ↪ models]
              if len(models) > 1:
                  # Flatten all values for all groups, ignore nan
                  all_values = np.concatenate([d for d in data if len(d) > 0])
                  if len(all_values) == 0:
                      print(f" {group_var.capitalize()} '{group_val}': No data
                      ↪ available.")
                      continue
                  # Check if all values are identical
```

```

        if np.all(all_values == all_values[0]):
            print(f" {group_var.capitalize()} '{group_val}': All values_
↳ identical; cannot perform Kruskal-Wallis test.")
            continue
        stat, p = kruskal(*data)
        print(f"{group_var.capitalize()} '{group_val}': Kruskal-Wallis_
↳ H={stat:.3f}, p={p:.4g}")
        if p < 0.05:
            print(" Significant: Pairwise comparisons:")
            for m1, m2 in combinations(models, 2):
                d1 = sub[sub['model'] == m1][metric_col].dropna()
                d2 = sub[sub['model'] == m2][metric_col].dropna()
                # Mann-Whitney also needs at least one unique value
                if len(d1) > 0 and len(d2) > 0 and (not np.all(d1 == d1.
↳ iloc[0]) or not np.all(d2 == d2.iloc[0])):
                    try:
                        u, p_mw = mannwhitneyu(d1, d2,
↳ alternative='two-sided')
                        print(f" {m1} vs {m2}: U={u:.1f}, p={p_mw:.4g}")
                    except ValueError as e:
                        print(f" {m1} vs {m2}: Mann-Whitney not possible_
↳ ({e})")
                else:
                    print(" No significant difference among LLMs.")
            else:
                print(f"{group_var.capitalize()} '{group_val}': Only one model_
↳ present.")

```

==== Kruskal-Wallis and Mann-Whitney U for levenshtein_similarity (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=6.463, p=0.09113

No significant difference among LLMs.

Augmentation_type 'expand': Kruskal-Wallis H=0.509, p=0.917

No significant difference among LLMs.

Augmentation_type 'explain_simple': Kruskal-Wallis H=1.656, p=0.6467

No significant difference among LLMs.

Augmentation_type 'identity': Kruskal-Wallis H=21.072, p=0.0001017

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=17.0, p=0.1304

qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0004099

qwen2:0.5b vs granite3.1-moe:1b: U=40.0, p=0.4418

smollm2:360m vs gemma3:1b: U=0.0, p=0.0004099

smollm2:360m vs granite3.1-moe:1b: U=54.0, p=0.02067

gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0004099

Augmentation_type 'negation': Kruskal-Wallis H=5.327, p=0.1494

No significant difference among LLMs.

Augmentation_type 'noise': Kruskal-Wallis H=14.105, p=0.002766

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=19.0, p=0.1949
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=6.0, p=0.7111
- smollm2:360m vs gemma3:1b: U=8.0, p=0.01041
- smollm2:360m vs granite3.1-moe:1b: U=9.0, p=0.8889
- gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444

Augmentation_type 'paraphrase': Kruskal-Wallis H=3.991, p=0.2624

No significant difference among LLMs.

Augmentation_type 'question_gen': Kruskal-Wallis H=13.592, p=0.003516

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=37.0, p=0.6454
- qwen2:0.5b vs gemma3:1b: U=61.0, p=0.001088
- qwen2:0.5b vs granite3.1-moe:1b: U=36.0, p=0.1419
- smollm2:360m vs gemma3:1b: U=60.0, p=0.001865
- smollm2:360m vs granite3.1-moe:1b: U=32.0, p=0.345
- gemma3:1b vs granite3.1-moe:1b: U=8.0, p=0.04262

Augmentation_type 'shuffle': Kruskal-Wallis H=20.262, p=0.0001498

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=5.0, p=0.002953
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000931
- qwen2:0.5b vs granite3.1-moe:1b: U=12.0, p=0.03792
- smollm2:360m vs gemma3:1b: U=1.0, p=0.001348
- smollm2:360m vs granite3.1-moe:1b: U=36.0, p=0.7209
- gemma3:1b vs granite3.1-moe:1b: U=58.0, p=0.007362

Augmentation_type 'summarize': Kruskal-Wallis H=18.702, p=0.0003151

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=9.0, p=0.01476
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.8785
- smollm2:360m vs gemma3:1b: U=18.0, p=0.1605
- smollm2:360m vs granite3.1-moe:1b: U=55.0, p=0.01476
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

Augmentation_type 'synonym': Kruskal-Wallis H=2.645, p=0.4497

No significant difference among LLMs.

==== Kruskal-Wallis and Mann-Whitney U for jaccard_similarity (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=15.628, p=0.001352

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=24.0, p=0.4306
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000931
- qwen2:0.5b vs granite3.1-moe:1b: U=14.0, p=0.06588
- smollm2:360m vs gemma3:1b: U=3.0, p=0.001088
- smollm2:360m vs granite3.1-moe:1b: U=17.0, p=0.1304

gemma3:1b vs granite3.1-moe:1b: U=49.0, p=0.08298

Augmentation_type 'expand': Kruskal-Wallis H=6.976, p=0.07267

No significant difference among LLMs.

Augmentation_type 'explain_simple': Kruskal-Wallis H=2.332, p=0.5064

No significant difference among LLMs.

Augmentation_type 'identity': Kruskal-Wallis H=18.281, p=0.0003848

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=25.0, p=0.5054
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0004099
- qwen2:0.5b vs granite3.1-moe:1b: U=23.0, p=0.3823
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0004099
- smollm2:360m vs granite3.1-moe:1b: U=33.0, p=0.9591
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0004099

Augmentation_type 'negation': Kruskal-Wallis H=8.592, p=0.03524

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=41.0, p=0.3823
- qwen2:0.5b vs gemma3:1b: U=14.0, p=0.06496
- qwen2:0.5b vs granite3.1-moe:1b: U=26.5, p=0.5992
- smollm2:360m vs gemma3:1b: U=7.0, p=0.006993
- smollm2:360m vs granite3.1-moe:1b: U=20.0, p=0.2345
- gemma3:1b vs granite3.1-moe:1b: U=50.0, p=0.06496

Augmentation_type 'noise': Kruskal-Wallis H=12.904, p=0.004849

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=22.0, p=0.3282
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000931
- qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.4
- smollm2:360m vs gemma3:1b: U=11.0, p=0.0312
- smollm2:360m vs granite3.1-moe:1b: U=5.0, p=0.5333
- gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04949

Augmentation_type 'paraphrase': Kruskal-Wallis H=17.713, p=0.000504

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=64.0, p=0.000931
- qwen2:0.5b vs gemma3:1b: U=35.0, p=0.7927
- qwen2:0.5b vs granite3.1-moe:1b: U=46.0, p=0.1556
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
- smollm2:360m vs granite3.1-moe:1b: U=4.5, p=0.004427
- gemma3:1b vs granite3.1-moe:1b: U=44.0, p=0.2268

Augmentation_type 'question_gen': Kruskal-Wallis H=13.291, p=0.004047

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=46.0, p=0.1559
- qwen2:0.5b vs gemma3:1b: U=9.0, p=0.01796
- qwen2:0.5b vs granite3.1-moe:1b: U=12.0, p=0.1372
- smollm2:360m vs gemma3:1b: U=3.0, p=0.002742
- smollm2:360m vs granite3.1-moe:1b: U=5.0, p=0.01265
- gemma3:1b vs granite3.1-moe:1b: U=22.0, p=0.8463

Augmentation_type 'shuffle': Kruskal-Wallis H=12.812, p=0.005061

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=26.0, p=0.5737

qwen2:0.5b vs gemma3:1b: U=5.0, p=0.005351
 qwen2:0.5b vs granite3.1-moe:1b: U=17.0, p=0.1304
 smollm2:360m vs gemma3:1b: U=6.0, p=0.007362
 smollm2:360m vs granite3.1-moe:1b: U=24.0, p=0.4418
 gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01352
 Augmentation_type 'summarize': Kruskal-Wallis H=10.275, p=0.01637
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=54.0, p=0.02385
 qwen2:0.5b vs gemma3:1b: U=38.0, p=0.5632
 qwen2:0.5b vs granite3.1-moe:1b: U=32.0, p=1
 smollm2:360m vs gemma3:1b: U=9.5, p=0.02058
 smollm2:360m vs granite3.1-moe:1b: U=5.0, p=0.005284
 gemma3:1b vs granite3.1-moe:1b: U=23.5, p=0.3991
 Augmentation_type 'synonym': Kruskal-Wallis H=5.595, p=0.133
 No significant difference among LLMs.

==== Kruskal-Wallis and Mann-Whitney U for length_ratio (by augmentation_type)
 ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=7.583, p=0.05546
 No significant difference among LLMs.
 Augmentation_type 'expand': Kruskal-Wallis H=20.349, p=0.0001437
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=19.0, p=0.1949
 qwen2:0.5b vs gemma3:1b: U=60.0, p=0.001865
 qwen2:0.5b vs granite3.1-moe:1b: U=58.0, p=0.004662
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=63.0, p=0.0003108
 gemma3:1b vs granite3.1-moe:1b: U=30.0, p=0.8785
 Augmentation_type 'explain_simple': Kruskal-Wallis H=8.957, p=0.02986
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=34.0, p=0.8785
 qwen2:0.5b vs gemma3:1b: U=54.0, p=0.02067
 qwen2:0.5b vs granite3.1-moe:1b: U=52.0, p=0.03792
 smollm2:360m vs gemma3:1b: U=51.0, p=0.04988
 smollm2:360m vs granite3.1-moe:1b: U=50.0, p=0.06496
 gemma3:1b vs granite3.1-moe:1b: U=30.0, p=0.8785
 Augmentation_type 'identity': Kruskal-Wallis H=6.772, p=0.07954
 No significant difference among LLMs.
 Augmentation_type 'negation': Kruskal-Wallis H=5.442, p=0.1421
 No significant difference among LLMs.
 Augmentation_type 'noise': Kruskal-Wallis H=9.853, p=0.01986
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=51.5, p=0.04584
 qwen2:0.5b vs gemma3:1b: U=28.0, p=0.7209
 qwen2:0.5b vs granite3.1-moe:1b: U=11.0, p=0.5333
 smollm2:360m vs gemma3:1b: U=5.0, p=0.002953
 smollm2:360m vs granite3.1-moe:1b: U=4.0, p=0.4

gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444

Augmentation_type 'paraphrase': Kruskal-Wallis H=17.026, p=0.0006982

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=62.0, p=0.0006216
- qwen2:0.5b vs gemma3:1b: U=60.0, p=0.001865
- qwen2:0.5b vs granite3.1-moe:1b: U=63.0, p=0.0003108
- smollm2:360m vs gemma3:1b: U=16.0, p=0.1049
- smollm2:360m vs granite3.1-moe:1b: U=20.0, p=0.2345
- gemma3:1b vs granite3.1-moe:1b: U=39.0, p=0.5054

Augmentation_type 'question_gen': Kruskal-Wallis H=20.380, p=0.0001416

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=28.0, p=0.7209
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=2.0, p=0.002664
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
- smollm2:360m vs granite3.1-moe:1b: U=3.0, p=0.004662
- gemma3:1b vs granite3.1-moe:1b: U=33.0, p=0.2824

Augmentation_type 'shuffle': Kruskal-Wallis H=4.444, p=0.2173

No significant difference among LLMs.

Augmentation_type 'summarize': Kruskal-Wallis H=20.957, p=0.0001074

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=60.0, p=0.001865
- qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=59.0, p=0.002953
- smollm2:360m vs gemma3:1b: U=36.0, p=0.7209
- smollm2:360m vs granite3.1-moe:1b: U=12.0, p=0.03792
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554

Augmentation_type 'synonym': Kruskal-Wallis H=6.611, p=0.08539

No significant difference among LLMs.

==== Kruskal-Wallis and Mann-Whitney U for bleu (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=16.699, p=0.000815

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=28.0, p=0.7209
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=17.0, p=0.1304
- smollm2:360m vs gemma3:1b: U=1.0, p=0.0003108
- smollm2:360m vs granite3.1-moe:1b: U=24.0, p=0.4418
- gemma3:1b vs granite3.1-moe:1b: U=58.0, p=0.004662

Augmentation_type 'expand': Kruskal-Wallis H=1.363, p=0.7142

No significant difference among LLMs.

Augmentation_type 'explain_simple': Kruskal-Wallis H=4.384, p=0.2228

No significant difference among LLMs.

Augmentation_type 'identity': Kruskal-Wallis H=18.515, p=0.0003443

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=22.0, p=0.3282
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0004099

qwen2:0.5b vs granite3.1-moe:1b: U=26.0, p=0.5737
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0004099
 smollm2:360m vs granite3.1-moe:1b: U=39.0, p=0.5054
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0004099
 Augmentation_type 'negation': Kruskal-Wallis H=12.764, p=0.005175
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=41.0, p=0.3823
 qwen2:0.5b vs gemma3:1b: U=13.0, p=0.04988
 qwen2:0.5b vs granite3.1-moe:1b: U=34.0, p=0.8785
 smollm2:360m vs gemma3:1b: U=1.0, p=0.0003108
 smollm2:360m vs granite3.1-moe:1b: U=20.0, p=0.2345
 gemma3:1b vs granite3.1-moe:1b: U=59.0, p=0.002953
 Augmentation_type 'noise': Kruskal-Wallis H=13.708, p=0.003331
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=23.0, p=0.3823
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000931
 qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.4
 smollm2:360m vs gemma3:1b: U=8.0, p=0.01352
 smollm2:360m vs granite3.1-moe:1b: U=7.0, p=0.8889
 gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04949
 Augmentation_type 'paraphrase': Kruskal-Wallis H=11.199, p=0.0107
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=54.0, p=0.02067
 qwen2:0.5b vs gemma3:1b: U=23.0, p=0.3823
 qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.8785
 smollm2:360m vs gemma3:1b: U=5.0, p=0.002953
 smollm2:360m vs granite3.1-moe:1b: U=8.0, p=0.01041
 gemma3:1b vs granite3.1-moe:1b: U=43.0, p=0.2786
 Augmentation_type 'question_gen': Kruskal-Wallis H=9.025, p=0.02896
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=44.0, p=0.2345
 qwen2:0.5b vs gemma3:1b: U=18.0, p=0.1605
 qwen2:0.5b vs granite3.1-moe:1b: U=12.0, p=0.1419
 smollm2:360m vs gemma3:1b: U=8.0, p=0.01041
 smollm2:360m vs granite3.1-moe:1b: U=7.0, p=0.0293
 gemma3:1b vs granite3.1-moe:1b: U=21.0, p=0.7546
 Augmentation_type 'shuffle': Kruskal-Wallis H=12.465, p=0.005948
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=20.0, p=0.2345
 qwen2:0.5b vs gemma3:1b: U=8.0, p=0.01352
 qwen2:0.5b vs granite3.1-moe:1b: U=13.0, p=0.04988
 smollm2:360m vs gemma3:1b: U=8.0, p=0.01352
 smollm2:360m vs granite3.1-moe:1b: U=21.0, p=0.2786
 gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01352
 Augmentation_type 'summarize': Kruskal-Wallis H=6.739, p=0.08071
 No significant difference among LLMs.
 Augmentation_type 'synonym': Kruskal-Wallis H=5.062, p=0.1673
 No significant difference among LLMs.

```

==== Kruskal-Wallis and Mann-Whitney U for cosine_similarity (by
augmentation_type) ====
Augmentation_type 'entity_replace': Kruskal-Wallis H=10.707, p=0.01342
    Significant: Pairwise comparisons:
        qwen2:0.5b vs smollm2:360m: U=43.0, p=0.2786
        qwen2:0.5b vs gemma3:1b: U=15.0, p=0.08298
        qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.8785
        smollm2:360m vs gemma3:1b: U=2.0, p=0.0006216
        smollm2:360m vs granite3.1-moe:1b: U=16.0, p=0.1049
        gemma3:1b vs granite3.1-moe:1b: U=50.0, p=0.06496
Augmentation_type 'expand': Kruskal-Wallis H=0.474, p=0.9245
    No significant difference among LLMs.
Augmentation_type 'explain_simple': Kruskal-Wallis H=1.310, p=0.7268
    No significant difference among LLMs.
Augmentation_type 'identity': Kruskal-Wallis H=18.302, p=0.0003811
    Significant: Pairwise comparisons:
        qwen2:0.5b vs smollm2:360m: U=25.0, p=0.5054
        qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0004099
        qwen2:0.5b vs granite3.1-moe:1b: U=24.0, p=0.4418
        smollm2:360m vs gemma3:1b: U=0.0, p=0.0004099
        smollm2:360m vs granite3.1-moe:1b: U=27.0, p=0.6454
        gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0004099
Augmentation_type 'negation': Kruskal-Wallis H=8.043, p=0.04514
    Significant: Pairwise comparisons:
        qwen2:0.5b vs smollm2:360m: U=35.0, p=0.7984
        qwen2:0.5b vs gemma3:1b: U=9.0, p=0.01476
        qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.8785
        smollm2:360m vs gemma3:1b: U=11.0, p=0.02813
        smollm2:360m vs granite3.1-moe:1b: U=25.0, p=0.5054
        gemma3:1b vs granite3.1-moe:1b: U=52.0, p=0.03792
Augmentation_type 'noise': Kruskal-Wallis H=14.062, p=0.002822
    Significant: Pairwise comparisons:
        qwen2:0.5b vs smollm2:360m: U=17.0, p=0.1304
        qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
        qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.4
        smollm2:360m vs gemma3:1b: U=10.0, p=0.02067
        smollm2:360m vs granite3.1-moe:1b: U=9.0, p=0.8889
        gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444
Augmentation_type 'paraphrase': Kruskal-Wallis H=18.014, p=0.0004369
    Significant: Pairwise comparisons:
        qwen2:0.5b vs smollm2:360m: U=56.0, p=0.01041
        qwen2:0.5b vs gemma3:1b: U=8.0, p=0.01041
        qwen2:0.5b vs granite3.1-moe:1b: U=43.0, p=0.2786
        smollm2:360m vs gemma3:1b: U=5.0, p=0.002953
        smollm2:360m vs granite3.1-moe:1b: U=9.0, p=0.01476
        gemma3:1b vs granite3.1-moe:1b: U=63.0, p=0.0003108

```

Augmentation_type 'question_gen': Kruskal-Wallis H=14.934, p=0.001874

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=45.0, p=0.1949

qwen2:0.5b vs gemma3:1b: U=7.0, p=0.006993

qwen2:0.5b vs granite3.1-moe:1b: U=10.0, p=0.08125

smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554

smollm2:360m vs granite3.1-moe:1b: U=6.0, p=0.01998

gemma3:1b vs granite3.1-moe:1b: U=24.0, p=1

Augmentation_type 'shuffle': Kruskal-Wallis H=12.857, p=0.004955

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591

qwen2:0.5b vs gemma3:1b: U=6.0, p=0.007362

qwen2:0.5b vs granite3.1-moe:1b: U=21.0, p=0.2786

smollm2:360m vs gemma3:1b: U=5.0, p=0.005351

smollm2:360m vs granite3.1-moe:1b: U=19.0, p=0.1949

gemma3:1b vs granite3.1-moe:1b: U=57.0, p=0.01003

Augmentation_type 'summarize': Kruskal-Wallis H=10.281, p=0.01632

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=48.0, p=0.1049

qwen2:0.5b vs gemma3:1b: U=34.0, p=0.8785

qwen2:0.5b vs granite3.1-moe:1b: U=26.0, p=0.5737

smollm2:360m vs gemma3:1b: U=6.0, p=0.004662

smollm2:360m vs granite3.1-moe:1b: U=4.0, p=0.001865

gemma3:1b vs granite3.1-moe:1b: U=21.0, p=0.2786

Augmentation_type 'synonym': Kruskal-Wallis H=5.770, p=0.1234

No significant difference among LLMs.

==== Kruskal-Wallis and Mann-Whitney U for wer (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=11.207, p=0.01066

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=25.0, p=0.4942

qwen2:0.5b vs gemma3:1b: U=35.0, p=0.7923

qwen2:0.5b vs granite3.1-moe:1b: U=58.0, p=0.004662

smollm2:360m vs gemma3:1b: U=41.0, p=0.3706

smollm2:360m vs granite3.1-moe:1b: U=61.0, p=0.001088

gemma3:1b vs granite3.1-moe:1b: U=50.0, p=0.06588

Augmentation_type 'expand': Kruskal-Wallis H=12.001, p=0.00738

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=24.5, p=0.4619

qwen2:0.5b vs gemma3:1b: U=40.0, p=0.4306

qwen2:0.5b vs granite3.1-moe:1b: U=59.5, p=0.004545

smollm2:360m vs gemma3:1b: U=41.5, p=0.3439

smollm2:360m vs granite3.1-moe:1b: U=57.0, p=0.006993

gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01352

Augmentation_type 'explain_simple': Kruskal-Wallis H=7.563, p=0.05596

No significant difference among LLMs.

Augmentation_type 'identity': Kruskal-Wallis H=26.338, p=8.104e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=63.0, p=0.001348
- qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0004099
- qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0004054
- smollm2:360m vs granite3.1-moe:1b: U=29.5, p=0.8334
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0004099

Augmentation_type 'negation': Kruskal-Wallis H=12.225, p=0.006649

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=29.5, p=0.8335
- qwen2:0.5b vs gemma3:1b: U=20.0, p=0.2345
- qwen2:0.5b vs granite3.1-moe:1b: U=50.5, p=0.05852
- smollm2:360m vs gemma3:1b: U=17.0, p=0.1304
- smollm2:360m vs granite3.1-moe:1b: U=52.5, p=0.03556
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

Augmentation_type 'noise': Kruskal-Wallis H=13.909, p=0.003032

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=53.0, p=0.0312
- qwen2:0.5b vs gemma3:1b: U=62.0, p=0.001933
- qwen2:0.5b vs granite3.1-moe:1b: U=13.0, p=0.2667
- smollm2:360m vs gemma3:1b: U=52.0, p=0.04027
- smollm2:360m vs granite3.1-moe:1b: U=10.0, p=0.6944
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.04949

Augmentation_type 'paraphrase': Kruskal-Wallis H=12.062, p=0.007174

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=54.5, p=0.02077
- qwen2:0.5b vs gemma3:1b: U=51.0, p=0.04988
- qwen2:0.5b vs granite3.1-moe:1b: U=59.0, p=0.005351
- smollm2:360m vs gemma3:1b: U=20.5, p=0.2476
- smollm2:360m vs granite3.1-moe:1b: U=41.0, p=0.3706
- gemma3:1b vs granite3.1-moe:1b: U=51.0, p=0.05134

Augmentation_type 'question_gen': Kruskal-Wallis H=16.093, p=0.001085

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=17.5, p=0.1353
- qwen2:0.5b vs gemma3:1b: U=3.0, p=0.002742
- qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.4772
- smollm2:360m vs gemma3:1b: U=6.0, p=0.007232
- smollm2:360m vs granite3.1-moe:1b: U=41.0, p=0.03239
- gemma3:1b vs granite3.1-moe:1b: U=45.5, p=0.006646

Augmentation_type 'shuffle': Kruskal-Wallis H=14.024, p=0.002872

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=50.0, p=0.06496
- qwen2:0.5b vs gemma3:1b: U=61.0, p=0.002742
- qwen2:0.5b vs granite3.1-moe:1b: U=56.0, p=0.01041
- smollm2:360m vs gemma3:1b: U=52.0, p=0.04042
- smollm2:360m vs granite3.1-moe:1b: U=38.5, p=0.5283
- gemma3:1b vs granite3.1-moe:1b: U=13.0, p=0.05186

Augmentation_type 'summarize': Kruskal-Wallis H=12.826, p=0.005029

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=48.5, p=0.09265
- qwen2:0.5b vs gemma3:1b: U=63.0, p=0.001348
- qwen2:0.5b vs granite3.1-moe:1b: U=50.0, p=0.06496
- smollm2:360m vs gemma3:1b: U=59.0, p=0.005351
- smollm2:360m vs granite3.1-moe:1b: U=35.5, p=0.7525
- gemma3:1b vs granite3.1-moe:1b: U=24.0, p=0.4306

Augmentation_type 'synonym': Kruskal-Wallis H=17.847, p=0.000473

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=47.5, p=0.1146
- qwen2:0.5b vs gemma3:1b: U=15.0, p=0.0829
- qwen2:0.5b vs granite3.1-moe:1b: U=59.0, p=0.005351
- smollm2:360m vs gemma3:1b: U=4.5, p=0.004545
- smollm2:360m vs granite3.1-moe:1b: U=49.5, p=0.07399
- gemma3:1b vs granite3.1-moe:1b: U=62.5, p=0.001616

==== Kruskal-Wallis and Mann-Whitney U for char_diversity (by augmentation_type)
====

Augmentation_type 'entity_replace': Kruskal-Wallis H=16.813, p=0.0007721

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=44.0, p=0.2258
- qwen2:0.5b vs gemma3:1b: U=18.0, p=0.1559
- qwen2:0.5b vs granite3.1-moe:1b: U=58.0, p=0.007319
- smollm2:360m vs gemma3:1b: U=3.0, p=0.001088
- smollm2:360m vs granite3.1-moe:1b: U=42.5, p=0.2929
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.000931

Augmentation_type 'expand': Kruskal-Wallis H=17.293, p=0.000615

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=46.0, p=0.1559
- qwen2:0.5b vs gemma3:1b: U=7.0, p=0.01003
- qwen2:0.5b vs granite3.1-moe:1b: U=47.0, p=0.1256
- smollm2:360m vs gemma3:1b: U=0.5, p=0.001122
- smollm2:360m vs granite3.1-moe:1b: U=35.0, p=0.7923
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0009069

Augmentation_type 'explain_simple': Kruskal-Wallis H=13.851, p=0.003116

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=34.5, p=0.8335
- qwen2:0.5b vs gemma3:1b: U=8.0, p=0.01345
- qwen2:0.5b vs granite3.1-moe:1b: U=22.5, p=0.3442
- smollm2:360m vs gemma3:1b: U=2.5, p=0.002271
- smollm2:360m vs granite3.1-moe:1b: U=13.0, p=0.04988
- gemma3:1b vs granite3.1-moe:1b: U=57.0, p=0.009808

Augmentation_type 'identity': Kruskal-Wallis H=18.517, p=0.0003441

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=41.0, p=0.3696
- qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0004054
- qwen2:0.5b vs granite3.1-moe:1b: U=35.0, p=0.7921

smollm2:360m vs gemma3:1b: U=64.0, p=0.0004054
 smollm2:360m vs granite3.1-moe:1b: U=23.0, p=0.3696
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0004054
 Augmentation_type 'negation': Kruskal-Wallis H=14.712, p=0.00208
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=39.0, p=0.4847
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000891
 qwen2:0.5b vs granite3.1-moe:1b: U=27.0, p=0.6335
 smollm2:360m vs gemma3:1b: U=8.0, p=0.01324
 smollm2:360m vs granite3.1-moe:1b: U=25.0, p=0.4929
 gemma3:1b vs granite3.1-moe:1b: U=63.0, p=0.001348
 Augmentation_type 'noise': Kruskal-Wallis H=4.849, p=0.1832
 No significant difference among LLMs.
 Augmentation_type 'paraphrase': Kruskal-Wallis H=18.096, p=0.0004202
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=24.0, p=0.4278
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0009229
 qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.8743
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0009148
 smollm2:360m vs granite3.1-moe:1b: U=40.5, p=0.398
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.000891
 Augmentation_type 'question_gen': Kruskal-Wallis H=18.394, p=0.0003648
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=18.0, p=0.1553
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000931
 qwen2:0.5b vs granite3.1-moe:1b: U=15.0, p=0.272
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=23.0, p=0.9484
 gemma3:1b vs granite3.1-moe:1b: U=48.0, p=0.000666
 Augmentation_type 'shuffle': Kruskal-Wallis H=1.506, p=0.6808
 No significant difference among LLMs.
 Augmentation_type 'summarize': Kruskal-Wallis H=0.333, p=0.9536
 No significant difference among LLMs.
 Augmentation_type 'synonym': Kruskal-Wallis H=18.256, p=0.0003895
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=51.0, p=0.05117
 qwen2:0.5b vs gemma3:1b: U=6.0, p=0.007319
 qwen2:0.5b vs granite3.1-moe:1b: U=39.0, p=0.492
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0009229
 smollm2:360m vs granite3.1-moe:1b: U=20.0, p=0.2227
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.000891

 ==== Kruskal-Wallis and Mann-Whitney U for type_token_ratio (by
 augmentation_type) ====
 Augmentation_type 'entity_replace': Kruskal-Wallis H=19.190, p=0.0002498
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=18.0, p=0.1556

qwen2:0.5b vs gemma3:1b: U=60.0, p=0.001865
 qwen2:0.5b vs granite3.1-moe:1b: U=12.5, p=0.04441
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=24.5, p=0.4589
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009069
 Augmentation_type 'expand': Kruskal-Wallis H=20.286, p=0.0001481
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=32.5, p=1
 qwen2:0.5b vs gemma3:1b: U=61.0, p=0.002624
 qwen2:0.5b vs granite3.1-moe:1b: U=16.0, p=0.09796
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0008832
 smollm2:360m vs granite3.1-moe:1b: U=7.0, p=0.009431
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.000822
 Augmentation_type 'explain_simple': Kruskal-Wallis H=14.494, p=0.002304
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=30.5, p=0.9161
 qwen2:0.5b vs gemma3:1b: U=61.0, p=0.001088
 qwen2:0.5b vs granite3.1-moe:1b: U=30.5, p=0.9161
 smollm2:360m vs gemma3:1b: U=60.0, p=0.003798
 smollm2:360m vs granite3.1-moe:1b: U=27.0, p=0.6348
 gemma3:1b vs granite3.1-moe:1b: U=2.0, p=0.001933
 Augmentation_type 'identity': Kruskal-Wallis H=17.688, p=0.0005101
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=26.0, p=0.5624
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0009229
 qwen2:0.5b vs granite3.1-moe:1b: U=30.0, p=0.8745
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0009148
 smollm2:360m vs granite3.1-moe:1b: U=33.0, p=0.9578
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009148
 Augmentation_type 'negation': Kruskal-Wallis H=18.646, p=0.0003235
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=47.0, p=0.1304
 qwen2:0.5b vs gemma3:1b: U=63.0, p=0.001348
 qwen2:0.5b vs granite3.1-moe:1b: U=24.0, p=0.4295
 smollm2:360m vs gemma3:1b: U=63.0, p=0.001348
 smollm2:360m vs granite3.1-moe:1b: U=19.0, p=0.1857
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0008989
 Augmentation_type 'noise': Kruskal-Wallis H=11.094, p=0.01123
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=24.0, p=0.4295
 qwen2:0.5b vs gemma3:1b: U=60.0, p=0.00385
 qwen2:0.5b vs granite3.1-moe:1b: U=5.5, p=0.6004
 smollm2:360m vs gemma3:1b: U=56.0, p=0.01345
 smollm2:360m vs granite3.1-moe:1b: U=8.0, p=1
 gemma3:1b vs granite3.1-moe:1b: U=1.0, p=0.08868
 Augmentation_type 'paraphrase': Kruskal-Wallis H=13.738, p=0.003285
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=29.5, p=0.8333

qwen2:0.5b vs gemma3:1b: U=58.0, p=0.007276
 qwen2:0.5b vs granite3.1-moe:1b: U=21.0, p=0.2694
 smollm2:360m vs gemma3:1b: U=58.0, p=0.007276
 smollm2:360m vs granite3.1-moe:1b: U=25.0, p=0.4916
 gemma3:1b vs granite3.1-moe:1b: U=2.0, p=0.001903
 Augmentation_type 'question_gen': Kruskal-Wallis H=20.201, p=0.0001542
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=22.5, p=0.2143
 qwen2:0.5b vs gemma3:1b: U=63.5, p=0.0009342
 qwen2:0.5b vs granite3.1-moe:1b: U=27.0, p=0.7304
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0005479
 smollm2:360m vs granite3.1-moe:1b: U=36.0, p=0.08359
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.002362
 Augmentation_type 'shuffle': Kruskal-Wallis H=18.362, p=0.0003704
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=43.0, p=0.268
 qwen2:0.5b vs gemma3:1b: U=63.0, p=0.001326
 qwen2:0.5b vs granite3.1-moe:1b: U=20.0, p=0.2146
 smollm2:360m vs gemma3:1b: U=60.0, p=0.003824
 smollm2:360m vs granite3.1-moe:1b: U=13.5, p=0.05464
 gemma3:1b vs granite3.1-moe:1b: U=1.0, p=0.001241
 Augmentation_type 'summarize': Kruskal-Wallis H=3.078, p=0.3798
 No significant difference among LLMs.
 Augmentation_type 'synonym': Kruskal-Wallis H=15.644, p=0.001342
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=22.5, p=0.3314
 qwen2:0.5b vs gemma3:1b: U=59.0, p=0.005183
 qwen2:0.5b vs granite3.1-moe:1b: U=30.5, p=0.9155
 smollm2:360m vs gemma3:1b: U=63.0, p=0.001305
 smollm2:360m vs granite3.1-moe:1b: U=43.5, p=0.2441
 gemma3:1b vs granite3.1-moe:1b: U=3.0, p=0.002722

 ==== Kruskal-Wallis and Mann-Whitney U for bigram_overlap (by augmentation_type)
 ====
 Augmentation_type 'entity_replace': Kruskal-Wallis H=16.003, p=0.001132
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=28.0, p=0.7112
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.000931
 qwen2:0.5b vs granite3.1-moe:1b: U=18.0, p=0.1529
 smollm2:360m vs gemma3:1b: U=1.0, p=0.001348
 smollm2:360m vs granite3.1-moe:1b: U=23.0, p=0.3681
 gemma3:1b vs granite3.1-moe:1b: U=55.0, p=0.01796
 Augmentation_type 'expand': Kruskal-Wallis H=1.631, p=0.6524
 No significant difference among LLMs.
 Augmentation_type 'explain_simple': Kruskal-Wallis H=5.506, p=0.1383
 No significant difference among LLMs.
 Augmentation_type 'identity': Kruskal-Wallis H=18.329, p=0.0003762

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=22.5, p=0.341
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0003918
- qwen2:0.5b vs granite3.1-moe:1b: U=27.0, p=0.6355
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0004099
- smollm2:360m vs granite3.1-moe:1b: U=36.0, p=0.7209
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0004099

Augmentation_type 'negation': Kruskal-Wallis H=12.423, p=0.006065

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=41.5, p=0.341
- qwen2:0.5b vs gemma3:1b: U=12.0, p=0.03792
- qwen2:0.5b vs granite3.1-moe:1b: U=33.5, p=0.9163
- smollm2:360m vs gemma3:1b: U=1.0, p=0.001315
- smollm2:360m vs granite3.1-moe:1b: U=21.5, p=0.29
- gemma3:1b vs granite3.1-moe:1b: U=57.0, p=0.006993

Augmentation_type 'noise': Kruskal-Wallis H=13.667, p=0.003396

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=24.0, p=0.4275
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0009229
- qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.3593
- smollm2:360m vs gemma3:1b: U=8.0, p=0.01345
- smollm2:360m vs granite3.1-moe:1b: U=6.0, p=0.6944
- gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04949

Augmentation_type 'paraphrase': Kruskal-Wallis H=10.377, p=0.01562

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=57.0, p=0.008244
- qwen2:0.5b vs gemma3:1b: U=32.0, p=1
- qwen2:0.5b vs granite3.1-moe:1b: U=40.0, p=0.4418
- smollm2:360m vs gemma3:1b: U=10.0, p=0.01833
- smollm2:360m vs granite3.1-moe:1b: U=8.5, p=0.01306
- gemma3:1b vs granite3.1-moe:1b: U=44.0, p=0.2268

Augmentation_type 'question_gen': Kruskal-Wallis H=5.139, p=0.1619

No significant difference among LLMs.

Augmentation_type 'shuffle': Kruskal-Wallis H=13.706, p=0.003334

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=20.0, p=0.2149
- qwen2:0.5b vs gemma3:1b: U=6.0, p=0.006521
- qwen2:0.5b vs granite3.1-moe:1b: U=13.0, p=0.0486
- smollm2:360m vs gemma3:1b: U=7.0, p=0.009808
- smollm2:360m vs granite3.1-moe:1b: U=22.0, p=0.317
- gemma3:1b vs granite3.1-moe:1b: U=56.5, p=0.0116

Augmentation_type 'summarize': Kruskal-Wallis H=3.779, p=0.2863

No significant difference among LLMs.

Augmentation_type 'synonym': Kruskal-Wallis H=4.229, p=0.2377

No significant difference among LLMs.

==== Kruskal-Wallis and Mann-Whitney U for total_duration_ns (by

```

augmentation_type) ====
Augmentation_type 'entity_replace': Kruskal-Wallis H=25.997, p=9.55e-06
  Significant: Pairwise comparisons:
    qwen2:0.5b vs smollm2:360m: U=5.0, p=0.002953
    qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
    qwen2:0.5b vs granite3.1-moe:1b: U=1.0, p=0.0003108
    smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
    smollm2:360m vs granite3.1-moe:1b: U=56.0, p=0.01041
    gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'expand': Kruskal-Wallis H=29.091, p=2.143e-06
  Significant: Pairwise comparisons:
    qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
    qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
    qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
    smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
    smollm2:360m vs granite3.1-moe:1b: U=64.0, p=0.0001554
    gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'explain_simple': Kruskal-Wallis H=29.091, p=2.143e-06
  Significant: Pairwise comparisons:
    qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
    qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
    qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
    smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
    smollm2:360m vs granite3.1-moe:1b: U=64.0, p=0.0001554
    gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'identity': Kruskal-Wallis H=21.702, p=7.525e-05
  Significant: Pairwise comparisons:
    qwen2:0.5b vs smollm2:360m: U=6.0, p=0.004662
    qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
    qwen2:0.5b vs granite3.1-moe:1b: U=7.0, p=0.006993
    smollm2:360m vs gemma3:1b: U=7.0, p=0.006993
    smollm2:360m vs granite3.1-moe:1b: U=43.0, p=0.2786
    gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'negation': Kruskal-Wallis H=27.003, p=5.879e-06
  Significant: Pairwise comparisons:
    qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
    qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
    qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
    smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
    smollm2:360m vs granite3.1-moe:1b: U=49.0, p=0.08298
    gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'noise': Kruskal-Wallis H=18.876, p=0.00029
  Significant: Pairwise comparisons:
    qwen2:0.5b vs smollm2:360m: U=13.0, p=0.04988
    qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
    qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.04444
    smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
    smollm2:360m vs granite3.1-moe:1b: U=8.0, p=1

```

gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444

Augmentation_type 'paraphrase': Kruskal-Wallis H=16.690, p=0.0008183

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=18.0, p=0.1605
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=7.0, p=0.006993
- smollm2:360m vs gemma3:1b: U=7.0, p=0.006993
- smollm2:360m vs granite3.1-moe:1b: U=31.0, p=0.9591
- gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01041

Augmentation_type 'question_gen': Kruskal-Wallis H=26.972, p=5.969e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.000666
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
- smollm2:360m vs granite3.1-moe:1b: U=1.0, p=0.001332
- gemma3:1b vs granite3.1-moe:1b: U=48.0, p=0.000666

Augmentation_type 'shuffle': Kruskal-Wallis H=19.486, p=0.0002169

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=6.0, p=0.004662
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
- smollm2:360m vs gemma3:1b: U=15.0, p=0.08298
- smollm2:360m vs granite3.1-moe:1b: U=40.0, p=0.4418
- gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01041

Augmentation_type 'summarize': Kruskal-Wallis H=13.213, p=0.004198

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=42.0, p=0.3282
- qwen2:0.5b vs gemma3:1b: U=2.0, p=0.0006216
- qwen2:0.5b vs granite3.1-moe:1b: U=1.0, p=0.0003108
- smollm2:360m vs gemma3:1b: U=16.0, p=0.1049
- smollm2:360m vs granite3.1-moe:1b: U=16.0, p=0.1049
- gemma3:1b vs granite3.1-moe:1b: U=15.0, p=0.08298

Augmentation_type 'synonym': Kruskal-Wallis H=24.207, p=2.261e-05

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=8.0, p=0.01041
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
- smollm2:360m vs granite3.1-moe:1b: U=31.0, p=0.9591
- gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

==== Kruskal-Wallis and Mann-Whitney U for load_duration_ns (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=27.318, p=5.049e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554

qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=12.0, p=0.03792
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'expand': Kruskal-Wallis H=26.227, p=8.547e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=28.0, p=0.7209
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'explain_simple': Kruskal-Wallis H=26.364, p=8.003e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=40.0, p=0.4418
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'identity': Kruskal-Wallis H=28.253, p=3.214e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=5.0, p=0.002953
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'negation': Kruskal-Wallis H=26.207, p=8.63e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=29.0, p=0.7984
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'noise': Kruskal-Wallis H=22.538, p=5.039e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=16.0, p=0.04444
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=14.0, p=0.1778
 gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444
 Augmentation_type 'paraphrase': Kruskal-Wallis H=15.372, p=0.001525
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=56.0, p=0.01041

qwen2:0.5b vs gemma3:1b: U=7.0, p=0.006993
 qwen2:0.5b vs granite3.1-moe:1b: U=56.0, p=0.01041
 smollm2:360m vs gemma3:1b: U=7.0, p=0.006993
 smollm2:360m vs granite3.1-moe:1b: U=40.0, p=0.4418
 gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01041
 Augmentation_type 'question_gen': Kruskal-Wallis H=25.015, p=1.533e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=48.0, p=0.000666
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=21.0, p=0.7546
 gemma3:1b vs granite3.1-moe:1b: U=48.0, p=0.000666
 Augmentation_type 'shuffle': Kruskal-Wallis H=26.466, p=7.618e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=42.0, p=0.3282
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'summarize': Kruskal-Wallis H=27.003, p=5.879e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=49.0, p=0.08298
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'synonym': Kruskal-Wallis H=26.412, p=7.819e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=41.0, p=0.3823
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

==== Kruskal-Wallis and Mann-Whitney U for prompt_eval_count (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=26.267, p=8.385e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
 qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
 qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554

gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'expand': Kruskal-Wallis H=26.267, p=8.385e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
- qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'explain_simple': Kruskal-Wallis H=26.450, p=7.677e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
- qwen2:0.5b vs gemma3:1b: U=23.5, p=0.3963
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'identity': Kruskal-Wallis H=26.267, p=8.385e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
- qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'negation': Kruskal-Wallis H=26.290, p=8.293e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
- qwen2:0.5b vs gemma3:1b: U=36.0, p=0.7105
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'noise': Kruskal-Wallis H=18.577, p=0.0003343

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
- qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0488
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.04444
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0488

Augmentation_type 'paraphrase': Kruskal-Wallis H=26.267, p=8.385e-06

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
- qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
- smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554

gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229
Augmentation_type 'question_gen': Kruskal-Wallis H=23.945, p=2.565e-05

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.002335
smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.002388
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.002335

Augmentation_type 'shuffle': Kruskal-Wallis H=26.290, p=8.293e-06

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
qwen2:0.5b vs gemma3:1b: U=36.0, p=0.7105
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'summarize': Kruskal-Wallis H=26.290, p=8.293e-06

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
qwen2:0.5b vs gemma3:1b: U=36.0, p=0.7105
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

Augmentation_type 'synonym': Kruskal-Wallis H=26.267, p=8.385e-06

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0009229
qwen2:0.5b vs gemma3:1b: U=29.5, p=0.8319
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0009229
smollm2:360m vs gemma3:1b: U=64.0, p=0.0009229
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0009229

==== Kruskal-Wallis and Mann-Whitney U for prompt_eval_duration_ns (by augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=24.136, p=2.339e-05

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=5.0, p=0.002953
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=3.0, p=0.001088
gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

Augmentation_type 'expand': Kruskal-Wallis H=25.389, p=1.28e-05

Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=30.0, p=0.8785

qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=2.0, p=0.0006216
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=1.0, p=0.0003108
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'explain_simple': Kruskal-Wallis H=24.136, p=2.339e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=5.0, p=0.002953
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=3.0, p=0.001088
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'identity': Kruskal-Wallis H=25.918, p=9.924e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=30.0, p=0.8785
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=1.0, p=0.0003108
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'negation': Kruskal-Wallis H=25.128, p=1.452e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=2.0, p=0.0006216
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=2.0, p=0.0006216
 gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 Augmentation_type 'noise': Kruskal-Wallis H=18.173, p=0.0004051
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=30.0, p=0.8785
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=1.0, p=0.08889
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.04444
 gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444
 Augmentation_type 'paraphrase': Kruskal-Wallis H=22.438, p=5.289e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=28.0, p=0.7209
 qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=3.0, p=0.001088
 smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=7.0, p=0.006993
 gemma3:1b vs granite3.1-moe:1b: U=57.0, p=0.006993
 Augmentation_type 'question_gen': Kruskal-Wallis H=23.572, p=3.069e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=30.0, p=0.8785

```

qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=1.0, p=0.001332
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.000666
gemma3:1b vs granite3.1-moe:1b: U=48.0, p=0.000666
Augmentation_type 'shuffle': Kruskal-Wallis H=24.378, p=2.083e-05
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=5.0, p=0.002953
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=2.0, p=0.0006216
gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'summarize': Kruskal-Wallis H=25.125, p=1.454e-05
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=3.0, p=0.001088
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=1.0, p=0.0003108
gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'synonym': Kruskal-Wallis H=24.622, p=1.852e-05
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=31.0, p=0.9591
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.001865
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=2.0, p=0.0006216
gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

==== Kruskal-Wallis and Mann-Whitney U for eval_count (by augmentation_type)
====
Augmentation_type 'entity_replace': Kruskal-Wallis H=7.167, p=0.06675
No significant difference among LLMs.
Augmentation_type 'expand': All values identical; cannot perform
Kruskal-Wallis test.
Augmentation_type 'explain_simple': All values identical; cannot perform
Kruskal-Wallis test.
Augmentation_type 'identity': Kruskal-Wallis H=7.916, p=0.04778
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=50.5, p=0.04809
qwen2:0.5b vs gemma3:1b: U=48.0, p=0.09413
qwen2:0.5b vs granite3.1-moe:1b: U=35.0, p=0.7496
smollm2:360m vs gemma3:1b: U=34.5, p=0.8332
smollm2:360m vs granite3.1-moe:1b: U=14.5, p=0.06657
gemma3:1b vs granite3.1-moe:1b: U=11.5, p=0.03261
Augmentation_type 'negation': Kruskal-Wallis H=17.854, p=0.0004715

```

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=42.5, p=0.2897
- qwen2:0.5b vs gemma3:1b: U=16.0, p=0.03247
- qwen2:0.5b vs granite3.1-moe:1b: U=18.0, p=0.08474
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0004099
- smollm2:360m vs granite3.1-moe:1b: U=2.0, p=0.00122
- gemma3:1b vs granite3.1-moe:1b: U=36.0, p=0.3816

Augmentation_type 'noise': Kruskal-Wallis H=10.862, p=0.0125

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=45.0, p=0.177
- qwen2:0.5b vs gemma3:1b: U=16.0, p=0.03247
- qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.3032
- smollm2:360m vs gemma3:1b: U=8.0, p=0.004569
- smollm2:360m vs granite3.1-moe:1b: U=2.0, p=0.1384

Augmentation_type 'paraphrase': Kruskal-Wallis H=23.092, p=3.865e-05

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=59.5, p=0.002435
- qwen2:0.5b vs gemma3:1b: U=28.0, p=0.3816
- qwen2:0.5b vs granite3.1-moe:1b: U=28.0, p=0.3816
- smollm2:360m vs gemma3:1b: U=4.0, p=0.001446
- smollm2:360m vs granite3.1-moe:1b: U=4.0, p=0.001446

Augmentation_type 'question_gen': Kruskal-Wallis H=23.453, p=3.249e-05

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=31.5, p=1
- qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0004099
- qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.002157
- smollm2:360m vs gemma3:1b: U=0.0, p=0.0004008
- smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.002107
- gemma3:1b vs granite3.1-moe:1b: U=32.0, p=0.1121

Augmentation_type 'shuffle': Kruskal-Wallis H=4.933, p=0.1768

No significant difference among LLMs.

Augmentation_type 'summarize': Kruskal-Wallis H=19.974, p=0.0001719

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=57.0, p=0.007185
- qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0006754
- qwen2:0.5b vs granite3.1-moe:1b: U=34.0, p=0.8481
- smollm2:360m vs gemma3:1b: U=35.0, p=0.7918
- smollm2:360m vs granite3.1-moe:1b: U=6.5, p=0.006978
- gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0007781

Augmentation_type 'synonym': Kruskal-Wallis H=17.606, p=0.0005304

Significant: Pairwise comparisons:

- qwen2:0.5b vs smollm2:360m: U=44.0, p=0.2202
- qwen2:0.5b vs gemma3:1b: U=16.0, p=0.03247
- qwen2:0.5b vs granite3.1-moe:1b: U=16.0, p=0.03247
- smollm2:360m vs gemma3:1b: U=4.0, p=0.00146
- smollm2:360m vs granite3.1-moe:1b: U=4.0, p=0.00146

```

==== Kruskal-Wallis and Mann-Whitney U for eval_duration_ns (by
augmentation_type) ====
Augmentation_type 'entity_replace': Kruskal-Wallis H=25.912, p=9.951e-06
Significant: Pairwise comparisons:
  qwen2:0.5b vs smollm2:360m: U=4.0, p=0.001865
  qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
  qwen2:0.5b vs granite3.1-moe:1b: U=4.0, p=0.001865
  smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
  smollm2:360m vs granite3.1-moe:1b: U=57.0, p=0.006993
  gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'expand': Kruskal-Wallis H=29.091, p=2.143e-06
Significant: Pairwise comparisons:
  qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
  qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
  qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
  smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
  smollm2:360m vs granite3.1-moe:1b: U=64.0, p=0.0001554
  gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'explain_simple': Kruskal-Wallis H=29.091, p=2.143e-06
Significant: Pairwise comparisons:
  qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
  qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
  qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
  smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
  smollm2:360m vs granite3.1-moe:1b: U=64.0, p=0.0001554
  gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'identity': Kruskal-Wallis H=18.659, p=0.0003216
Significant: Pairwise comparisons:
  qwen2:0.5b vs smollm2:360m: U=6.0, p=0.004662
  qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
  qwen2:0.5b vs granite3.1-moe:1b: U=12.0, p=0.03792
  smollm2:360m vs gemma3:1b: U=22.0, p=0.3282
  smollm2:360m vs granite3.1-moe:1b: U=46.0, p=0.1605
  gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'negation': Kruskal-Wallis H=27.818, p=3.966e-06
Significant: Pairwise comparisons:
  qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
  qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
  qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
  smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
  smollm2:360m vs granite3.1-moe:1b: U=56.0, p=0.01041
  gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'noise': Kruskal-Wallis H=19.731, p=0.000193
Significant: Pairwise comparisons:
  qwen2:0.5b vs smollm2:360m: U=9.0, p=0.01476
  qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
  qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.04444
  smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554

```

smollm2:360m vs granite3.1-moe:1b: U=8.0, p=1
gemma3:1b vs granite3.1-moe:1b: U=16.0, p=0.04444
Augmentation_type 'paraphrase': Kruskal-Wallis H=22.366, p=5.472e-05
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=16.0, p=0.1049
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=41.0, p=0.3823
gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554
Augmentation_type 'question_gen': Kruskal-Wallis H=26.640, p=7.003e-06
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=0.0, p=0.0001554
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.000666
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=3.0, p=0.004662
gemma3:1b vs granite3.1-moe:1b: U=48.0, p=0.000666
Augmentation_type 'shuffle': Kruskal-Wallis H=16.656, p=0.0008316
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=6.0, p=0.004662
qwen2:0.5b vs gemma3:1b: U=6.0, p=0.004662
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=17.0, p=0.1304
smollm2:360m vs granite3.1-moe:1b: U=42.0, p=0.3282
gemma3:1b vs granite3.1-moe:1b: U=56.0, p=0.01041
Augmentation_type 'summarize': Kruskal-Wallis H=10.261, p=0.01647
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=42.0, p=0.3282
qwen2:0.5b vs gemma3:1b: U=39.0, p=0.5054
qwen2:0.5b vs granite3.1-moe:1b: U=5.0, p=0.002953
smollm2:360m vs gemma3:1b: U=24.0, p=0.4418
smollm2:360m vs granite3.1-moe:1b: U=16.0, p=0.1049
gemma3:1b vs granite3.1-moe:1b: U=3.0, p=0.001088
Augmentation_type 'synonym': Kruskal-Wallis H=24.182, p=2.289e-05
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=8.0, p=0.01041
qwen2:0.5b vs gemma3:1b: U=0.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=0.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=0.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=32.0, p=1
gemma3:1b vs granite3.1-moe:1b: U=64.0, p=0.0001554

==== Kruskal-Wallis and Mann-Whitney U for tokens_per_second (by
augmentation_type) ====

Augmentation_type 'entity_replace': Kruskal-Wallis H=29.091, p=2.143e-06
Significant: Pairwise comparisons:

qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
 Augmentation_type 'expand': Kruskal-Wallis H=29.091, p=2.143e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
 Augmentation_type 'explain_simple': Kruskal-Wallis H=29.091, p=2.143e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
 Augmentation_type 'identity': Kruskal-Wallis H=29.091, p=2.143e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
 Augmentation_type 'negation': Kruskal-Wallis H=29.091, p=2.143e-06
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
 Augmentation_type 'noise': Kruskal-Wallis H=22.838, p=4.366e-05
 Significant: Pairwise comparisons:
 qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
 qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
 qwen2:0.5b vs granite3.1-moe:1b: U=16.0, p=0.04444
 smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
 smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.04444
 gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.04444
 Augmentation_type 'paraphrase': Kruskal-Wallis H=29.091, p=2.143e-06
 Significant: Pairwise comparisons:

```

qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
Augmentation_type 'question_gen': Kruskal-Wallis H=27.148, p=5.48e-06
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=48.0, p=0.000666
smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.000666
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.000666
Augmentation_type 'shuffle': Kruskal-Wallis H=29.091, p=2.143e-06
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
Augmentation_type 'summarize': Kruskal-Wallis H=29.091, p=2.143e-06
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554
Augmentation_type 'synonym': Kruskal-Wallis H=29.091, p=2.143e-06
Significant: Pairwise comparisons:
qwen2:0.5b vs smollm2:360m: U=64.0, p=0.0001554
qwen2:0.5b vs gemma3:1b: U=64.0, p=0.0001554
qwen2:0.5b vs granite3.1-moe:1b: U=64.0, p=0.0001554
smollm2:360m vs gemma3:1b: U=64.0, p=0.0001554
smollm2:360m vs granite3.1-moe:1b: U=0.0, p=0.0001554
gemma3:1b vs granite3.1-moe:1b: U=0.0, p=0.0001554

```

```

[52]: # Friedman's Test (by augmentation_type)

from scipy.stats import friedmanchisquare

metrics = [
    'levenshtein_similarity', 'jaccard_similarity', 'length_ratio', 'bleu',
    'cosine_similarity', 'wer', 'char_diversity', 'type_token_ratio',
    'bigram_overlap',

```



```

    'total_duration_ns', 'load_duration_ns', 'prompt_eval_count',
    ↪ 'prompt_eval_duration_ns',
    'eval_count', 'eval_duration_ns', 'tokens_per_second'
]

group_var = 'augmentation_type'    # Or use 'prompt_id' if you want

for metric_col in metrics:
    print(f"\n=== Friedman's Test for {metric_col} (by {group_var}) ===")
    for group_val in sorted(df[group_var].unique()):
        sub = df[df[group_var] == group_val]
        pivot = sub.pivot(index='prompt_id', columns='model', values=metric_col)
        # Only run if all models have data for all prompt_id
        if pivot.notnull().all(axis=1).any() and pivot.shape[1] > 1:
            # Only keep rows (prompts) where all models present
            data = [pivot[m].dropna().values for m in pivot.columns]
            if all(len(x) == len(data[0]) for x in data):
                stat, p = friedmanchisquare(*data)
                print(f"{group_var.capitalize()} '{group_val}': Friedman's
                ↪ chi2={stat:.3f}, p={p:.4g}")
            else:
                print(f"{group_var.capitalize()} '{group_val}': Not all models
                ↪ have data for all prompts.")
        else:
            print(f"{group_var.capitalize()} '{group_val}': Not enough data for
            ↪ Friedman's test.")
#

```

```

=== Friedman's Test for levenshtein_similarity (by augmentation_type) ===
Augmentation_type 'entity_replace': Friedman's chi2=6.450, p=0.09166
Augmentation_type 'expand': Friedman's chi2=1.350, p=0.7173
Augmentation_type 'explain_simple': Friedman's chi2=2.550, p=0.4663
Augmentation_type 'identity': Friedman's chi2=19.950, p=0.0001738
Augmentation_type 'negation': Friedman's chi2=4.200, p=0.2407
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=1.050, p=0.7892
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=20.250, p=0.0001506
Augmentation_type 'summarize': Friedman's chi2=17.250, p=0.0006278
Augmentation_type 'synonym': Friedman's chi2=1.950, p=0.5828

```

```

=== Friedman's Test for jaccard_similarity (by augmentation_type) ===
Augmentation_type 'entity_replace': Friedman's chi2=12.300, p=0.006423
Augmentation_type 'expand': Friedman's chi2=4.443, p=0.2174
Augmentation_type 'explain_simple': Friedman's chi2=2.544, p=0.4673
Augmentation_type 'identity': Friedman's chi2=14.550, p=0.002245

```

```

Augmentation_type 'negation': Friedman's chi2=8.250, p=0.04112
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=14.700, p=0.002092
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=12.450, p=0.00599
Augmentation_type 'summarize': Friedman's chi2=4.950, p=0.1755
Augmentation_type 'synonym': Friedman's chi2=4.650, p=0.1993

```

=== Friedman's Test for length_ratio (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=8.089, p=0.04422
Augmentation_type 'expand': Friedman's chi2=20.550, p=0.0001305
Augmentation_type 'explain_simple': Friedman's chi2=14.850, p=0.001949
Augmentation_type 'identity': Friedman's chi2=4.950, p=0.1755
Augmentation_type 'negation': Friedman's chi2=5.430, p=0.1429
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=15.450, p=0.00147
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=3.300, p=0.3476
Augmentation_type 'summarize': Friedman's chi2=17.100, p=0.000674
Augmentation_type 'synonym': Friedman's chi2=5.550, p=0.1357

```

=== Friedman's Test for bleu (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=14.700, p=0.002092
Augmentation_type 'expand': Friedman's chi2=1.650, p=0.6481
Augmentation_type 'explain_simple': Friedman's chi2=4.050, p=0.2561
Augmentation_type 'identity': Friedman's chi2=15.000, p=0.001817
Augmentation_type 'negation': Friedman's chi2=12.450, p=0.00599
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=11.400, p=0.009748
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=9.150, p=0.02736
Augmentation_type 'summarize': Friedman's chi2=4.200, p=0.2407
Augmentation_type 'synonym': Friedman's chi2=4.050, p=0.2561

```

=== Friedman's Test for cosine_similarity (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=11.250, p=0.01045
Augmentation_type 'expand': Friedman's chi2=0.750, p=0.8614
Augmentation_type 'explain_simple': Friedman's chi2=1.800, p=0.6149
Augmentation_type 'identity': Friedman's chi2=15.000, p=0.001817
Augmentation_type 'negation': Friedman's chi2=6.600, p=0.0858
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=13.950, p=0.002974
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=13.050, p=0.00453
Augmentation_type 'summarize': Friedman's chi2=5.850, p=0.1191
Augmentation_type 'synonym': Friedman's chi2=3.450, p=0.3273

```

=== Friedman's Test for wer (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=8.423, p=0.03803
Augmentation_type 'expand': Friedman's chi2=17.416, p=0.0005804
Augmentation_type 'explain_simple': Friedman's chi2=13.720, p=0.003312
Augmentation_type 'identity': Friedman's chi2=21.911, p=6.806e-05
Augmentation_type 'negation': Friedman's chi2=12.797, p=0.005096
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=12.038, p=0.007253
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=12.600, p=0.005587
Augmentation_type 'summarize': Friedman's chi2=11.962, p=0.007514
Augmentation_type 'synonym': Friedman's chi2=16.269, p=0.0009986

```

=== Friedman's Test for char_diversity (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=13.709, p=0.003329
Augmentation_type 'expand': Friedman's chi2=14.316, p=0.002505
Augmentation_type 'explain_simple': Friedman's chi2=8.250, p=0.04112
Augmentation_type 'identity': Friedman's chi2=15.450, p=0.00147
Augmentation_type 'negation': Friedman's chi2=14.760, p=0.002034
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=14.400, p=0.002408
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=0.797, p=0.8501
Augmentation_type 'summarize': Friedman's chi2=0.375, p=0.9454
Augmentation_type 'synonym': Friedman's chi2=14.392, p=0.002417

```

=== Friedman's Test for type_token_ratio (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=19.105, p=0.00026
Augmentation_type 'expand': Friedman's chi2=16.350, p=0.0009612
Augmentation_type 'explain_simple': Friedman's chi2=14.924, p=0.001883
Augmentation_type 'identity': Friedman's chi2=15.759, p=0.00127
Augmentation_type 'negation': Friedman's chi2=16.063, p=0.001101
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=12.342, p=0.0063
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=18.974, p=0.0002768
Augmentation_type 'summarize': Friedman's chi2=4.378, p=0.2234
Augmentation_type 'synonym': Friedman's chi2=13.937, p=0.002993

```

=== Friedman's Test for bigram_overlap (by augmentation_type) ===

```

Augmentation_type 'entity_replace': Friedman's chi2=15.000, p=0.001817
Augmentation_type 'expand': Friedman's chi2=1.481, p=0.6867
Augmentation_type 'explain_simple': Friedman's chi2=5.388, p=0.1455
Augmentation_type 'identity': Friedman's chi2=14.550, p=0.002245
Augmentation_type 'negation': Friedman's chi2=8.924, p=0.03032
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=9.115, p=0.0278
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=11.960, p=0.007521

```

Augmentation_type 'summarize': Friedman's chi2=1.481, p=0.6867
Augmentation_type 'synonym': Friedman's chi2=4.038, p=0.2573

=== Friedman's Test for total_duration_ns (by augmentation_type) ===
Augmentation_type 'entity_replace': Friedman's chi2=22.950, p=4.136e-05
Augmentation_type 'expand': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'explain_simple': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'identity': Friedman's chi2=18.300, p=0.0003814
Augmentation_type 'negation': Friedman's chi2=22.950, p=4.136e-05
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=16.950, p=0.0007237
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=16.650, p=0.0008341
Augmentation_type 'summarize': Friedman's chi2=11.100, p=0.0112
Augmentation_type 'synonym': Friedman's chi2=19.950, p=0.0001738

=== Friedman's Test for load_duration_ns (by augmentation_type) ===
Augmentation_type 'entity_replace': Friedman's chi2=22.950, p=4.136e-05
Augmentation_type 'expand': Friedman's chi2=21.600, p=7.9e-05
Augmentation_type 'explain_simple': Friedman's chi2=22.200, p=5.927e-05
Augmentation_type 'identity': Friedman's chi2=22.950, p=4.136e-05
Augmentation_type 'negation': Friedman's chi2=21.600, p=7.9e-05
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=15.750, p=0.001276
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=22.200, p=5.927e-05
Augmentation_type 'summarize': Friedman's chi2=22.200, p=5.927e-05
Augmentation_type 'synonym': Friedman's chi2=21.750, p=7.353e-05

=== Friedman's Test for prompt_eval_count (by augmentation_type) ===
Augmentation_type 'entity_replace': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'expand': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'explain_simple': Friedman's chi2=23.734, p=2.838e-05
Augmentation_type 'identity': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'negation': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'summarize': Friedman's chi2=22.792, p=4.462e-05
Augmentation_type 'synonym': Friedman's chi2=22.792, p=4.462e-05

=== Friedman's Test for prompt_eval_duration_ns (by augmentation_type) ===
Augmentation_type 'entity_replace': Friedman's chi2=21.750, p=7.353e-05
Augmentation_type 'expand': Friedman's chi2=21.600, p=7.9e-05
Augmentation_type 'explain_simple': Friedman's chi2=21.750, p=7.353e-05
Augmentation_type 'identity': Friedman's chi2=21.600, p=7.9e-05
Augmentation_type 'negation': Friedman's chi2=21.600, p=7.9e-05

```
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=22.200, p=5.927e-05
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=21.750, p=7.353e-05
Augmentation_type 'summarize': Friedman's chi2=21.750, p=7.353e-05
Augmentation_type 'synonym': Friedman's chi2=21.750, p=7.353e-05
```

```
=== Friedman's Test for eval_count (by augmentation_type) ===
```

```
Augmentation_type 'entity_replace': Friedman's chi2=9.279, p=0.0258
Augmentation_type 'expand': Friedman's chi2=nan, p=nan
Augmentation_type 'explain_simple': Friedman's chi2=nan, p=nan
Augmentation_type 'identity': Friedman's chi2=5.211, p=0.157
Augmentation_type 'negation': Friedman's chi2=13.476, p=0.003712
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=19.800, p=0.0001867
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=7.163, p=0.06687
Augmentation_type 'summarize': Friedman's chi2=18.917, p=0.0002845
Augmentation_type 'synonym': Friedman's chi2=14.053, p=0.002834
```

```
=== Friedman's Test for eval_duration_ns (by augmentation_type) ===
```

```
Augmentation_type 'entity_replace': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'expand': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'explain_simple': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'identity': Friedman's chi2=15.600, p=0.001369
Augmentation_type 'negation': Friedman's chi2=22.950, p=4.136e-05
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=18.450, p=0.0003552
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=13.650, p=0.003422
Augmentation_type 'summarize': Friedman's chi2=8.700, p=0.03356
Augmentation_type 'synonym': Friedman's chi2=19.950, p=0.0001738
```

```
=== Friedman's Test for tokens_per_second (by augmentation_type) ===
```

```
Augmentation_type 'entity_replace': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'expand': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'explain_simple': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'identity': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'negation': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'noise': Not all models have data for all prompts.
Augmentation_type 'paraphrase': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'question_gen': Not all models have data for all prompts.
Augmentation_type 'shuffle': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'summarize': Friedman's chi2=24.000, p=2.498e-05
Augmentation_type 'synonym': Friedman's chi2=24.000, p=2.498e-05
```

```
C:\Users\parth\anaconda3\Lib\site-packages\scipy\stats\_stats_py.py:9427:
RuntimeWarning: invalid value encountered in scalar divide
```

```

chisq = (12.0 / (k*n*(k+1)) * ssbn - 3*n*(k+1)) / c
C:\Users\parth\anaconda3\Lib\site-packages\scipy\stats\_stats_py.py:9427:
RuntimeWarning: invalid value encountered in scalar divide
chisq = (12.0 / (k*n*(k+1)) * ssbn - 3*n*(k+1)) / c

```

```

[53]: # Kolmogorov-Smirnov (K-S) Test (for all pairs of models)

from scipy.stats import ks_2samp
from itertools import combinations

for metric_col in metrics:
    print(f"\n=== Kolmogorov-Smirnov Test for {metric_col} (All Model Pairs)␣
    ␣====")
    models = df['model'].unique()
    for m1, m2 in combinations(models, 2):
        data1 = df[df['model'] == m1][metric_col].dropna()
        data2 = df[df['model'] == m2][metric_col].dropna()
        if len(data1) > 0 and len(data2) > 0:
            ks_stat, ks_p = ks_2samp(data1, data2)
            print(f" {m1} vs {m2}: KS stat={ks_stat:.3f}, p={ks_p:.4g}")
        else:
            print(f" {m1} vs {m2}: Not enough data.")

```

```

=== Kolmogorov-Smirnov Test for levenshtein_similarity (All Model Pairs) ===

```

```

qwen2:0.5b vs smollm2:360m: KS stat=0.239, p=0.01307
qwen2:0.5b vs gemma3:1b: KS stat=0.364, p=1.447e-05
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.175, p=0.1334
smollm2:360m vs gemma3:1b: KS stat=0.273, p=0.002746
smollm2:360m vs granite3.1-moe:1b: KS stat=0.328, p=0.0001615
gemma3:1b vs granite3.1-moe:1b: KS stat=0.430, p=1.772e-07

```

```

=== Kolmogorov-Smirnov Test for jaccard_similarity (All Model Pairs) ===

```

```

qwen2:0.5b vs smollm2:360m: KS stat=0.125, p=0.4999
qwen2:0.5b vs gemma3:1b: KS stat=0.386, p=3.018e-06
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.261, p=0.005102
smollm2:360m vs gemma3:1b: KS stat=0.432, p=9.537e-08
smollm2:360m vs granite3.1-moe:1b: KS stat=0.350, p=4.452e-05
gemma3:1b vs granite3.1-moe:1b: KS stat=0.309, p=0.0004753

```

```

=== Kolmogorov-Smirnov Test for length_ratio (All Model Pairs) ===

```

```

qwen2:0.5b vs smollm2:360m: KS stat=0.352, p=3.046e-05
qwen2:0.5b vs gemma3:1b: KS stat=0.352, p=3.046e-05
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.552, p=2.635e-12
smollm2:360m vs gemma3:1b: KS stat=0.250, p=0.007959
smollm2:360m vs granite3.1-moe:1b: KS stat=0.234, p=0.01648
gemma3:1b vs granite3.1-moe:1b: KS stat=0.263, p=0.004845

```

```

=== Kolmogorov-Smirnov Test for bleu (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.114, p=0.6237
qwen2:0.5b vs gemma3:1b: KS stat=0.409, p=5.664e-07
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.266, p=0.004142
smollm2:360m vs gemma3:1b: KS stat=0.432, p=9.537e-08
smollm2:360m vs granite3.1-moe:1b: KS stat=0.250, p=0.00845
gemma3:1b vs granite3.1-moe:1b: KS stat=0.419, p=3.87e-07

=== Kolmogorov-Smirnov Test for cosine_similarity (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.159, p=0.2161
qwen2:0.5b vs gemma3:1b: KS stat=0.398, p=1.325e-06
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.188, p=0.09014
smollm2:360m vs gemma3:1b: KS stat=0.489, p=6.79e-10
smollm2:360m vs granite3.1-moe:1b: KS stat=0.289, p=0.001383
gemma3:1b vs granite3.1-moe:1b: KS stat=0.307, p=0.0005372

=== Kolmogorov-Smirnov Test for wer (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.284, p=0.001555
qwen2:0.5b vs gemma3:1b: KS stat=0.330, p=0.000125
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.474, p=4.679e-09
smollm2:360m vs gemma3:1b: KS stat=0.261, p=0.004732
smollm2:360m vs granite3.1-moe:1b: KS stat=0.266, p=0.004142
gemma3:1b vs granite3.1-moe:1b: KS stat=0.317, p=0.0003074

=== Kolmogorov-Smirnov Test for char_diversity (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.216, p=0.03278
qwen2:0.5b vs gemma3:1b: KS stat=0.523, p=2.457e-11
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.165, p=0.1801
smollm2:360m vs gemma3:1b: KS stat=0.602, p=3.508e-15
smollm2:360m vs granite3.1-moe:1b: KS stat=0.170, p=0.1528
gemma3:1b vs granite3.1-moe:1b: KS stat=0.562, p=8.939e-13

=== Kolmogorov-Smirnov Test for type_token_ratio (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.125, p=0.4999
qwen2:0.5b vs gemma3:1b: KS stat=0.716, p=5.142e-22
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.309, p=0.0004753
smollm2:360m vs gemma3:1b: KS stat=0.773, p=3.786e-26
smollm2:360m vs granite3.1-moe:1b: KS stat=0.230, p=0.01977
gemma3:1b vs granite3.1-moe:1b: KS stat=0.815, p=1.146e-28

=== Kolmogorov-Smirnov Test for bigram_overlap (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.068, p=0.9876
qwen2:0.5b vs gemma3:1b: KS stat=0.409, p=5.664e-07
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.211, p=0.03958
smollm2:360m vs gemma3:1b: KS stat=0.432, p=9.537e-08
smollm2:360m vs granite3.1-moe:1b: KS stat=0.223, p=0.02583
gemma3:1b vs granite3.1-moe:1b: KS stat=0.339, p=8.87e-05

```

```

=== Kolmogorov-Smirnov Test for total_duration_ns (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.693, p=1.659e-20
qwen2:0.5b vs gemma3:1b: KS stat=0.977, p=5.354e-48
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.926, p=2.924e-39
smollm2:360m vs gemma3:1b: KS stat=0.795, p=5.799e-28
smollm2:360m vs granite3.1-moe:1b: KS stat=0.533, p=1.9e-11
gemma3:1b vs granite3.1-moe:1b: KS stat=0.897, p=1.893e-36

=== Kolmogorov-Smirnov Test for load_duration_ns (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.989, p=6.119e-50
qwen2:0.5b vs gemma3:1b: KS stat=0.989, p=6.119e-50
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.988, p=1.766e-47
smollm2:360m vs gemma3:1b: KS stat=0.989, p=6.119e-50
smollm2:360m vs granite3.1-moe:1b: KS stat=0.116, p=0.578
gemma3:1b vs granite3.1-moe:1b: KS stat=0.988, p=1.766e-47

=== Kolmogorov-Smirnov Test for prompt_eval_count (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=1.000, p=3.477e-52
qwen2:0.5b vs gemma3:1b: KS stat=0.068, p=0.9876
qwen2:0.5b vs granite3.1-moe:1b: KS stat=1.000, p=1.051e-49
smollm2:360m vs gemma3:1b: KS stat=0.989, p=6.119e-50
smollm2:360m vs granite3.1-moe:1b: KS stat=1.000, p=1.051e-49
gemma3:1b vs granite3.1-moe:1b: KS stat=1.000, p=1.051e-49

=== Kolmogorov-Smirnov Test for prompt_eval_duration_ns (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.250, p=0.007959
qwen2:0.5b vs gemma3:1b: KS stat=1.000, p=3.477e-52
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.875, p=3.991e-34
smollm2:360m vs gemma3:1b: KS stat=1.000, p=3.477e-52
smollm2:360m vs granite3.1-moe:1b: KS stat=0.898, p=1.456e-36
gemma3:1b vs granite3.1-moe:1b: KS stat=0.988, p=1.766e-47

=== Kolmogorov-Smirnov Test for eval_count (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.341, p=6.249e-05
qwen2:0.5b vs gemma3:1b: KS stat=0.136, p=0.3883
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.244, p=0.01078
smollm2:360m vs gemma3:1b: KS stat=0.455, p=1.437e-08
smollm2:360m vs granite3.1-moe:1b: KS stat=0.535, p=1.513e-11
gemma3:1b vs granite3.1-moe:1b: KS stat=0.155, p=0.2388

=== Kolmogorov-Smirnov Test for eval_duration_ns (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=0.739, p=1.327e-23
qwen2:0.5b vs gemma3:1b: KS stat=0.920, p=3.198e-40
qwen2:0.5b vs granite3.1-moe:1b: KS stat=0.938, p=1.104e-40
smollm2:360m vs gemma3:1b: KS stat=0.784, p=4.822e-27
smollm2:360m vs granite3.1-moe:1b: KS stat=0.625, p=6.761e-16
gemma3:1b vs granite3.1-moe:1b: KS stat=0.898, p=1.456e-36

```



```

=== Kolmogorov-Smirnov Test for tokens_per_second (All Model Pairs) ===
qwen2:0.5b vs smollm2:360m: KS stat=1.000, p=3.477e-52
qwen2:0.5b vs gemma3:1b: KS stat=1.000, p=3.477e-52
qwen2:0.5b vs granite3.1-moe:1b: KS stat=1.000, p=1.051e-49
smollm2:360m vs gemma3:1b: KS stat=1.000, p=3.477e-52
smollm2:360m vs granite3.1-moe:1b: KS stat=1.000, p=1.051e-49
gemma3:1b vs granite3.1-moe:1b: KS stat=1.000, p=1.051e-49

```

```

[54]: # Spearman's Rank Correlation for All Metrics

# How to interpret:

# Correlation coefficient in [-1, 1]:
# Closer to 1 = strong positive monotonic relation
# Closer to -1 = strong negative
# Close to 0 = no monotonic association

# P-value:
# p < 0.05 means significant correlation
# p > 0.05 means correlation is not statistically significant

import pandas as pd
from scipy.stats import spearmanr

# List of your numeric metrics
metrics = [
    'levenshtein_similarity', 'jaccard_similarity', 'length_ratio', 'bleu',
    'cosine_similarity', 'wer', 'char_diversity', 'type_token_ratio',
    'bigram_overlap',
    'total_duration_ns', 'load_duration_ns', 'prompt_eval_count',
    'prompt_eval_duration_ns',
    'eval_count', 'eval_duration_ns', 'tokens_per_second'
]

# Subset dataframe to just these columns, drop rows with any NaN
metrics_df = df[metrics].dropna()

# Compute Spearman correlation (returns correlation matrix and p-value matrix)
corr, pval = spearmanr(metrics_df, axis=0)

# To DataFrame for easier reading:
corr_df = pd.DataFrame(corr, index=metrics, columns=metrics)
pval_df = pd.DataFrame(pval, index=metrics, columns=metrics)

# Print or display the correlation matrix
print("Spearman Correlation Matrix:")
print(corr_df.round(3))

```

```
print("\nSpearman Correlation P-value Matrix:")
print(pval_df.round(3))
```

Spearman Correlation Matrix:

	levenshtein_similarity	jaccard_similarity \
levenshtein_similarity	1.000	0.329
jaccard_similarity	0.329	1.000
length_ratio	-0.240	0.301
bleu	0.403	0.921
cosine_similarity	0.325	0.843
wer	-0.538	-0.454
char_diversity	-0.226	-0.262
type_token_ratio	-0.060	-0.352
bigram_overlap	0.447	0.911
total_duration_ns	-0.079	0.340
load_duration_ns	0.131	0.198
prompt_eval_count	-0.183	-0.053
prompt_eval_duration_ns	0.041	0.369
eval_count	-0.440	0.245
eval_duration_ns	-0.115	0.305
tokens_per_second	-0.172	-0.227

	length_ratio	bleu	cosine_similarity	wer \
levenshtein_similarity	-0.240	0.403	0.325	-0.538
jaccard_similarity	0.301	0.921	0.843	-0.454
length_ratio	1.000	0.253	0.279	0.494
bleu	0.253	1.000	0.816	-0.448
cosine_similarity	0.279	0.816	1.000	-0.399
wer	0.494	-0.448	-0.399	1.000
char_diversity	0.168	-0.173	-0.087	0.456
type_token_ratio	-0.270	-0.379	-0.507	-0.069
bigram_overlap	0.218	0.972	0.794	-0.462
total_duration_ns	0.218	0.365	0.368	0.031
load_duration_ns	0.127	0.253	0.273	0.082
prompt_eval_count	-0.277	-0.101	-0.137	-0.257
prompt_eval_duration_ns	-0.142	0.373	0.360	-0.238
eval_count	0.625	0.203	0.238	0.413
eval_duration_ns	0.298	0.335	0.341	0.112
tokens_per_second	0.023	-0.273	-0.263	0.072

	char_diversity	type_token_ratio	bigram_overlap \
levenshtein_similarity	-0.226	-0.060	0.447
jaccard_similarity	-0.262	-0.352	0.911
length_ratio	0.168	-0.270	0.218
bleu	-0.173	-0.379	0.972
cosine_similarity	-0.087	-0.507	0.794
wer	0.456	-0.069	-0.462

char_diversity	1.000	-0.238	-0.213
type_token_ratio	-0.238	1.000	-0.342
bigram_overlap	-0.213	-0.342	1.000
total_duration_ns	0.324	-0.561	0.302
load_duration_ns	0.307	-0.512	0.247
prompt_eval_count	-0.242	0.394	-0.111
prompt_eval_duration_ns	0.223	-0.445	0.345
eval_count	0.397	-0.294	0.145
eval_duration_ns	0.372	-0.568	0.268
tokens_per_second	-0.266	0.511	-0.246

	total_duration_ns	load_duration_ns \
levenshtein_similarity	-0.079	0.131
jaccard_similarity	0.340	0.198
length_ratio	0.218	0.127
bleu	0.365	0.253
cosine_similarity	0.368	0.273
wer	0.031	0.082
char_diversity	0.324	0.307
type_token_ratio	-0.561	-0.512
bigram_overlap	0.302	0.247
total_duration_ns	1.000	0.285
load_duration_ns	0.285	1.000
prompt_eval_count	0.078	-0.728
prompt_eval_duration_ns	0.703	0.457
eval_count	0.505	0.117
eval_duration_ns	0.971	0.239
tokens_per_second	-0.841	-0.292

	prompt_eval_count	prompt_eval_duration_ns \
levenshtein_similarity	-0.183	0.041
jaccard_similarity	-0.053	0.369
length_ratio	-0.277	-0.142
bleu	-0.101	0.373
cosine_similarity	-0.137	0.360
wer	-0.257	-0.238
char_diversity	-0.242	0.223
type_token_ratio	0.394	-0.445
bigram_overlap	-0.111	0.345
total_duration_ns	0.078	0.703
load_duration_ns	-0.728	0.457
prompt_eval_count	1.000	0.061
prompt_eval_duration_ns	0.061	1.000
eval_count	0.091	0.314
eval_duration_ns	0.052	0.604
tokens_per_second	0.043	-0.608

eval_count eval_duration_ns tokens_per_second

levenshtein_similarity	-0.440	-0.115	-0.172
jaccard_similarity	0.245	0.305	-0.227
length_ratio	0.625	0.298	0.023
bleu	0.203	0.335	-0.273
cosine_similarity	0.238	0.341	-0.263
wer	0.413	0.112	0.072
char_diversity	0.397	0.372	-0.266
type_token_ratio	-0.294	-0.568	0.511
bigram_overlap	0.145	0.268	-0.246
total_duration_ns	0.505	0.971	-0.841
load_duration_ns	0.117	0.239	-0.292
prompt_eval_count	0.091	0.052	0.043
prompt_eval_duration_ns	0.314	0.604	-0.608
eval_count	1.000	0.548	-0.119
eval_duration_ns	0.548	1.000	-0.830
tokens_per_second	-0.119	-0.830	1.000

Spearman Correlation P-value Matrix:

	levenshtein_similarity	jaccard_similarity \
levenshtein_similarity	0.000	0.000
jaccard_similarity	0.000	0.000
length_ratio	0.000	0.000
bleu	0.000	0.000
cosine_similarity	0.000	0.000
wer	0.000	0.000
char_diversity	0.000	0.000
type_token_ratio	0.263	0.000
bigram_overlap	0.000	0.000
total_duration_ns	0.145	0.000
load_duration_ns	0.015	0.000
prompt_eval_count	0.001	0.326
prompt_eval_duration_ns	0.450	0.000
eval_count	0.000	0.000
eval_duration_ns	0.033	0.000
tokens_per_second	0.001	0.000

	length_ratio	bleu	cosine_similarity	wer \
levenshtein_similarity	0.000	0.000	0.000	0.000
jaccard_similarity	0.000	0.000	0.000	0.000
length_ratio	0.000	0.000	0.000	0.000
bleu	0.000	0.000	0.000	0.000
cosine_similarity	0.000	0.000	0.000	0.000
wer	0.000	0.000	0.000	0.000
char_diversity	0.002	0.001	0.106	0.000
type_token_ratio	0.000	0.000	0.000	0.201
bigram_overlap	0.000	0.000	0.000	0.000
total_duration_ns	0.000	0.000	0.000	0.565
load_duration_ns	0.018	0.000	0.000	0.128

prompt_eval_count	0.000	0.061	0.011	0.000
prompt_eval_duration_ns	0.008	0.000	0.000	0.000
eval_count	0.000	0.000	0.000	0.000
eval_duration_ns	0.000	0.000	0.000	0.039
tokens_per_second	0.674	0.000	0.000	0.184

	char_diversity	type_token_ratio	bigram_overlap	\
levenshtein_similarity	0.000	0.263	0.000	
jaccard_similarity	0.000	0.000	0.000	
length_ratio	0.002	0.000	0.000	
bleu	0.001	0.000	0.000	
cosine_similarity	0.106	0.000	0.000	
wer	0.000	0.201	0.000	
char_diversity	0.000	0.000	0.000	
type_token_ratio	0.000	0.000	0.000	
bigram_overlap	0.000	0.000	0.000	
total_duration_ns	0.000	0.000	0.000	
load_duration_ns	0.000	0.000	0.000	
prompt_eval_count	0.000	0.000	0.040	
prompt_eval_duration_ns	0.000	0.000	0.000	
eval_count	0.000	0.000	0.007	
eval_duration_ns	0.000	0.000	0.000	
tokens_per_second	0.000	0.000	0.000	

	total_duration_ns	load_duration_ns	\
levenshtein_similarity	0.145	0.015	
jaccard_similarity	0.000	0.000	
length_ratio	0.000	0.018	
bleu	0.000	0.000	
cosine_similarity	0.000	0.000	
wer	0.565	0.128	
char_diversity	0.000	0.000	
type_token_ratio	0.000	0.000	
bigram_overlap	0.000	0.000	
total_duration_ns	0.000	0.000	
load_duration_ns	0.000	0.000	
prompt_eval_count	0.150	0.000	
prompt_eval_duration_ns	0.000	0.000	
eval_count	0.000	0.031	
eval_duration_ns	0.000	0.000	
tokens_per_second	0.000	0.000	

	prompt_eval_count	prompt_eval_duration_ns	\
levenshtein_similarity	0.001	0.450	
jaccard_similarity	0.326	0.000	
length_ratio	0.000	0.008	
bleu	0.061	0.000	
cosine_similarity	0.011	0.000	

wer	0.000	0.000
char_diversity	0.000	0.000
type_token_ratio	0.000	0.000
bigram_overlap	0.040	0.000
total_duration_ns	0.150	0.000
load_duration_ns	0.000	0.000
prompt_eval_count	0.000	0.261
prompt_eval_duration_ns	0.261	0.000
eval_count	0.091	0.000
eval_duration_ns	0.338	0.000
tokens_per_second	0.425	0.000

	eval_count	eval_duration_ns	tokens_per_second
levenshtein_similarity	0.000	0.033	0.001
jaccard_similarity	0.000	0.000	0.000
length_ratio	0.000	0.000	0.674
bleu	0.000	0.000	0.000
cosine_similarity	0.000	0.000	0.000
wer	0.000	0.039	0.184
char_diversity	0.000	0.000	0.000
type_token_ratio	0.000	0.000	0.000
bigram_overlap	0.007	0.000	0.000
total_duration_ns	0.000	0.000	0.000
load_duration_ns	0.031	0.000	0.000
prompt_eval_count	0.091	0.338	0.425
prompt_eval_duration_ns	0.000	0.000	0.000
eval_count	0.000	0.000	0.027
eval_duration_ns	0.000	0.000	0.000
tokens_per_second	0.027	0.000	0.000

```
[55]: import pandas as pd

metrics = [
    'levenshtein_similarity', 'jaccard_similarity', 'length_ratio', 'bleu',
    'cosine_similarity', 'wer', 'char_diversity', 'type_token_ratio',
    ↪ 'bigram_overlap',
    'total_duration_ns', 'load_duration_ns', 'prompt_eval_count',
    ↪ 'prompt_eval_duration_ns',
    'eval_count', 'eval_duration_ns', 'tokens_per_second'
]

# Filter to relevant columns, drop rows with any missing value in those columns
df_manova = df[['model', 'augmentation_type'] + metrics].dropna()
print(df_manova.shape)

## MANOVA by Model Only
```

```

# Interpretation
# Look at the line for model in Wilks' lambda, Pillai's trace, Hotelling-Lawley
  ↳ trace, and Roy's largest root.
# The Pr > F column gives the p-value for the global null: "All LLMs have the
  ↳ same mean vector of metrics."
# p < 0.05 means at least one LLM is different on at least one combination of
  ↳ metrics.

from statsmodels.multivariate.manova import MANOVA

formula = ' + '.join(metrics) + ' ~ model'
maov = MANOVA.from_formula(formula, data=df_manova)
print(maov.mv_test())

```

(344, 18)

Multivariate linear model

=====						

Intercept	Value	Num DF	Den DF	F Value	Pr > F	

Wilks' lambda	0.0002	4.0000	337.0000	397563.2644	0.0000	
Pillai's trace	0.9998	4.0000	337.0000	397563.2644	0.0000	
Hotelling-Lawley trace	4718.8518	4.0000	337.0000	397563.2644	0.0000	
Roy's greatest root	4718.8518	4.0000	337.0000	397563.2644	0.0000	

model	Value	Num DF	Den DF	F Value	Pr > F	

Wilks' lambda	0.0000	12.0000	891.9097	14012.8080	0.0000	
Pillai's trace	2.7918	12.0000	1017.0000	1136.1569	0.0000	
Hotelling-Lawley trace	4114.5409	12.0000	585.4521	115258.0057	0.0000	
Roy's greatest root	4062.1679	4.0000	339.0000	344268.7302	0.0000	
=====						

[56]: #MANOVA by Model and Augmentation Type (Factorial MANOVA)

```

formula = ' + '.join(metrics) + ' ~ model + augmentation_type'
maov2 = MANOVA.from_formula(formula, data=df_manova)
print(maov2.mv_test())

```

Multivariate linear model

Intercept	Value	Num DF	Den DF	F Value	Pr > F
Wilks' lambda	0.0007	4.0000	327.0000	118474.4430	0.0000
Pillai's trace	0.9993	4.0000	327.0000	118474.4430	0.0000
Hotelling-Lawley trace	1449.2287	4.0000	327.0000	118474.4430	0.0000
Roy's greatest root	1449.2287	4.0000	327.0000	118474.4430	0.0000

model	Value	Num DF	Den DF	F Value	Pr > F
Wilks' lambda	0.0000	12.0000	865.4522	18710.9331	0.0000
Pillai's trace	2.8161	12.0000	987.0000	1259.7218	0.0000
Hotelling-Lawley trace	4490.8920	12.0000	567.9526	122058.1435	0.0000
Roy's greatest root	4389.9069	4.0000	329.0000	361069.8449	0.0000

augmentation_type	Value	Num DF	Den DF	F Value	Pr > F
Wilks' lambda	0.0624	40.0000	1241.8012	33.4889	0.0000
Pillai's trace	2.1361	40.0000	1320.0000	37.8209	0.0000
Hotelling-Lawley trace	3.8299	40.0000	903.4338	31.1871	0.0000
Roy's greatest root	1.6174	10.0000	330.0000	53.3743	0.0000

[57]: *# MANOVA With Interaction Term*

```
formula = ' + '.join(metrics) + ' ~ model * augmentation_type'
maov3 = MANOVA.from_formula(formula, data=df_manova)
print(maov3.mv_test())
```

Multivariate linear model

Intercept	Value	Num DF	Den DF	F Value	Pr > F
Wilks' lambda	0.0017	4.0000	297.0000	42399.6329	0.0000
Pillai's trace	0.9983	4.0000	297.0000	42399.6329	0.0000
Hotelling-Lawley trace	571.0388	4.0000	297.0000	42399.6329	0.0000
Roy's greatest root	571.0388	4.0000	297.0000	42399.6329	0.0000

model	Value	Num DF	Den DF	F Value	Pr > F
-------	-------	--------	--------	---------	--------

Wilks' lambda	0.0001	12.0000	786.0796	1896.7015	0.0000
Pillai's trace	2.2510	12.0000	897.0000	224.6345	0.0000
Hotelling-Lawley trace	475.3388	12.0000	515.4541	11730.9296	0.0000
Roy's greatest root	464.4037	4.0000	299.0000	34714.1753	0.0000

augmentation_type	Value	Num DF	Den DF	F Value	Pr > F
Wilks' lambda	0.0221	40.0000	1128.0447	48.8563	0.0000
Pillai's trace	2.4002	40.0000	1200.0000	45.0087	0.0000
Hotelling-Lawley trace	7.5527	40.0000	819.4412	55.8374	0.0000
Roy's greatest root	4.4835	10.0000	300.0000	134.5065	0.0000

model:augmentation_type	Value	Num DF	Den DF	F Value	Pr > F
Wilks' lambda	0.0071	120.0000	1183.2579	24.4081	0.0000
Pillai's trace	3.4426	120.0000	1200.0000	61.7649	0.0000
Hotelling-Lawley trace	8.2202	120.0000	1016.8090	20.2479	0.0000
Roy's greatest root	3.2635	30.0000	300.0000	32.6351	0.0000

[58]: *#Post-hoc Exploration: Which Metrics Are Responsible?*

#If MANOVA is significant, run univariate ANOVA for each metric:

```
import statsmodels.api as sm
from statsmodels.formula.api import ols

for metric in metrics:
    model = ols(f'{metric} ~ model', data=df_manova).fit()
    anova_table = sm.stats.anova_lm(model, typ=2)
    print(f"ANOVA for {metric}:\n", anova_table, "\n")
```

ANOVA for levenshtein_similarity:

	sum_sq	df	F	PR(>F)
model	3.910936	3.0	26.939258	1.177641e-15
Residual	16.453291	340.0	NaN	NaN

ANOVA for jaccard_similarity:

	sum_sq	df	F	PR(>F)
model	2.841742	3.0	27.623562	5.199714e-16
Residual	11.659036	340.0	NaN	NaN

ANOVA for length_ratio:

	sum_sq	df	F	PR(>F)
model	1.649220	3.0	7.683006	0.000056
Residual	24.327932	340.0	NaN	NaN

ANOVA for bleu:

	sum_sq	df	F	PR(>F)
model	3.799517	3.0	32.509959	1.687200e-18
Residual	13.245540	340.0	NaN	NaN

ANOVA for cosine_similarity:

	sum_sq	df	F	PR(>F)
model	3.033348	3.0	26.071969	3.336757e-15
Residual	13.185790	340.0	NaN	NaN

ANOVA for wer:

	sum_sq	df	F	PR(>F)
model	4.486417	3.0	15.780305	1.249106e-09
Residual	32.221214	340.0	NaN	NaN

ANOVA for char_diversity:

	sum_sq	df	F	PR(>F)
model	0.773419	3.0	25.9462	3.882705e-15
Residual	3.378306	340.0	NaN	NaN

ANOVA for type_token_ratio:

	sum_sq	df	F	PR(>F)
model	0.456416	3.0	107.263221	6.956917e-49
Residual	0.482245	340.0	NaN	NaN

ANOVA for bigram_overlap:

	sum_sq	df	F	PR(>F)
model	3.389387	3.0	30.396176	1.967544e-17
Residual	12.637461	340.0	NaN	NaN

ANOVA for total_duration_ns:

	sum_sq	df	F	PR(>F)
model	3.833423e+21	3.0	114.62624	2.664217e-51
Residual	3.790185e+21	340.0	NaN	NaN

ANOVA for load_duration_ns:

	sum_sq	df	F	PR(>F)
model	6.248601e+18	3.0	0.399618	0.753364
Residual	1.772131e+21	340.0	NaN	NaN

ANOVA for prompt_eval_count:

	sum_sq	df	F	PR(>F)
model	145085.989429	3.0	1572.665804	6.740158e-199

Residual	10455.545455	340.0		NaN	NaN
----------	--------------	-------	--	-----	-----

ANOVA for prompt_eval_duration_ns:

	sum_sq	df	F	PR(>F)
model	1.914607e+20	3.0	825.644883	1.072120e-155
Residual	2.628112e+19	340.0	NaN	NaN

ANOVA for eval_count:

	sum_sq	df	F	PR(>F)
model	7829.789852	3.0	19.476164	1.124172e-11
Residual	45562.163636	340.0	NaN	NaN

ANOVA for eval_duration_ns:

	sum_sq	df	F	PR(>F)
model	2.324837e+21	3.0	158.921245	2.226701e-64
Residual	1.657938e+21	340.0	NaN	NaN

ANOVA for tokens_per_second:

	sum_sq	df	F	PR(>F)
model	1418.920496	3.0	97736.33144	0.0
Residual	1.645355	340.0	NaN	NaN

```
[59]: # Interpreting Model-wise Effect Sizes
#1. Mean Difference (mean_diff):
#This is simply the difference in mean value of a metric between two models.

#Positive value: model_1 has a higher mean than model_2 for that metric.

#Negative value: model_2 has a higher mean.

#2. Cohen's d (cohens_d):
#Cohen's d quantifies the magnitude of the difference (in standard deviations)
↪ between two distributions.

#Conventional thresholds for interpretation:

# |d| < 0.2: Negligible difference (practically equivalent)

# 0.2 |d| < 0.5: Small effect (slight difference)

# 0.5 |d| < 0.8: Medium effect (moderate difference)

# |d| 0.8: Large effect (substantial/practically meaningful difference)

# Sign of d:
```

```

# Positive: model_1 performs better (higher mean for positively oriented
↳ metrics, e.g., similarity, tokens/sec).

# Negative: model_2 performs better.

# 3. Metric Orientation
# For similarity/accuracy metrics (bleu, levenshtein_similarity,
↳ cosine_similarity, etc.): Higher is better.

# For error or duration metrics (wer, total_duration_ns, etc.): Lower is better.

# If Cohen's d is positive and the metric is an error/duration, model_1 has
↳ worse performance.

```

```

[60]: import pandas as pd
import numpy as np
from itertools import combinations
from scipy.stats import ttest_ind

def cohens_d(x, y):
    nx, ny = len(x), len(y)
    pooled_std = np.sqrt(((nx-1)*np.var(x, ddof=1) + (ny-1)*np.var(y, ddof=1)) /
↳ (nx + ny - 2))
    return (np.mean(x) - np.mean(y)) / pooled_std

metrics = [
    "total_duration_ns", "load_duration_ns", "prompt_eval_count",
↳ "prompt_eval_duration_ns", "eval_count",
    "eval_duration_ns", "tokens_per_second", "levenshtein_similarity",
↳ "jaccard_similarity", "length_ratio",
    "bleu", "cosine_similarity", "wer", "char_diversity", "type_token_ratio",
↳ "bigram_overlap"
]

model_list = sorted(df['model'].unique())
results = []

for metric in metrics:
    for m1, m2 in combinations(model_list, 2):
        x = df[df['model'] == m1][metric].dropna()
        y = df[df['model'] == m2][metric].dropna()
        if len(x) > 1 and len(y) > 1:
            d = cohens_d(x, y)
            mean_diff = x.mean() - y.mean()
            t_stat, p_val = ttest_ind(x, y, equal_var=False)
            results.append({

```

```

        'metric': metric,
        'model_1': m1,
        'model_2': m2,
        'mean_model_1': x.mean(),
        'mean_model_2': y.mean(),
        'std_model_1': x.std(ddof=1),
        'std_model_2': y.std(ddof=1),
        'n_model_1': len(x),
        'n_model_2': len(y),
        'mean_diff': mean_diff,
        'cohens_d': d,
        'ttest_pval': p_val
    })

effectsize_df = pd.DataFrame(results)

pd.set_option('display.max_rows', None)
print("Table: Pairwise Model Comparisons Across All Metrics-Means, Standard_
Deviations, Sample Sizes, Mean Differences, Cohen's d, and t-Test p-values")
display(effectsize_df)

```

Table: Pairwise Model Comparisons Across All Metrics-Means, Standard Deviations, Sample Sizes, Mean Differences, Cohen's d, and t-Test p-values

	metric	model_1	model_2 \
0	total_duration_ns	gemma3:1b	granite3.1-moe:1b
1	total_duration_ns	gemma3:1b	qwen2:0.5b
2	total_duration_ns	gemma3:1b	smollm2:360m
3	total_duration_ns	granite3.1-moe:1b	qwen2:0.5b
4	total_duration_ns	granite3.1-moe:1b	smollm2:360m
5	total_duration_ns	qwen2:0.5b	smollm2:360m
6	load_duration_ns	gemma3:1b	granite3.1-moe:1b
7	load_duration_ns	gemma3:1b	qwen2:0.5b
8	load_duration_ns	gemma3:1b	smollm2:360m
9	load_duration_ns	granite3.1-moe:1b	qwen2:0.5b
10	load_duration_ns	granite3.1-moe:1b	smollm2:360m
11	load_duration_ns	qwen2:0.5b	smollm2:360m
12	prompt_eval_count	gemma3:1b	granite3.1-moe:1b
13	prompt_eval_count	gemma3:1b	qwen2:0.5b
14	prompt_eval_count	gemma3:1b	smollm2:360m
15	prompt_eval_count	granite3.1-moe:1b	qwen2:0.5b
16	prompt_eval_count	granite3.1-moe:1b	smollm2:360m
17	prompt_eval_count	qwen2:0.5b	smollm2:360m
18	prompt_eval_duration_ns	gemma3:1b	granite3.1-moe:1b
19	prompt_eval_duration_ns	gemma3:1b	qwen2:0.5b
20	prompt_eval_duration_ns	gemma3:1b	smollm2:360m
21	prompt_eval_duration_ns	granite3.1-moe:1b	qwen2:0.5b
22	prompt_eval_duration_ns	granite3.1-moe:1b	smollm2:360m

23	prompt_eval_duration_ns	qwen2:0.5b	smollm2:360m
24	eval_count	gemma3:1b	granite3.1-moe:1b
25	eval_count	gemma3:1b	qwen2:0.5b
26	eval_count	gemma3:1b	smollm2:360m
27	eval_count	granite3.1-moe:1b	qwen2:0.5b
28	eval_count	granite3.1-moe:1b	smollm2:360m
29	eval_count	qwen2:0.5b	smollm2:360m
30	eval_duration_ns	gemma3:1b	granite3.1-moe:1b
31	eval_duration_ns	gemma3:1b	qwen2:0.5b
32	eval_duration_ns	gemma3:1b	smollm2:360m
33	eval_duration_ns	granite3.1-moe:1b	qwen2:0.5b
34	eval_duration_ns	granite3.1-moe:1b	smollm2:360m
35	eval_duration_ns	qwen2:0.5b	smollm2:360m
36	tokens_per_second	gemma3:1b	granite3.1-moe:1b
37	tokens_per_second	gemma3:1b	qwen2:0.5b
38	tokens_per_second	gemma3:1b	smollm2:360m
39	tokens_per_second	granite3.1-moe:1b	qwen2:0.5b
40	tokens_per_second	granite3.1-moe:1b	smollm2:360m
41	tokens_per_second	qwen2:0.5b	smollm2:360m
42	levenshtein_similarity	gemma3:1b	granite3.1-moe:1b
43	levenshtein_similarity	gemma3:1b	qwen2:0.5b
44	levenshtein_similarity	gemma3:1b	smollm2:360m
45	levenshtein_similarity	granite3.1-moe:1b	qwen2:0.5b
46	levenshtein_similarity	granite3.1-moe:1b	smollm2:360m
47	levenshtein_similarity	qwen2:0.5b	smollm2:360m
48	jaccard_similarity	gemma3:1b	granite3.1-moe:1b
49	jaccard_similarity	gemma3:1b	qwen2:0.5b
50	jaccard_similarity	gemma3:1b	smollm2:360m
51	jaccard_similarity	granite3.1-moe:1b	qwen2:0.5b
52	jaccard_similarity	granite3.1-moe:1b	smollm2:360m
53	jaccard_similarity	qwen2:0.5b	smollm2:360m
54	length_ratio	gemma3:1b	granite3.1-moe:1b
55	length_ratio	gemma3:1b	qwen2:0.5b
56	length_ratio	gemma3:1b	smollm2:360m
57	length_ratio	granite3.1-moe:1b	qwen2:0.5b
58	length_ratio	granite3.1-moe:1b	smollm2:360m
59	length_ratio	qwen2:0.5b	smollm2:360m
60	bleu	gemma3:1b	granite3.1-moe:1b
61	bleu	gemma3:1b	qwen2:0.5b
62	bleu	gemma3:1b	smollm2:360m
63	bleu	granite3.1-moe:1b	qwen2:0.5b
64	bleu	granite3.1-moe:1b	smollm2:360m
65	bleu	qwen2:0.5b	smollm2:360m
66	cosine_similarity	gemma3:1b	granite3.1-moe:1b
67	cosine_similarity	gemma3:1b	qwen2:0.5b
68	cosine_similarity	gemma3:1b	smollm2:360m
69	cosine_similarity	granite3.1-moe:1b	qwen2:0.5b
70	cosine_similarity	granite3.1-moe:1b	smollm2:360m

71	cosine_similarity	qwen2:0.5b	smollm2:360m
72	wer	gemma3:1b	granite3.1-moe:1b
73	wer	gemma3:1b	qwen2:0.5b
74	wer	gemma3:1b	smollm2:360m
75	wer	granite3.1-moe:1b	qwen2:0.5b
76	wer	granite3.1-moe:1b	smollm2:360m
77	wer	qwen2:0.5b	smollm2:360m
78	char_diversity	gemma3:1b	granite3.1-moe:1b
79	char_diversity	gemma3:1b	qwen2:0.5b
80	char_diversity	gemma3:1b	smollm2:360m
81	char_diversity	granite3.1-moe:1b	qwen2:0.5b
82	char_diversity	granite3.1-moe:1b	smollm2:360m
83	char_diversity	qwen2:0.5b	smollm2:360m
84	type_token_ratio	gemma3:1b	granite3.1-moe:1b
85	type_token_ratio	gemma3:1b	qwen2:0.5b
86	type_token_ratio	gemma3:1b	smollm2:360m
87	type_token_ratio	granite3.1-moe:1b	qwen2:0.5b
88	type_token_ratio	granite3.1-moe:1b	smollm2:360m
89	type_token_ratio	qwen2:0.5b	smollm2:360m
90	bigram_overlap	gemma3:1b	granite3.1-moe:1b
91	bigram_overlap	gemma3:1b	qwen2:0.5b
92	bigram_overlap	gemma3:1b	smollm2:360m
93	bigram_overlap	granite3.1-moe:1b	qwen2:0.5b
94	bigram_overlap	granite3.1-moe:1b	smollm2:360m
95	bigram_overlap	qwen2:0.5b	smollm2:360m

	mean_model_1	mean_model_2	std_model_1	std_model_2	n_model_1	\
0	1.637960e+10	1.051084e+10	3.626593e+09	3.985375e+09	88	
1	1.637960e+10	7.160551e+09	3.626593e+09	1.498693e+09	88	
2	1.637960e+10	1.105877e+10	3.626593e+09	3.707344e+09	88	
3	1.051084e+10	7.160551e+09	3.985375e+09	1.498693e+09	80	
4	1.051084e+10	1.105877e+10	3.985375e+09	3.707344e+09	80	
5	7.160551e+09	1.105877e+10	1.498693e+09	3.707344e+09	88	
6	3.576625e+08	4.339573e+08	2.094949e+09	3.690988e+09	88	
7	3.576625e+08	7.873249e+07	2.094949e+09	2.526810e+08	88	
8	3.576625e+08	2.220176e+08	2.094949e+09	1.883081e+09	88	
9	4.339573e+08	7.873249e+07	3.690988e+09	2.526810e+08	80	
10	4.339573e+08	2.220176e+08	3.690988e+09	1.883081e+09	80	
11	7.873249e+07	2.220176e+08	2.526810e+08	1.883081e+09	88	
12	6.051136e+01	1.113750e+02	4.651027e+00	7.350837e+00	88	
13	6.051136e+01	6.019318e+01	4.651027e+00	4.977785e+00	88	
14	6.051136e+01	8.168182e+01	4.651027e+00	4.970130e+00	88	
15	1.113750e+02	6.019318e+01	7.350837e+00	4.977785e+00	80	
16	1.113750e+02	8.168182e+01	7.350837e+00	4.970130e+00	80	
17	6.019318e+01	8.168182e+01	4.977785e+00	4.970130e+00	88	
18	3.247818e+09	1.986857e+09	3.790890e+08	3.220527e+08	88	
19	3.247818e+09	1.438463e+09	3.790890e+08	1.865101e+08	88	
20	3.247818e+09	1.446045e+09	3.790890e+08	1.714840e+08	88	

21	1.986857e+09	1.438463e+09	3.220527e+08	1.865101e+08	80
22	1.986857e+09	1.446045e+09	3.220527e+08	1.714840e+08	80
23	1.438463e+09	1.446045e+09	1.865101e+08	1.714840e+08	88
24	5.450000e+01	5.842500e+01	1.127911e+01	5.204124e+00	88
25	5.450000e+01	5.244318e+01	1.127911e+01	1.289735e+01	88
26	5.450000e+01	4.521591e+01	1.127911e+01	1.433703e+01	88
27	5.842500e+01	5.244318e+01	5.204124e+00	1.289735e+01	80
28	5.842500e+01	4.521591e+01	5.204124e+00	1.433703e+01	80
29	5.244318e+01	4.521591e+01	1.289735e+01	1.433703e+01	88
30	1.277295e+10	8.088102e+09	2.695345e+09	7.398549e+08	88
31	1.277295e+10	5.641965e+09	2.695345e+09	1.415511e+09	88
32	1.277295e+10	9.388878e+09	2.695345e+09	3.048139e+09	88
33	8.088102e+09	5.641965e+09	7.398549e+08	1.415511e+09	80
34	8.088102e+09	9.388878e+09	7.398549e+08	3.048139e+09	80
35	5.641965e+09	9.388878e+09	1.415511e+09	3.048139e+09	88
36	4.273367e+00	7.226096e+00	3.561606e-02	4.246905e-02	88
37	4.273367e+00	9.319944e+00	3.561606e-02	1.127637e-01	88
38	4.273367e+00	4.833095e+00	3.561606e-02	5.736037e-02	88
39	7.226096e+00	9.319944e+00	4.246905e-02	1.127637e-01	80
40	7.226096e+00	4.833095e+00	4.246905e-02	5.736037e-02	80
41	9.319944e+00	4.833095e+00	1.127637e-01	5.736037e-02	88
42	3.679158e-01	9.707118e-02	3.630858e-01	1.097015e-01	88
43	3.679158e-01	1.198860e-01	3.630858e-01	1.267990e-01	88
44	3.679158e-01	1.812696e-01	3.630858e-01	1.740149e-01	88
45	9.707118e-02	1.198860e-01	1.097015e-01	1.267990e-01	80
46	9.707118e-02	1.812696e-01	1.097015e-01	1.740149e-01	80
47	1.198860e-01	1.812696e-01	1.267990e-01	1.740149e-01	88
48	3.546248e-01	1.860765e-01	3.107514e-01	1.117166e-01	88
49	3.546248e-01	1.367926e-01	3.107514e-01	8.982222e-02	88
50	3.546248e-01	1.347398e-01	3.107514e-01	1.343297e-01	88
51	1.860765e-01	1.367926e-01	1.117166e-01	8.982222e-02	80
52	1.860765e-01	1.347398e-01	1.117166e-01	1.343297e-01	80
53	1.367926e-01	1.347398e-01	8.982222e-02	1.343297e-01	88
54	1.095413e+00	1.026081e+00	2.090215e-01	1.511895e-01	88
55	1.095413e+00	1.198417e+00	2.090215e-01	3.287786e-01	88
56	1.095413e+00	1.033046e+00	2.090215e-01	3.272457e-01	88
57	1.026081e+00	1.198417e+00	1.511895e-01	3.287786e-01	80
58	1.026081e+00	1.033046e+00	1.511895e-01	3.272457e-01	80
59	1.198417e+00	1.033046e+00	3.287786e-01	3.272457e-01	88
60	3.051032e-01	7.407307e-02	3.461640e-01	9.534950e-02	88
61	3.051032e-01	5.953811e-02	3.461640e-01	8.590007e-02	88
62	3.051032e-01	6.096194e-02	3.461640e-01	1.295520e-01	88
63	7.407307e-02	5.953811e-02	9.534950e-02	8.590007e-02	80
64	7.407307e-02	6.096194e-02	9.534950e-02	1.295520e-01	80
65	5.953811e-02	6.096194e-02	8.590007e-02	1.295520e-01	88
66	5.728996e-01	3.985247e-01	2.586751e-01	1.549181e-01	88
67	5.728996e-01	3.658572e-01	2.586751e-01	1.736955e-01	88
68	5.728996e-01	3.319037e-01	2.586751e-01	1.807902e-01	88

69	3.985247e-01	3.658572e-01	1.549181e-01	1.736955e-01	80
70	3.985247e-01	3.319037e-01	1.549181e-01	1.807902e-01	80
71	3.658572e-01	3.319037e-01	1.736955e-01	1.807902e-01	88
72	8.781957e-01	9.144531e-01	4.386598e-01	2.030163e-01	88
73	8.781957e-01	1.165567e+00	4.386598e-01	2.397264e-01	88
74	8.781957e-01	1.046012e+00	4.386598e-01	2.881699e-01	88
75	9.144531e-01	1.165567e+00	2.030163e-01	2.397264e-01	80
76	9.144531e-01	1.046012e+00	2.030163e-01	2.881699e-01	80
77	1.165567e+00	1.046012e+00	2.397264e-01	2.881699e-01	88
78	2.378138e-01	1.262054e-01	1.270707e-01	8.308520e-02	88
79	2.378138e-01	1.519883e-01	1.270707e-01	9.223367e-02	88
80	2.378138e-01	1.196035e-01	1.270707e-01	8.893092e-02	88
81	1.262054e-01	1.519883e-01	8.308520e-02	9.223367e-02	80
82	1.262054e-01	1.196035e-01	8.308520e-02	8.893092e-02	80
83	1.519883e-01	1.196035e-01	9.223367e-02	8.893092e-02	88
84	8.790624e-01	9.697264e-01	5.318228e-02	2.606863e-02	88
85	8.790624e-01	9.547864e-01	5.318228e-02	3.235066e-02	88
86	8.790624e-01	9.612689e-01	5.318228e-02	3.241982e-02	88
87	9.697264e-01	9.547864e-01	2.606863e-02	3.235066e-02	80
88	9.697264e-01	9.612689e-01	2.606863e-02	3.241982e-02	80
89	9.547864e-01	9.612689e-01	3.235066e-02	3.241982e-02	88
90	2.815986e-01	6.352941e-02	3.447873e-01	8.072161e-02	88
91	2.815986e-01	4.801648e-02	3.447873e-01	6.967926e-02	88
92	2.815986e-01	5.264021e-02	3.447873e-01	1.249316e-01	88
93	6.352941e-02	4.801648e-02	8.072161e-02	6.967926e-02	80
94	6.352941e-02	5.264021e-02	8.072161e-02	1.249316e-01	80
95	4.801648e-02	5.264021e-02	6.967926e-02	1.249316e-01	88

	n_model_2	mean_diff	cohens_d	ttest_pval
0	80	5.868757e+09	1.543774	1.851807e-18
1	88	9.219049e+09	3.322503	2.961777e-43
2	88	5.320826e+09	1.450927	7.612526e-18
3	88	3.350292e+09	1.133510	2.123763e-10
4	88	-5.479307e+08	-0.142610	3.589552e-01
5	88	-3.898222e+09	-1.378640	2.599088e-15
6	80	-7.629477e+07	-0.025743	8.711013e-01
7	88	2.789301e+08	0.186939	2.182100e-01
8	88	1.356449e+08	0.068101	6.520340e-01
9	88	3.552248e+08	0.139150	3.929292e-01
10	88	2.119397e+08	0.073381	6.450656e-01
11	88	-1.432851e+08	-0.106653	4.811095e-01
12	80	-5.086364e+01	-8.355988	1.886816e-90
13	88	3.181818e-01	0.066052	6.618335e-01
14	88	-2.117045e+01	-4.398394	9.215735e-69
15	88	5.118182e+01	8.227185	2.036954e-92
16	88	2.969318e+01	4.775472	1.143564e-62
17	88	-2.148864e+01	-4.320228	8.524955e-68
18	80	1.260961e+09	3.571155	4.332445e-54

19	88	1.809355e+09	6.056565	5.553946e-74
20	88	1.801773e+09	6.124173	1.976495e-72
21	88	5.483942e+08	2.109349	1.038791e-25
22	88	5.408122e+08	2.124975	1.991987e-25
23	88	-7.582020e+06	-0.042321	7.792561e-01
24	80	-3.925000e+00	-0.440030	3.929526e-03
25	88	2.056818e+00	0.169771	2.616884e-01
26	88	9.284091e+00	0.719752	3.956313e-06
27	88	5.981818e+00	0.597979	1.086813e-04
28	88	1.320909e+01	1.202731	8.598580e-13
29	88	7.227273e+00	0.530006	5.608076e-04
30	80	4.684852e+09	2.322765	8.042706e-29
31	88	7.130989e+09	3.312522	7.247881e-46
32	88	3.384076e+09	1.176189	5.715397e-13
33	88	2.446137e+09	2.136695	1.594519e-28
34	88	-1.300776e+09	-0.574309	1.895015e-04
35	88	-3.746913e+09	-1.576699	1.068502e-18
36	80	-2.952729e+00	-75.657094	2.351837e-248
37	88	-5.046577e+00	-60.352309	7.143902e-168
38	88	-5.597273e-01	-11.723850	2.343432e-120
39	88	-2.093848e+00	-24.141375	7.374049e-136
40	88	2.393001e+00	47.087189	1.213217e-223
41	88	4.486849e+00	50.155324	2.266405e-191
42	80	2.708446e-01	0.990177	1.254957e-09
43	88	2.480298e-01	0.912055	2.128155e-08
44	88	1.866462e-01	0.655580	2.813334e-05
45	88	-2.281481e-02	-0.191771	2.130662e-01
46	88	-8.419844e-02	-0.572931	2.215139e-04
47	88	-6.138362e-02	-0.403181	8.268438e-03
48	80	1.685483e-01	0.708777	5.856547e-06
49	88	2.178322e-01	0.952357	7.156117e-09
50	88	2.198850e-01	0.918539	1.420898e-08
51	88	4.928390e-02	0.488751	2.087465e-03
52	88	5.133669e-02	0.413728	7.620339e-03
53	88	2.052788e-03	0.017965	9.052994e-01
54	80	6.933222e-02	0.377251	1.422422e-02
55	88	-1.030037e-01	-0.373898	1.425727e-02
56	88	6.236747e-02	0.227144	1.340200e-01
57	88	-1.723359e-01	-0.663170	2.045843e-05
58	88	-6.964743e-03	-0.026906	8.577026e-01
59	88	1.653711e-01	0.504160	1.010546e-03
60	80	2.310301e-01	0.891690	2.877416e-08
61	88	2.455651e-01	0.973697	4.123433e-09
62	88	2.441413e-01	0.934135	1.008213e-08
63	88	1.453496e-02	0.160571	3.025711e-01
64	88	1.311113e-02	0.114452	4.534489e-01
65	88	-1.423834e-03	-0.012954	9.316383e-01
66	80	1.743749e-01	0.808729	3.295645e-07

67	88	2.070424e-01	0.939730	4.274439e-09
68	88	2.409959e-01	1.079940	2.947308e-11
69	88	3.266754e-02	0.197954	1.993847e-01
70	88	6.662102e-02	0.394272	1.101874e-02
71	88	3.395348e-02	0.191526	2.056260e-01
72	80	-3.625732e-02	-0.104465	4.867596e-01
73	88	-2.873716e-01	-0.812987	3.021430e-07
74	88	-1.678158e-01	-0.452184	3.166699e-03
75	88	-2.511143e-01	-1.126018	8.851207e-12
76	88	-1.315585e-01	-0.523574	7.346167e-04
77	88	1.195558e-01	0.451057	3.188480e-03
78	80	1.116084e-01	1.029719	2.334366e-10
79	88	8.582543e-02	0.773014	8.445384e-07
80	88	1.182103e-01	1.077858	3.172278e-11
81	88	-2.578296e-02	-0.292993	5.836048e-02
82	88	6.601907e-03	0.076590	6.195587e-01
83	88	3.238487e-02	0.357459	1.882925e-02
84	80	-9.066401e-02	-2.133571	3.319137e-28
85	88	-7.572400e-02	-1.720351	7.128313e-22
86	88	-8.220651e-02	-1.866547	2.014489e-24
87	88	1.494001e-02	0.505958	1.152241e-03
88	88	8.457504e-03	0.286037	6.316345e-02
89	88	-6.482510e-03	-0.200169	1.859940e-01
90	80	2.180692e-01	0.852688	9.730077e-08
91	88	2.335821e-01	0.939098	1.306125e-08
92	88	2.289584e-01	0.882943	5.028529e-08
93	88	1.551292e-02	0.206462	1.863744e-01
94	88	1.088920e-02	0.102523	4.995303e-01
95	88	-4.623723e-03	-0.045711	7.621899e-01

[61]: *#What does Cliff's Delta calculate?*
#It measures the probability that a randomly selected value from group A will
↪ be larger than a randomly selected value from group B, minus the reverse
↪ probability.

#Values range from -1 (all values in A are less than all in B) to +1 (all
↪ values in A are greater than all in B).

#0 means the two groups completely overlap (no effect).

#Positive values: A tends to be greater than B.

#Negative values: A tends to be less than B.

Interpretation

```

# Large effect ( $|\Delta| \geq 0.474$ ): Models are substantially different for this
↪ metric

# Medium/small effect: Moderate/small practical differences

# Negligible: Distributions largely overlap

```

```

[62]: import pandas as pd
import numpy as np
from itertools import combinations

# --- Pure Python Cliff's Delta implementation ---
def cliffs_delta(x, y):
    """
    Computes Cliff's delta and magnitude for two arrays/lists x, y.
    Returns delta (float) and magnitude (str: 'negligible', 'small', 'medium',
    ↪ 'large')
    """
    x, y = np.asarray(x), np.asarray(y)
    n_x, n_y = len(x), len(y)
    more = sum(xi > yj for xi in x for yj in y)
    less = sum(xi < yj for xi in x for yj in y)
    delta = (more - less) / (n_x * n_y)
    # Magnitude thresholds (Romano et al., 2006)
    adelta = abs(delta)
    if adelta < 0.147:
        magnitude = "negligible"
    elif adelta < 0.33:
        magnitude = "small"
    elif adelta < 0.474:
        magnitude = "medium"
    else:
        magnitude = "large"
    return delta, magnitude

# --- Specify metrics for which to compute Cliff's delta ---
metrics = [
    "total_duration_ns", "load_duration_ns", "prompt_eval_count",
    ↪ "prompt_eval_duration_ns", "eval_count",
    "eval_duration_ns", "tokens_per_second", "levenshtein_similarity",
    ↪ "jaccard_similarity", "length_ratio",
    "bleu", "cosine_similarity", "wer", "char_diversity", "type_token_ratio",
    ↪ "bigram_overlap"
]

```

```

# --- Get unique models ---
model_list = sorted(df['model'].unique())

# --- Compute Cliff's delta for all pairs and all metrics ---
results = []
for metric in metrics:
    for m1, m2 in combinations(model_list, 2):
        x = df[df['model'] == m1][metric].dropna()
        y = df[df['model'] == m2][metric].dropna()
        if len(x) > 1 and len(y) > 1:
            delta, magnitude = cliffs_delta(x, y)
            results.append({
                'metric': metric,
                'model_1': m1,
                'model_2': m2,
                "cliffs_delta": delta,
                "magnitude": magnitude,
                "mean_model_1": x.mean(),
                "mean_model_2": y.mean(),
                "n_model_1": len(x),
                "n_model_2": len(y)
            })

cliffs_df = pd.DataFrame(results)

# Show all rows (optional, for Jupyter)
pd.set_option("display.max_rows", None)
print("Table: Cliff's Delta Effect Sizes for Pairwise Model Comparisons Across_
↳All Metrics")
display(cliffs_df)

```

Table: Cliff's Delta Effect Sizes for Pairwise Model Comparisons Across All Metrics

	metric	model_1	model_2 \
0	total_duration_ns	gemma3:1b	granite3.1-moe:1b
1	total_duration_ns	gemma3:1b	qwen2:0.5b
2	total_duration_ns	gemma3:1b	smollm2:360m
3	total_duration_ns	granite3.1-moe:1b	qwen2:0.5b
4	total_duration_ns	granite3.1-moe:1b	smollm2:360m
5	total_duration_ns	qwen2:0.5b	smollm2:360m
6	load_duration_ns	gemma3:1b	granite3.1-moe:1b
7	load_duration_ns	gemma3:1b	qwen2:0.5b
8	load_duration_ns	gemma3:1b	smollm2:360m
9	load_duration_ns	granite3.1-moe:1b	qwen2:0.5b
10	load_duration_ns	granite3.1-moe:1b	smollm2:360m
11	load_duration_ns	qwen2:0.5b	smollm2:360m
12	prompt_eval_count	gemma3:1b	granite3.1-moe:1b

13	prompt_eval_count	gemma3:1b	qwen2:0.5b
14	prompt_eval_count	gemma3:1b	smollm2:360m
15	prompt_eval_count	granite3.1-moe:1b	qwen2:0.5b
16	prompt_eval_count	granite3.1-moe:1b	smollm2:360m
17	prompt_eval_count	qwen2:0.5b	smollm2:360m
18	prompt_eval_duration_ns	gemma3:1b	granite3.1-moe:1b
19	prompt_eval_duration_ns	gemma3:1b	qwen2:0.5b
20	prompt_eval_duration_ns	gemma3:1b	smollm2:360m
21	prompt_eval_duration_ns	granite3.1-moe:1b	qwen2:0.5b
22	prompt_eval_duration_ns	granite3.1-moe:1b	smollm2:360m
23	prompt_eval_duration_ns	qwen2:0.5b	smollm2:360m
24	eval_count	gemma3:1b	granite3.1-moe:1b
25	eval_count	gemma3:1b	qwen2:0.5b
26	eval_count	gemma3:1b	smollm2:360m
27	eval_count	granite3.1-moe:1b	qwen2:0.5b
28	eval_count	granite3.1-moe:1b	smollm2:360m
29	eval_count	qwen2:0.5b	smollm2:360m
30	eval_duration_ns	gemma3:1b	granite3.1-moe:1b
31	eval_duration_ns	gemma3:1b	qwen2:0.5b
32	eval_duration_ns	gemma3:1b	smollm2:360m
33	eval_duration_ns	granite3.1-moe:1b	qwen2:0.5b
34	eval_duration_ns	granite3.1-moe:1b	smollm2:360m
35	eval_duration_ns	qwen2:0.5b	smollm2:360m
36	tokens_per_second	gemma3:1b	granite3.1-moe:1b
37	tokens_per_second	gemma3:1b	qwen2:0.5b
38	tokens_per_second	gemma3:1b	smollm2:360m
39	tokens_per_second	granite3.1-moe:1b	qwen2:0.5b
40	tokens_per_second	granite3.1-moe:1b	smollm2:360m
41	tokens_per_second	qwen2:0.5b	smollm2:360m
42	levenshtein_similarity	gemma3:1b	granite3.1-moe:1b
43	levenshtein_similarity	gemma3:1b	qwen2:0.5b
44	levenshtein_similarity	gemma3:1b	smollm2:360m
45	levenshtein_similarity	granite3.1-moe:1b	qwen2:0.5b
46	levenshtein_similarity	granite3.1-moe:1b	smollm2:360m
47	levenshtein_similarity	qwen2:0.5b	smollm2:360m
48	jaccard_similarity	gemma3:1b	granite3.1-moe:1b
49	jaccard_similarity	gemma3:1b	qwen2:0.5b
50	jaccard_similarity	gemma3:1b	smollm2:360m
51	jaccard_similarity	granite3.1-moe:1b	qwen2:0.5b
52	jaccard_similarity	granite3.1-moe:1b	smollm2:360m
53	jaccard_similarity	qwen2:0.5b	smollm2:360m
54	length_ratio	gemma3:1b	granite3.1-moe:1b
55	length_ratio	gemma3:1b	qwen2:0.5b
56	length_ratio	gemma3:1b	smollm2:360m
57	length_ratio	granite3.1-moe:1b	qwen2:0.5b
58	length_ratio	granite3.1-moe:1b	smollm2:360m
59	length_ratio	qwen2:0.5b	smollm2:360m
60	bleu	gemma3:1b	granite3.1-moe:1b

61		bleu	gemma3:1b	qwen2:0.5b
62		bleu	gemma3:1b	smollm2:360m
63		bleu	granite3.1-moe:1b	qwen2:0.5b
64		bleu	granite3.1-moe:1b	smollm2:360m
65		bleu	qwen2:0.5b	smollm2:360m
66	cosine_similarity		gemma3:1b	granite3.1-moe:1b
67	cosine_similarity		gemma3:1b	qwen2:0.5b
68	cosine_similarity		gemma3:1b	smollm2:360m
69	cosine_similarity		granite3.1-moe:1b	qwen2:0.5b
70	cosine_similarity		granite3.1-moe:1b	smollm2:360m
71	cosine_similarity		qwen2:0.5b	smollm2:360m
72		wer	gemma3:1b	granite3.1-moe:1b
73		wer	gemma3:1b	qwen2:0.5b
74		wer	gemma3:1b	smollm2:360m
75		wer	granite3.1-moe:1b	qwen2:0.5b
76		wer	granite3.1-moe:1b	smollm2:360m
77		wer	qwen2:0.5b	smollm2:360m
78	char_diversity		gemma3:1b	granite3.1-moe:1b
79	char_diversity		gemma3:1b	qwen2:0.5b
80	char_diversity		gemma3:1b	smollm2:360m
81	char_diversity		granite3.1-moe:1b	qwen2:0.5b
82	char_diversity		granite3.1-moe:1b	smollm2:360m
83	char_diversity		qwen2:0.5b	smollm2:360m
84	type_token_ratio		gemma3:1b	granite3.1-moe:1b
85	type_token_ratio		gemma3:1b	qwen2:0.5b
86	type_token_ratio		gemma3:1b	smollm2:360m
87	type_token_ratio		granite3.1-moe:1b	qwen2:0.5b
88	type_token_ratio		granite3.1-moe:1b	smollm2:360m
89	type_token_ratio		qwen2:0.5b	smollm2:360m
90	bigram_overlap		gemma3:1b	granite3.1-moe:1b
91	bigram_overlap		gemma3:1b	qwen2:0.5b
92	bigram_overlap		gemma3:1b	smollm2:360m
93	bigram_overlap		granite3.1-moe:1b	qwen2:0.5b
94	bigram_overlap		granite3.1-moe:1b	smollm2:360m
95	bigram_overlap		qwen2:0.5b	smollm2:360m

	cliffs_delta	magnitude	mean_model_1	mean_model_2	n_model_1	n_model_2
0	0.809943	large	1.637960e+10	1.051084e+10	88	80
1	0.993027	large	1.637960e+10	7.160551e+09	88	88
2	0.778151	large	1.637960e+10	1.105877e+10	88	88
3	0.934943	large	1.051084e+10	7.160551e+09	80	88
4	-0.219034	small	1.051084e+10	1.105877e+10	80	88
5	-0.594525	large	7.160551e+09	1.105877e+10	88	88
6	0.975000	large	3.576625e+08	4.339573e+08	88	80
7	0.977531	large	3.576625e+08	7.873249e+07	88	88
8	0.977531	large	3.576625e+08	2.220176e+08	88	88
9	-0.975000	large	4.339573e+08	7.873249e+07	80	88
10	-0.015909	negligible	4.339573e+08	2.220176e+08	80	88

11	0.977273	large	7.873249e+07	2.220176e+08	88	88
12	-1.000000	large	6.051136e+01	1.113750e+02	88	80
13	0.030088	negligible	6.051136e+01	6.019318e+01	88	88
14	-0.999742	large	6.051136e+01	8.168182e+01	88	88
15	1.000000	large	1.113750e+02	6.019318e+01	80	88
16	1.000000	large	1.113750e+02	8.168182e+01	80	88
17	-1.000000	large	6.019318e+01	8.168182e+01	88	88
18	0.978125	large	3.247818e+09	1.986857e+09	88	80
19	1.000000	large	3.247818e+09	1.438463e+09	88	88
20	1.000000	large	3.247818e+09	1.446045e+09	88	88
21	0.906250	large	1.986857e+09	1.438463e+09	80	88
22	0.935227	large	1.986857e+09	1.446045e+09	80	88
23	-0.059143	negligible	1.438463e+09	1.446045e+09	88	88
24	-0.100000	negligible	5.450000e+01	5.842500e+01	88	80
25	0.121772	negligible	5.450000e+01	5.244318e+01	88	88
26	0.428848	medium	5.450000e+01	4.521591e+01	88	88
27	0.237358	small	5.842500e+01	5.244318e+01	80	88
28	0.580966	large	5.842500e+01	4.521591e+01	80	88
29	0.324768	small	5.244318e+01	4.521591e+01	88	88
30	0.803693	large	1.277295e+10	8.088102e+09	88	80
31	0.887138	large	1.277295e+10	5.641965e+09	88	88
32	0.715909	large	1.277295e+10	9.388878e+09	88	88
33	0.909943	large	8.088102e+09	5.641965e+09	80	88
34	-0.284091	small	8.088102e+09	9.388878e+09	80	88
35	-0.608471	large	5.641965e+09	9.388878e+09	88	88
36	-1.000000	large	4.273367e+00	7.226096e+00	88	80
37	-1.000000	large	4.273367e+00	9.319944e+00	88	88
38	-1.000000	large	4.273367e+00	4.833095e+00	88	88
39	-1.000000	large	7.226096e+00	9.319944e+00	80	88
40	1.000000	large	7.226096e+00	4.833095e+00	80	88
41	1.000000	large	9.319944e+00	4.833095e+00	88	88
42	0.417898	medium	3.679158e-01	9.707118e-02	88	80
43	0.362087	medium	3.679158e-01	1.198860e-01	88	88
44	0.196152	small	3.679158e-01	1.812696e-01	88	88
45	-0.105256	negligible	9.707118e-02	1.198860e-01	80	88
46	-0.385085	medium	9.707118e-02	1.812696e-01	80	88
47	-0.251033	small	1.198860e-01	1.812696e-01	88	88
48	0.249148	small	3.546248e-01	1.860765e-01	88	80
49	0.462939	medium	3.546248e-01	1.367926e-01	88	88
50	0.522082	large	3.546248e-01	1.347398e-01	88	88
51	0.292330	small	1.860765e-01	1.367926e-01	80	88
52	0.377273	medium	1.860765e-01	1.347398e-01	80	88
53	0.098011	negligible	1.367926e-01	1.347398e-01	88	88
54	0.263494	small	1.095413e+00	1.026081e+00	88	80
55	-0.307076	small	1.095413e+00	1.198417e+00	88	88
56	0.133135	negligible	1.095413e+00	1.033046e+00	88	88
57	-0.453409	medium	1.026081e+00	1.198417e+00	80	88
58	-0.008381	negligible	1.026081e+00	1.033046e+00	80	88

59	0.319473	small	1.198417e+00	1.033046e+00	88	88
60	0.381960	medium	3.051032e-01	7.407307e-02	88	80
61	0.506586	large	3.051032e-01	5.953811e-02	88	88
62	0.537319	large	3.051032e-01	6.096194e-02	88	88
63	0.231534	small	7.407307e-02	5.953811e-02	80	88
64	0.279403	small	7.407307e-02	6.096194e-02	80	88
65	0.048425	negligible	5.953811e-02	6.096194e-02	88	88
66	0.388068	medium	5.728996e-01	3.985247e-01	88	80
67	0.510072	large	5.728996e-01	3.658572e-01	88	88
68	0.584065	large	5.728996e-01	3.319037e-01	88	88
69	0.177273	small	3.985247e-01	3.658572e-01	80	88
70	0.297301	small	3.985247e-01	3.319037e-01	80	88
71	0.131973	negligible	3.658572e-01	3.319037e-01	88	88
72	0.149290	small	8.781957e-01	9.144531e-01	88	80
73	-0.323735	small	8.781957e-01	1.165567e+00	88	88
74	-0.070377	negligible	8.781957e-01	1.046012e+00	88	88
75	-0.598580	large	9.144531e-01	1.165567e+00	80	88
76	-0.296733	small	9.144531e-01	1.046012e+00	80	88
77	0.299458	small	1.165567e+00	1.046012e+00	88	88
78	0.514773	large	2.378138e-01	1.262054e-01	88	80
79	0.427686	medium	2.378138e-01	1.519883e-01	88	88
80	0.525956	large	2.378138e-01	1.196035e-01	88	88
81	-0.140341	negligible	1.262054e-01	1.519883e-01	80	88
82	0.097869	negligible	1.262054e-01	1.196035e-01	80	88
83	0.228564	small	1.519883e-01	1.196035e-01	88	88
84	-0.903835	large	8.790624e-01	9.697264e-01	88	80
85	-0.819473	large	8.790624e-01	9.547864e-01	88	88
86	-0.845558	large	8.790624e-01	9.612689e-01	88	88
87	0.266051	small	9.697264e-01	9.547864e-01	80	88
88	0.148438	small	9.697264e-01	9.612689e-01	80	88
89	-0.113895	negligible	9.547864e-01	9.612689e-01	88	88
90	0.371165	medium	2.815986e-01	6.352941e-02	88	80
91	0.477402	large	2.815986e-01	4.801648e-02	88	88
92	0.509168	large	2.815986e-01	5.264021e-02	88	88
93	0.175142	small	6.352941e-02	4.801648e-02	80	88
94	0.225426	small	6.352941e-02	5.264021e-02	80	88
95	0.051911	negligible	4.801648e-02	5.264021e-02	88	88

[63]: *#Association between Text Augmentation with Prompts*

```
import pandas as pd
import numpy as np
from scipy.stats import chi2_contingency
from sklearn.metrics import mutual_info_score

# Optional: If you don't have `theils_u` function, define it:
def theils_u(x, y):
```

```

"""Theil's U (uncertainty coefficient  $U(x|y)$ ), asymmetric"""
s_xy = mutual_info_score(x, y)
s_x = entropy_from_series(x)
return 0 if s_x == 0 else s_xy / s_x

def entropy_from_series(s):
    p = s.value_counts(normalize=True)
    return -np.sum(p * np.log2(p + 1e-12))

# Build contingency table
contingency = pd.crosstab(df['augmentation_type'], df['prompt_id'])

# 1. Chi-square Test
chi2_stat, p_chi2, dof, expected = chi2_contingency(contingency)

# 2. Cramér's V
n = contingency.values.sum()
cramers_v = np.sqrt(chi2_stat / (n * (min(contingency.shape) - 1)))

# 3. Mutual Information Score
mi = mutual_info_score(df['augmentation_type'], df['prompt_id'])

# 4. Theil's U (both directions, asymmetric)
u_aug_given_prompt = theils_u(df['augmentation_type'], df['prompt_id'])
u_prompt_given_aug = theils_u(df['prompt_id'], df['augmentation_type'])

# 5. Summary Table
summary = pd.DataFrame({
    'Test': [
        "Chi-square Test of Independence",
        "Cramér's V (Effect Size)",
        "Mutual Information Score",
        "Theil's U (Augmentation|Prompt)",
        "Theil's U (Prompt|Augmentation)"
    ],
    'Statistic/Value': [
        f"Chi2={chi2_stat:.5g}, dof={dof}, p={p_chi2:.5g}",
        f"{cramers_v:.5f}",
        f"{mi:.5f}",
        f"{u_aug_given_prompt:.5f}",
        f"{u_prompt_given_aug:.5f}"
    ]
})

print("Association between Text Augmentation with Prompts")
display(summary)

```

Association between Text Augmentation with Prompts

	Test	Statistic/Value
0	Chi-square Test of Independence	Chi2=0.86154, dof=70, p=1
1	Cramér's V (Effect Size)	0.01892
2	Mutual Information Score	0.00125
3	Theil's U (Augmentation Prompt)	0.00036
4	Theil's U (Prompt Augmentation)	0.00042

```
[69]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import seaborn as sns

# 2. Select Numeric Columns for PCA
numeric_cols = [
    'total_duration_ns', 'load_duration_ns', 'prompt_eval_count',
    'prompt_eval_duration_ns', 'eval_count', 'eval_duration_ns',
    'tokens_per_second', 'levenshtein_similarity', 'jaccard_similarity',
    'length_ratio', 'bleu', 'cosine_similarity', 'wer', 'char_diversity',
    'type_token_ratio', 'bigram_overlap'
]
df_numeric = df[numeric_cols].dropna()

# 3. Standardize
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df_numeric)

# 4. PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
df_pca = pd.DataFrame(X_pca, columns=['PC1', 'PC2'])

# 5. KMeans
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
clusters = kmeans.fit_predict(X_pca)
df_pca['Cluster'] = clusters

# 6. Plotting
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df_pca, x='PC1', y='PC2', hue='Cluster', palette='Set2',
               s=80)
plt.title('KMeans Clustering on PCA-Reduced Augmentation Metrics')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.grid(True)
plt.tight_layout()
```

```

plt.show()

# Print the cluster assignments for the first few rows and the centroids
print("Cluster assignments (first 10):")
print(df_pca['Cluster'].head(10))


print("\nKMeans cluster centers in PCA space:")
print(kmeans.cluster_centers_)

# Count the number of points in each cluster
print("\nNumber of points in each cluster:")
print(df_pca['Cluster'].value_counts())

loadings = pd.DataFrame(
    pca.components_.T, # shape: (n_features, n_components)
    index=numeric_cols,
    columns=['PC1', 'PC2']
)
print(loadings)

```

C:\Users\parth\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1436:
 UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
 there are less chunks than available threads. You can avoid it by setting the
 environment variable OMP_NUM_THREADS=2.
 warnings.warn(



Cluster assignments (first 10):

```
0    0
1    0
2    0
3    0
4    0
5    0
6    0
7    0
8    0
9    0
```

Name: Cluster, dtype: int32

KMeans cluster centers in PCA space:

```
[[-1.10796476 -0.66622526]
 [ 6.82928688 -1.79858054]
 [ 0.9291741   2.67848547]]
```

Number of points in each cluster:

```
Cluster
0      238
```

```

2      78
1      28
Name: count, dtype: int64

```

	PC1	PC2
total_duration_ns	0.226049	0.368875
load_duration_ns	0.028109	0.105715
prompt_eval_count	-0.091955	-0.058772
prompt_eval_duration_ns	0.271370	0.210033
eval_count	0.062268	0.354597
eval_duration_ns	0.230526	0.390617
tokens_per_second	-0.216960	-0.162517
levenshtein_similarity	0.313847	-0.260757
jaccard_similarity	0.370734	-0.155544
length_ratio	0.048719	0.257775
bleu	0.374915	-0.149463
cosine_similarity	0.342834	-0.069804
wer	-0.272953	0.314114
char_diversity	0.010643	0.364538
type_token_ratio	-0.228997	-0.236165
bigram_overlap	0.372093	-0.165159

```
[ ]:
```